

Synthetic Data Compliance Framework (SDCF), Version 1.95

A Purpose-Bounded Methodology for Assessing Privacy, Fidelity, and Fairness in Synthetic Data

Wayne Kearns
Kaionix Labs & Nortec Consulting, Ireland
`wayne.kearns@nortecconsulting.com`

December 2025

Abstract

Synthetic data is increasingly adopted to enable privacy-preserving analytics, data sharing, and artificial intelligence (AI) model development in regulated environments. However, organisations lack standardised, audit-ready methodologies for determining whether a given synthetic dataset satisfies regulatory expectations for privacy protection, statistical utility, and algorithmic fairness. This paper presents the *Synthetic Data Compliance Framework* (SDCF), a purpose-bounded, three-pillar assessment methodology that connects quantitative technical metrics to regulatory requirements under the GDPR, the EU AI Act, and relevant ISO/IEC and NIST standards.

SDCF introduces a tiered assessment architecture (Gold, Silver, Bronze) based on the degree of source data accessibility, together with composite metrics for Privacy Risk (PRS), Fidelity (FI), and Fairness Variance (FV). Version 1.95 reports a preliminary retrospective empirical evaluation of the Bronze Tier methodology across ten heterogeneous synthetic datasets spanning healthcare, demographics, e-commerce, AI training, and business domains. The results provide indicative support for conservative privacy risk classification, meaningful quality discrimination under source-data-absent conditions, and cross-domain applicability, with observed B-PRS values ranging from 0.09 to 0.79 and B-FI consistently exceeding 0.92. A comparative synthesis experiment shows lower observed privacy risk for TVAE relative to CTGAN under equivalent fidelity.

The framework further specifies formal mathematical definitions, provisional thresholds, governance control sets, assessment workflows, regulatory mapping tables, and reference Python implementations to support reproducibility and auditability. SDCF is intended for data protection officers, AI governance teams, data scientists, legal practitioners, and regulators seeking structured, purpose-bounded validation of synthetic data in compliance-critical contexts.

Keywords: Synthetic Data, Privacy Assessment, GDPR Compliance, EU AI Act, Fairness Metrics, Data Protection, Regulatory Technology, AI Governance

Contents

0.1	Version 1.95	10
0.2	Copyright and Licensing	10
0.3	Document Status	11
0.4	Table of Contents	11
0.4.1	Core Framework	11
0.4.2	Technical Appendices	12
0.4.3	Supporting Materials	13
1	Introduction	13
1.1	What is SDCF?	13
1.2	Why SDCF is Needed	13
1.2.1	The Synthetic Data Promise and Challenge	13
1.2.2	The Standards Gap	14
1.2.3	Empirical Validation	14
1.3	Research Contributions	15
1.4	Scope and Limitations	15
1.4.1	What SDCF Covers	15
1.4.2	What SDCF Does NOT Do	16
1.5	Target Audience	16
1.6	How to Use This Document	17
1.6.1	Reading Paths by Role	17
1.6.2	Document Conventions	17
1.7	Document Structure	17
2	Related Work	18
2.1	Synthetic Data Generation and Evaluation	18
2.2	Privacy Risks and Attacks	18
2.3	Fairness, Representation, and Responsible AI	18
2.4	Governance, Regulation, and Standards	19
2.5	Positioning of SDCF	19
3	Regulatory Context	19
3.1	GDPR and Synthetic Data	19
3.1.1	The Anonymisation Question	19
3.1.2	GDPR Does Not Provide a Simple Answer	20
3.1.3	EDPB Guidelines 01/2025 on Pseudonymisation	20
3.1.4	GDPR Articles Implicated in Synthetic Data Use	21
3.1.5	Practical Implications for Organisations	21
3.2	EU AI Act Article 10	22
3.2.1	Article 10: Data and Data Governance	22
3.2.2	Synthetic Data as Article 10 Compliance Strategy	22
3.2.3	What Article 10 Does NOT Specify	22
3.2.4	SDCF as Article 10 Implementation Framework	22
3.3	Sector-Specific Standards	23
3.3.1	Healthcare	23
3.3.2	Financial Services	23

3.3.3	Insurance	24
3.3.4	Public Sector	24
3.4	International Alignment	24
3.4.1	ISO/IEC Standards	24
3.4.2	NIST Frameworks	25
3.4.3	UK / Singapore / Other Jurisdictions	25
3.5	Use Cases Covered	25
3.5.1	Use Case 1: Internal Analytics and Reporting	25
3.5.2	Use Case 2: External Data Sharing and Collaboration	25
3.5.3	Use Case 3: AI/ML Model Training	26
3.5.4	Use Case 4: Software Testing and Development	26
3.5.5	Use Case 5: Regulatory Reporting and Compliance	26
3.5.6	Use Case 6: Training and Education	27
3.5.7	Purpose-Bounded Assessment in Practice	27
4	SDCF Architecture	27
4.1	Purpose-Bounded Assessment Philosophy	27
4.1.1	The Fitness-for-Purpose Principle	27
4.1.2	How Purpose-Bounded Assessment Works	28
4.2	The Three Pillars: Privacy, Fidelity, Fairness	29
4.2.1	Pillar 1: Privacy Risk	29
4.2.2	Pillar 2: Statistical Fidelity	29
4.2.3	Pillar 3: Algorithmic Fairness	30
4.2.4	The Three-Pillar Balance	30
4.3	Assessment Tiers: Gold, Silver, Bronze	31
4.3.1	Gold Tier: Full Source Data Access	31
4.3.2	Silver Tier: Partial Source Data Access	32
4.3.3	Bronze Tier: No Source Data Access (Synthetic-Only)	32
4.3.4	Tier Selection Decision Tree	33
4.4	Conformance Levels: SDCF-A, SDCF-P, SDCF-R	33
4.4.1	SDCF-A (Approved)	34
4.4.2	SDCF-P (Provisional)	34
4.4.3	SDCF-R (Restricted)	35
4.4.4	Conformance Level Decision Matrix	35
4.5	Integration Points	36
4.5.1	Integration with Data Governance	36
4.5.2	Integration with Risk Management	36
4.5.3	Integration with AI Governance	36
4.5.4	Integration with Vendor Management	36
4.5.5	Integration with Compliance and Audit	37
4.5.6	Integration with Existing Tools	37
4.5.7	Standards Alignment	37
5	Control Sets (C1-C7)	38
5.1	C1: Purpose Sheet	38
5.1.1	What It Is	38
5.1.2	Why It Matters	38
5.1.3	Required Elements	38

5.1.4	Implementation Guidance	39
5.1.5	Common Mistakes to Avoid	39
5.2	C2: Governance Record	39
5.2.1	What It Is	39
5.2.2	Why It Matters	40
5.2.3	Required Elements	40
5.2.4	Implementation Guidance	41
5.2.5	Example Governance Record Entry	41
5.3	C3: Privacy Risk Testing	41
5.3.1	What It Is	41
5.3.2	Why It Matters	41
5.3.3	Required Tests	42
5.3.4	Composite Privacy Risk Score (PRS)	43
5.3.5	Implementation Guidance	43
5.4	C4: Fidelity Testing	43
5.4.1	What It Is	43
5.4.2	Why It Matters	43
5.4.3	Required Tests	43
5.4.4	Composite Fidelity Index (FI)	44
5.4.5	Implementation Guidance	45
5.5	C5: Fairness Assessment	45
5.5.1	What It Is	45
5.5.2	Why It Matters	45
5.5.3	Required Tests	45
5.5.4	Composite Fairness Variance (FV)	46
5.5.5	Implementation Guidance	46
5.6	C6: Transparency Pack	47
5.6.1	What It Is	47
5.6.2	Why It Matters	47
5.6.3	Required Components	47
5.6.4	Implementation Guidance	47
5.6.5	Example Transparency Pack Structure	48
5.7	C7: Release Rules	49
5.7.1	What It Is	49
5.7.2	Why It Matters	49
5.7.3	Required Elements	49
5.7.4	Implementation Guidance	50
5.7.5	Example Release Rules	50
6	Assessment Process	51
6.1	Five-Step Workflow	51
6.1.1	Overview	51
6.1.2	Step 1: Define Purpose	52
6.1.3	Step 2: Select Tier	52
6.1.4	Step 3: Execute Tests	53
6.1.5	Step 4: Evaluate Results	54
6.1.6	Step 5: Issue Certificate	56
6.2	Total Assessment Timeline	57

6.3	Roles and Responsibilities	57
6.4	Documentation and Audit Trail	57
6.5	Common Challenges and Mitigations	58
7	Relationship to Existing Tools and Standards	59
7.1	Complementary Open-Source Tools	59
7.1.1	SDMetrics / SDV (Synthetic Data Vault)	59
7.1.2	mostlyai-qa (MOSTLY AI Quality Assurance)	60
7.1.3	Other Relevant Tools	60
7.2	Vendor Platform Integration	60
7.2.1	Gretel.ai	60
7.2.2	MOSTLY AI	61
7.2.3	Syntho	61
7.2.4	General Vendor Integration Pattern	61
7.3	Standards Alignment	61
7.3.1	ISO/IEC 27001:2022 (Information Security Management)	61
7.3.2	ISO/IEC 27701:2019 (Privacy Information Management)	62
7.3.3	ISO/IEC 23894:2023 (AI Risk Management)	62
7.3.4	ISO/IEC 42001:2023 (AI Management System)	62
7.3.5	NIST Privacy Framework (2020)	62
7.3.6	NIST AI Risk Management Framework (2023)	62
7.4	What SDCF Adds to the Ecosystem	63
7.4.1	The Gap SDCF Fills	63
7.4.2	SDCF Value Proposition Summary	63
7.4.3	How to Use SDCF in Practice	63
8	Empirical Validation Study: Bronze Tier Retrospective Evaluation	65
8.1	Study Design and Motivation	65
8.1.1	Validation Objective	65
8.1.2	Why Validate Bronze Tier First	65
8.2	Dataset Portfolio	66
8.2.1	Selection Criteria	66
8.2.2	Portfolio Summary	66
8.3	Methodology	67
8.3.1	Assessment Procedure	67
8.3.2	Implementation Details	67
8.4	Results	68
8.4.1	Conformance Distribution	68
8.4.2	Privacy Risk Score (B-PRS) Performance	69
8.4.3	Fidelity Index (B-FI) Performance	69
8.4.4	Fairness Variance (B-FV) Performance	70
8.4.5	Cross-Domain Patterns	71
8.4.6	Synthesis Method Comparison	72
8.5	Framework Alignment Assessment	73
8.6	Key Findings	74
8.7	Limitations	74
8.8	Implications for Practitioners	75
8.8.1	When to Use Bronze Tier	75

8.8.2	Expected Performance Ranges	76
8.8.3	Synthesis Method Recommendations	76
8.8.4	Threshold Interpretation Guidance	77
8.8.5	Community Validation Invitation	77
9	APPENDICES	78
10	Appendix A: Mathematical Definitions	79
10.1	A.1 Privacy Risk Score (PRS)	79
10.1.1	Overview	79
10.1.2	Component 1: Membership Inference Risk (MIR)	79
10.1.3	Component 2: Record Similarity Risk (RSR)	80
10.1.4	Component 3: Attribute Disclosure Risk (ADR)	81
10.1.5	Composite Privacy Risk Score (PRS)	82
10.1.6	Confidence Intervals	83
10.1.7	Provisional Thresholds	83
10.2	A.2 Fidelity Index (FI)	84
10.2.1	Overview	84
10.2.2	Component 1: Distribution Similarity (DS)	84
10.2.3	Component 2: Dependency Preservation (DP)	85
10.2.4	Component 3: Predictive Utility (PU)	86
10.2.5	Composite Fidelity Index (FI)	87
10.2.6	Confidence Intervals	88
10.2.7	Provisional Thresholds	88
10.3	A.3 Fairness Variance (FV)	88
10.3.1	Overview	88
10.3.2	Component 1: Representation Variance (RV)	88
10.3.3	Component 2: Predictive Parity Violation (PPV)	89
10.3.4	Composite Fairness Variance (FV)	91
10.3.5	Confidence Intervals	91
10.3.6	Provisional Thresholds	91
10.4	A.4 Normalization and Weighting Rationale	92
10.4.1	Why Composite Metrics?	92
10.4.2	Default Weight Selection	92
10.4.3	Normalization Approach	92
10.5	A.5 Handling Missing Components	92
10.6	A.6 Conflicting Signals	92
11	Appendix B: Legal and Regulatory Disclaimers	93
11.1	B.1 Synthetic Data and Anonymisation	93
11.2	B.1a EU AI Act Mapping	93
11.2.1	Critical Disclaimer	94
11.2.2	What SDCF Does vs. What It Does Not Do	94
11.2.3	Recommended Legal Review Points	94
11.2.4	Liability Limitation	95
11.3	B.2 GDPR Interpretation Guidance	95
11.3.1	The Anonymisation Question	95
11.3.2	Recital 26 Analysis	95

11.3.3	How SDCF Supports GDPR Analysis	95
11.3.4	EDPB Guidelines 01/2025 Implications	96
11.3.5	Practical GDPR Compliance Posture	96
11.4	B.3 EU AI Act Article 10 Mapping	97
11.4.1	Article 10 Requirements	97
11.4.2	How SDCF Satisfies Article 10	97
11.4.3	What SDCF Does NOT Cover	98
11.5	B.4 Scope Limitations	98
11.5.1	What SDCF Assesses	98
11.5.2	What SDCF Does NOT Assess	98
11.6	B.5 Liability Framework and Indemnification	98
11.7	B.6 Warranty Disclaimers	99
11.8	B.7 Jurisdiction and Regulatory Change	99
11.9	B.8 Acknowledgment	99
12	Appendix C: Bronze Tier Guidance	100
12.1	C.1 Synthetic-Only Assessment Methodology	100
12.1.1	Why Bronze Tier Matters	100
12.1.2	Bronze Tier Philosophy	100
12.1.3	When to Use Bronze Tier	100
12.2	C.2 B-PRS: Privacy Risk Without Source Data	101
12.2.1	Component 1: B-MIR (Membership Inference Proxy)	101
12.2.2	Component 2: B-RSR (Record Similarity Proxy)	102
12.2.3	Component 3: B-ADR (Attribute Disclosure Proxy)	102
12.2.4	Composite B-PRS	103
12.3	C.3 B-FI: Fidelity Without Source Data	104
12.3.1	Component 1: B-DS (Distribution Validation)	104
12.3.2	Component 2: B-DP (Dependency Validation)	105
12.3.3	Component 3: B-PU (Predictive Utility Estimation)	106
12.3.4	Composite B-FI	107
12.4	C.4 B-FV: Fairness Assessment	107
12.4.1	Good News: Fairness is More Tractable in Bronze	107
12.4.2	B-RV: Representation Variance	107
12.4.3	B-FV: No Predictive Parity	108
12.5	C.5 Bronze Tier Certificate Templates	108
12.5.1	Template 1: Bronze Tier SDCF-P Certificate	108
12.5.2	Template 2: Bronze Tier SDCF-R Certificate	110
12.6	C.6 Risk Statement Examples	112
12.6.1	Risk Statement: Bronze Tier SDCF-P	112
12.7	C.7 Bronze Tier Case Studies	113
12.7.1	Case Study 1: AI Training Data Pre-Check (Success)	113
12.7.2	Case Study 2: Legacy Data Retirement (Failure Detected)	114
12.7.3	Case Study 3: Open-Source AI Training Dataset (Qualified Approval)	114
12.8	C.8 Observed Performance Ranges (Bronze Tier)	115
12.8.1	B-PRS (Privacy Risk Score)	115
12.8.2	B-FI (Fidelity Index)	115
12.8.3	B-FV (Fairness Variance)	115
12.8.4	Conformance Distribution	115

12.8.5	Interpretive Guidance	116
12.9	C.9 Validated Synthesis Method Benchmarks	116
12.9.1	For Demographic/Census Data	116
12.9.2	For Commercial/Transactional Data	116
12.9.3	For AI Training/Code Data	117
12.9.4	Cross-Method Insights	117
13	Appendix D: Regulatory Mapping Tables	118
13.1	D.1 SDCF → GDPR Article Mapping	118
13.1.1	Table D.1.1: GDPR Principles (Article 5)	118
13.1.2	Table D.1.2: GDPR Lawful Basis (Article 6)	118
13.1.3	Table D.1.3: GDPR Special Categories (Article 9)	119
13.1.4	Table D.1.4: GDPR Security and Accountability	119
13.1.5	Table D.1.5: Anonymisation vs. Pseudonymization	120
13.2	D.2 SDCF → EU AI Act Mapping	120
13.2.1	Table D.2.1: Article 10 (Training Data)	120
13.2.2	Table D.2.2: High-Risk AI System Requirements	121
13.2.3	Table D.2.3: Conformity Assessment Evidence	121
13.3	D.3 SDCF → ISO/IEC Standards Mapping	121
13.3.1	Table D.3.1: ISO/IEC 27001:2022 (Information Security)	121
13.3.2	Table D.3.2: ISO/IEC 27701:2019 (Privacy)	122
13.3.3	Table D.3.3: ISO/IEC 23894:2023 (AI Risk Management)	123
13.3.4	Table D.3.4: ISO/IEC 42001:2023 (AI Management System)	124
13.4	D.4 SDCF → NIST Frameworks Mapping	124
13.4.1	Table D.4.1: NIST Privacy Framework	124
13.4.2	Table D.4.2: NIST AI Risk Management Framework	125
13.5	D.5 Sector-Specific Regulatory Mapping	126
13.5.1	Table D.5.1: Healthcare (GDPR Article 9, MDR/IVDR)	126
13.5.2	Table D.5.2: Financial Services (Basel III, MiFID II)	126
13.5.3	Table D.5.3: Insurance (Solvency II, IDD)	127
13.6	D.6 Cross-Reference: SDCF Controls to All Regulations	127
13.6.1	Table D.6.1: C1 Purpose Sheet	127
13.6.2	Table D.6.2: C2 Governance Record	127
13.6.3	Table D.6.3: C3 Privacy Risk Testing	128
13.6.4	Table D.6.4: C4 Fidelity Testing	128
13.6.5	Table D.6.5: C5 Fairness Assessment	128
13.6.6	Table D.6.6: C6 Transparency Pack	128
13.6.7	Table D.6.7: C7 Release Rules	129
13.7	D.7 Using Mapping Tables for Compliance	129
13.7.1	How to Use These Mappings	129
13.7.2	Example: Demonstrating EU AI Act Article 10 Compliance	129
13.7.3	Gaps and Additional Requirements	130
14	Appendix E: Sample Outputs	130
14.1	E.1 Gold Tier SDCF-A Certificate (Abbreviated)	130
14.2	E.2 Assessment Report (Executive Summary)	131
14.3	E.3 Risk Statement (Gold Tier)	132
14.4	E.4 JSON Metadata Schema	133

15 Appendix F: Reference Implementations	134
15.1 F.1 Bronze Tier Assessment Using SDMetrics	134
15.1.1 F.1.1 Setup and Data Loading	134
15.1.2 F.1.2 B-PRS Component 1: Membership Inference Proxy (Outlier Analysis)	135
15.1.3 F.1.3 B-PRS Component 2: Record Similarity Proxy	135
15.1.4 F.1.4 B-PRS Component 3: Attribute Disclosure Proxy	136
15.1.5 F.1.5 Composite B-PRS	137
15.1.6 F.1.6 B-FI: Fidelity Assessment (Domain Validation)	138
15.1.7 F.1.7 Complete Bronze Assessment Function	140
15.2 F.2 Gold Tier Assessment Using SDMetrics	141
15.2.1 F.2.1 Complete Gold Tier Implementation	141
15.3 F.3 Tool Integration Patterns	143
15.3.1 F.3.1 mostlyai-qa Integration	143
15.3.2 F.3.2 Generating SDCF Certificate	144
16 Supporting Materials	145
16.1 Glossary	145
16.2 References	146
16.2.1 Regulatory Documents	146
16.2.2 Standards	146
16.2.3 Technical Literature	147
16.2.4 AI Training on Synthetic Data	147
16.2.5 Privacy and Security	147
16.2.6 Fairness and Bias	147
16.2.7 Tools and Software	148
16.3 Acknowledgments	148
16.3.1 Development and Contributions	148
16.3.2 Intellectual Foundations	148
16.3.3 Community Feedback	148
16.3.4 Contact and Contributions	148
16.4 Version History	148
16.4.1 Version 1.95 (December 2025) - Reference Corrections	148
16.4.2 Version 1.9 (December 2025) - HAL Optimised	149
16.4.3 Version 1.0 (November 2025) - Initial Public Release	151
17 Document Summary	151
18 Appendix G: Bronze Tier Validation Detailed Results	153
18.1 G.1 Complete Results Table	153
18.2 G.2 Statistical Summary	153
18.2.1 Descriptive Statistics by Metric	153
18.2.2 Distribution by Conformance Level	153
18.2.3 Correlation Analysis	154
18.3 G.3 Dataset-Specific Notes	154
18.3.1 D1: PLEIAs SYNTH	154
18.3.2 D2–D4: SDV Adult Variants	155
18.3.3 D5: Gretel Safety Alignment	155
18.3.4 D6: MostlyAI Census	156

18.3.5	D7: MostlyAI CDNOW Purchases	156
18.3.6	D8: CMS DE-SynPUF Demo	157
18.3.7	D9: US Census SynLBD Demo	157
18.3.8	D10: Jupyter Agent Dataset	157
18.4	G.4 Reproducibility Information	158
18.4.1	Code Availability	158
18.4.2	Data Access	158
18.4.3	Computational Requirements	158
18.4.4	Expected Results	159
18.5	G.5 Framework Alignment Evidence	159

Synthetic Data Compliance Framework (SDCF)

0.1 Version 1.95

A Purpose-Bounded Methodology for Assessing Privacy, Fidelity, and Fairness in Synthetic Data

Author: Wayne Kearns, Kaionix Labs

Contact: wayne.kearns@nortesconsulting.com

Version: 1.95

Date: December 2025

Licence: Creative Commons Attribution-ShareAlike 4.0 International (CC BY-SA 4.0)

0.2 Copyright and Licensing

Copyright © 2025 Wayne Kearns, Kaionix Labs. All rights reserved.

This work is licensed under a Creative Commons Attribution-ShareAlike 4.0 International Licence.

You are free to: - **Share** — copy and redistribute the material in any medium or format - **Adapt** — remix, transform, and build upon the material - **Use commercially** — organisations may use this framework for internal compliance, governance, and data protection activities

Under the following terms: - **Attribution** — You must give appropriate credit, provide a link to the licence, and indicate if changes were made - **ShareAlike** — If you remix, transform, or build upon the material, you must distribute your contributions under the same licence

This framework is freely available for use by organisations, researchers, and practitioners worldwide, including commercial entities using it for internal compliance and governance. Derivatives must be shared under the same licence.

Citation:

Kearns, W. (2025). Synthetic Data Compliance Framework (SDCF) Version 1.95.

Kaionix Labs. <https://www.kaionix.com/kaionix-labs>

0.3 Document Status

Version: 1.95 (Preprint - HAL Optimised)

Status: Request for Comments (RFC)

Audience: Data Protection Officers, AI Governance Teams, Data Scientists, Legal Counsel, Regulators

Note on Provisional Status: This framework represents best-practice guidance based on current regulatory interpretation and technical capabilities as of December 2025. Thresholds and metrics are provisional pending empirical validation across multiple domains and use cases. Organisations implementing SDCF should calibrate thresholds to their specific risk tolerance and regulatory environment.

Users are encouraged to provide feedback, share implementation experiences, and contribute to the ongoing development of this framework.

0.4 Table of Contents

0.4.1 Core Framework

1. Introduction

- 1.1 What is SDCF?
- 1.2 Why SDCF is Needed
- 1.3 Scope and Limitations
- 1.4 Target Audience
- 1.5 How to Use This Document
- 1.6 Document Structure

2. Regulatory Context

- 2.1 GDPR and Synthetic Data
- 2.2 EU AI Act Article 10
- 2.3 Sector-Specific Standards
- 2.4 International Alignment
- 2.5 Use Cases Covered

3. SDCF Architecture

- 3.1 Purpose-Bounded Assessment Philosophy
- 3.2 The Three Pillars: Privacy, Fidelity, Fairness
- 3.3 Assessment Tiers: Gold, Silver, Bronze
- 3.4 Conformance Levels: SDCF-A, SDCF-P, SDCF-R
- 3.5 Integration Points

4. Control Sets (C1-C7)

- 4.1 C1: Purpose Sheet
- 4.2 C2: Governance Record
- 4.3 C3: Privacy Risk Testing
- 4.4 C4: Fidelity Testing
- 4.5 C5: Fairness Assessment
- 4.6 C6: Transparency Pack
- 4.7 C7: Release Rules

5. Assessment Process

- 5.1 Five-Step Workflow
- 5.2 Tier Selection
- 5.3 Evidence Collection
- 5.4 Scoring and Interpretation
- 5.5 Certificate Issuance
- 6. **Relationship to Existing Tools and Standards**
 - 6.1 Complementary Open-Source Tools
 - 6.2 Vendor Platform Integration
 - 6.3 Standards Alignment (ISO, NIST)
 - 6.4 What SDCF Adds

0.4.2 Technical Appendices

- **Appendix A: Mathematical Definitions**
 - A.1 Privacy Risk Score (PRS)
 - A.2 Fidelity Index (FI)
 - A.3 Fairness Variance (FV)
 - A.4 Normalization and Weighting
 - A.5 Confidence Intervals
 - A.6 Provisional Thresholds
- **Appendix B: Legal and Regulatory Disclaimers**
 - B.1 Not Legal Advice
 - B.2 GDPR Interpretation Guidance
 - B.3 EU AI Act Mapping
 - B.4 Scope Limitations
 - B.5 Liability Framework
 - B.6 Indemnification
- **Appendix C: Bronze Tier Guidance**
 - C.1 Synthetic-Only Assessment Methodology
 - C.2 B-PRS: Privacy Risk Without Source Data
 - C.3 B-FI: Fidelity Without Source Data
 - C.4 B-FV: Fairness Assessment
 - C.5 Certificate Templates
 - C.6 Risk Statements
 - C.7 Case Studies
- **Appendix D: Regulatory Mapping Tables**
 - D.1 SDCF → GDPR Article Mapping
 - D.2 SDCF → EU AI Act Article 10 Mapping
 - D.3 SDCF → ISO 27001/27701 Mapping
 - D.4 SDCF → Sector Standards
- **Appendix E: Sample Outputs**
 - E.1 Assessment Report Example
 - E.2 Certificate Examples (Gold/Silver/Bronze)
 - E.3 Risk Statement Examples
 - E.4 Transparency Pack Example
- **Appendix F: Reference Implementations**
 - F.1 Bronze Tier Using SDMetrics
 - F.2 Silver Tier Using mostlyai-qa

- F.3 Gold Tier Full Assessment
- F.4 Tool Integration Patterns

0.4.3 Supporting Materials

- **Glossary**
 - **References**
 - **Acknowledgments**
 - **Version History**
-

1 Introduction

1.1 What is SDCF?

The Synthetic Data Compliance Framework (SDCF) is a practical, purpose-bounded methodology for assessing whether synthetic data is fit for a specific intended use from privacy, fidelity, and fairness perspectives.

SDCF is designed to answer a fundamental question that organisations face when considering synthetic data:

“Can we confidently use this synthetic dataset for our intended purpose while meeting our regulatory obligations and managing technical risks?”

This framework provides:

- **A structured assessment process** for evaluating synthetic data against regulatory requirements (GDPR, EU AI Act, sector standards)
- **Three complementary measurement pillars:** Privacy risk, statistical fidelity, and algorithmic fairness
- **Tiered methodologies** that work with full source data access (Gold), partial access (Silver), or synthetic-only scenarios (Bronze)
- **Evidence packs** that support regulatory compliance demonstrations, audit trails, and stakeholder communication
- **Honest guidance** about limitations, provisional thresholds, and areas requiring legal interpretation

SDCF is **not** a generic data quality scoring system. It is explicitly purpose-bounded: the same synthetic dataset may be suitable for one use case but unsuitable for another. SDCF requires organisations to define their intended purpose upfront and assesses fitness for that specific purpose.

1.2 Why SDCF is Needed

1.2.1 The Synthetic Data Promise and Challenge

Synthetic data, generated through techniques such as generative adversarial networks [17], variational autoencoders [21], and specialised tabular methods [27, 30], offers compelling benefits:

- **Privacy protection:** Enables data sharing and analysis without exposing real individuals
- **Data augmentation:** Addresses data scarcity in regulated domains (healthcare [7], finance)

- **Bias mitigation:** Opportunity to create more balanced, representative datasets
- **AI training efficiency:** Recent advances (e.g., Pleias Baguettotron, 2025) show high-quality synthetic training data can produce competitive models with dramatically lower compute requirements

However, these benefits materialise **only if the synthetic data is actually fit for purpose**. Poor-quality synthetic data can:

- Create false sense of privacy protection (re-identification risks remain [6, 28])
- Produce misleading analytical insights (distribution distortion)
- Amplify rather than mitigate algorithmic bias [5]
- Cause model collapse in AI training scenarios
- Generate regulatory liability (GDPR Article 32 security failures, EU AI Act Article 10 data governance violations [14, 15])

1.2.2 The Standards Gap

Organisations seeking to validate synthetic data currently face a fragmented landscape:

Available tools provide metrics but not interpretation:

- Open-source libraries [9, 27] compute privacy and fidelity metrics
- Vendor platforms (Gretel, MOSTLY AI, Syntho) generate quality reports
- Neither translates metrics into regulatory compliance evidence

Official standards remain principle-based:

- GDPR [14] establishes requirements for anonymisation (Recital 26) but not validation methods
- EU AI Act Article 10 [15] mandates data governance but not specific test procedures
- ISO/IEC standards [19, 20] define controls but not synthetic-data-specific implementation
- NIST frameworks [25] provide conceptual guidance but limited operational detail

Practitioners need actionable guidance that:

- Bridges the gap between metrics and compliance
- Works with real-world constraints (often no source data access)
- Provides defensible evidence for regulators and auditors
- Integrates privacy, fidelity, and fairness assessment
- Acknowledges limitations honestly rather than overpromising

SDCF fills this gap as an **interpretation and compliance mapping layer** that sits above metric computation tools and below principle-based regulations.

1.2.3 Empirical Validation

This framework presents preliminary empirical validation through a Bronze Tier retrospective study assessing 10 diverse synthetic datasets across 7 domains (demographic, healthcare, e-commerce, AI training, AI safety, business, code/data) using 5 synthesis methods (GaussianCopula, CTGAN, TVAE, commercial GANs, LLM-generated). Initial evidence suggests the methodology performs as designed, with conservative risk classification (60% Restricted conformance), quality discrimination capability (8.8x B-PRS range: 0.090 to 0.789), and cross-domain applicability.

The study identified that TVAE (Tabular Variational Autoencoder) synthesis [30] demonstrates superior privacy-quality balance for demographic data ($n=2$ datasets), achieving 19–20% lower privacy risk scores compared to CTGAN [30] or GaussianCopula methods while maintaining equivalent fidelity (> 0.99). Statistical expansion to $n>50$ datasets with adversarial validation (membership inference benchmarking, linkage attacks) is planned for v2.0. Complete validation results, including dataset-by-dataset analysis, framework alignment assessment, and practitioner guidance, are presented in Section 7, with detailed statistical analyses in Appendix G.

1.3 Research Contributions

This work makes four primary contributions to the field of synthetic data validation and governance:

1. Unified Compliance Architecture: SDCF provides the first comprehensive framework integrating privacy risk assessment, statistical fidelity evaluation, and algorithmic fairness testing with explicit regulatory mapping to GDPR Articles 25 & 32, EU AI Act Article 10, and ISO/IEC standards. Unlike existing approaches that address individual dimensions in isolation, SDCF operationalises the trade-offs between privacy, utility, and fairness through purpose-bounded assessment methodology.

2. Tiered Validation Methodology: The three-tier system (Gold/Silver/Bronze) addresses a critical real-world constraint: source data is often unavailable for validation due to privacy regulations, commercial restrictions, or legacy system limitations. The Bronze Tier methodology enables evidence-based assessment of third-party synthetic datasets without source access, filling a significant gap in existing validation frameworks.

3. Mathematical Formalisation and Provisional Thresholds: SDCF provides rigorous mathematical definitions for Privacy Risk Score (PRS), Fidelity Index (FI), and Fairness Variance (FV), along with provisional thresholds derived from conservative risk principles. These metrics connect technical measurements to compliance requirements, enabling auditable, repeatable assessments.

4. Empirical Validation and Reference Implementation: Version 1.8 presents preliminary empirical validation through retrospective assessment of 10 diverse synthetic datasets spanning 7 domains (demographic, healthcare, e-commerce, AI training, AI safety, business, code/data) and 5 synthesis methods. This initial evidence supports framework design principles; statistical expansion to $n>50$ with adversarial validation is planned for v2.0. The validation demonstrates framework effectiveness, cross-domain applicability, and conservative risk classification. A complete Python reference implementation with reproducibility materials enables practitioners to adopt and adapt the methodology.

These contributions collectively enable organisations to establish evidence-based synthetic data validation processes with clear audit trails for regulatory compliance review.

1.4 Scope and Limitations

1.4.1 What SDCF Covers

In Scope: - Structured tabular synthetic data (rows and columns) - Synthetic data derived from personal data or used as substitute for personal data - Use cases including: analytics, reporting, model training, software testing, data sharing - Privacy risk assessment (re-identification, attribute disclosure) - Statistical fidelity assessment (distribution similarity, utility preservation) - Fairness assessment (representation, predictive parity) - GDPR Article 6 (lawful basis) and Article 9 (special

categories) compliance support - EU AI Act Article 10 (data governance) compliance support - Assessment with full source access, partial access, or no source access

Out of Scope: - Unstructured data (images, video, audio, text) - different risk profiles - Fully synthetic data with no relationship to real individuals (not privacy-sensitive) - Federated or distributed synthetic data generation - Adversarial robustness testing - Synthetic data generation methodology assessment (SDCF assesses outputs, not generation processes) - Legal determination of “anonymisation” vs. “pseudonymization” (SDCF provides technical evidence; legal classification requires counsel)

1.4.2 What SDCF Does NOT Do

SDCF is not a substitute for: - **Legal advice:** Organisations must obtain independent legal opinion on GDPR/regulatory compliance - **Data Protection Impact Assessment (DPIA):** SDCF can inform DPIA but does not replace the requirement - **Security assessment:** SDCF assumes appropriate technical/organisational security measures are in place - **Domain expertise:** Interpreting fidelity and fairness requires subject matter knowledge - **Official certification:** SDCF certificates reflect technical assessment, not regulatory endorsement

SDCF does not guarantee: - That synthetic data is legally “anonymous” under GDPR (context-dependent, requires legal interpretation) - That use of synthetic data is appropriate for all purposes - That risks identified will not materialise (probabilistic assessment, not elimination) - That thresholds are universally applicable (provisional pending domain-specific calibration)

1.5 Target Audience

SDCF is designed for practitioners who need to make evidence-based decisions about synthetic data:

Primary Audience: - **Data Protection Officers (DPOs):** Need to assess GDPR compliance of synthetic data initiatives - **AI Governance Teams:** Responsible for EU AI Act Article 10 compliance, model risk management - **Data Scientists/ML Engineers:** Building and validating synthetic data pipelines, training AI models on synthetic data - **Legal Counsel:** Advising on data protection compliance, requiring technical evidence - **Internal Audit/Compliance:** Verifying data governance controls

Secondary Audience: - **Regulators:** Reference methodology for assessing organisational practices - **External Auditors:** Framework for independent validation - **Procurement Teams:** Evaluation criteria for synthetic data vendors - **Researchers:** Standardised methodology for reproducible assessments

Knowledge Prerequisites: - Basic understanding of GDPR principles (lawful basis, anonymisation, accountability) - Familiarity with statistical concepts (distributions, correlation, variance) - Awareness of algorithmic fairness concepts (protected attributes, disparate impact) - Ability to work with data analysis tools (Python/R or willingness to engage technical colleagues)

No prerequisites in: - Advanced mathematics (formulas explained with intuition) - Machine learning model development - Cryptography or privacy-enhancing technologies

1.6 How to Use This Document

1.6.1 Reading Paths by Role

If you are a DPO or Legal Counsel: 1. Read Section 1 (Introduction) and Section 2 (Regulatory Context) fully 2. Skim Section 3 (Architecture) to understand three pillars and tiers 3. Read Section 4 (Control Sets) focusing on C1, C2, C6, C7 4. Study Appendix B (Legal Disclaimers) carefully 5. Reference Appendix D (Regulatory Mapping) as needed

If you are a Data Scientist or ML Engineer: 1. Read Section 1 (Introduction) and Section 3 (Architecture) 2. Study Section 4 (Control Sets) focusing on C3, C4, C5 3. Work through Section 5 (Assessment Process) step-by-step 4. Dive deep into Appendix A (Mathematical Definitions) 5. Implement using Appendix F (Reference Implementations)

If you are an AI Governance Lead or Compliance Manager: 1. Read entire document sequentially (Sections 1-6) 2. Pay special attention to Section 3.4 (Conformance Levels) 3. Review Appendix C if you anticipate Bronze Tier scenarios 4. Use Appendix E (Sample Outputs) to understand deliverables

If you are evaluating SDCF for adoption: 1. Read Section 1 (Introduction) and Section 6 (Relationship to Existing Tools) 2. Review Quick Start Guide (separate document) 3. Examine sample outputs in Appendix E 4. Try reference implementation in Appendix F on test data 5. Assess organisational readiness and tooling requirements

1.6.2 Document Conventions

Terminology: - **MUST/REQUIRED:** Mandatory requirement for SDCF compliance - **SHOULD/RECOMMENDED:** Strongly advised but may be omitted with justification - **MAY/OPTIONAL:** Discretionary choice based on context - **Source data:** Original real-world data from which synthetic data is derived - **Synthetic data:** Artificially generated data intended to preserve statistical properties of source data - **Purpose:** The specific intended use case for the synthetic data

Typography: - **Code blocks** indicate technical commands or data structures - *Italics* emphasize key terms on first introduction - **Bold** highlights critical warnings or requirements

Examples: Throughout this document, examples use fictional scenarios to illustrate concepts: - IrishHealth Bank: Financial services synthetic transaction data - MedTech Research: Healthcare synthetic patient data - InsureCo: Insurance synthetic claims data

These examples are illustrative only and do not represent actual assessments.

1.7 Document Structure

This framework is organised in three parts:

Part I: Core Framework (Sections 1-6) Conceptual foundation, process overview, and practical guidance for all practitioners. Read sequentially for comprehensive understanding.

Part II: Technical Appendices (Appendices A-F) Mathematical rigor, legal considerations, and implementation details. Reference as needed based on role and use case.

Part III: Supporting Materials Glossary, references, and version history for ongoing framework evolution.

Companion Documents: - **SDCF Quick Start Guide:** 4-page practical introduction - **SDCF Templates Package:** Fillable Word/Excel templates for all control sets - **SDCF Reference Implementation:** Python scripts demonstrating Bronze/Silver/Gold tier assessments

All materials available at: <https://www.kaionix.com/kaionix-labs>

End of Section 1

Continue to Section 2: Regulatory Context to understand the compliance landscape that SDCF addresses.

2 Related Work

The SDCF framework intersects four major areas of research: synthetic data generation, privacy and security auditing, fairness and responsible AI, and AI governance. This section summarises the most relevant work and positions SDCF in relation to existing approaches.

2.1 Synthetic Data Generation and Evaluation

Synthetic data has been studied extensively through generative adversarial networks (GANs) [17] and variational autoencoders [21], with more recent work focusing on tabular and time-series data via models such as CTGAN [30], TVAE, and TimeGAN [33]. Practical tooling is provided by frameworks such as the Synthetic Data Vault (SDV) [27] and SDMetrics [9], which supply statistical similarity and utility measures. Domain-specific applications have demonstrated the potential of synthetic data in healthcare [7, 8] and other sensitive domains. These works demonstrate strong progress in generation and evaluation, but they generally lack a compliance or governance orientation.

2.2 Privacy Risks and Attacks

Privacy leakage from synthetic data has been measured using membership inference attacks [28, 32] and reconstruction attacks [6, 24], as well as differential privacy techniques that provide formal guarantees [1, 10, 11]. Practical implementations such as PrivBayes [34] demonstrate generation under privacy constraints. Comprehensive security analyses [26] contextualise these threats within broader ML security concerns. These contributions offer critical insights into privacy risks for generative models, yet most operate as isolated tests rather than components of an integrated assessment pipeline. SDCF incorporates such risks into a unified scoring model (Privacy Risk Score), where privacy tests are contextualised within a broader risk taxonomy and compliance framework.

2.3 Fairness, Representation, and Responsible AI

Fairness in machine learning is a large and multidisciplinary field [4], with work examining representational harms [5], calibration trade-offs [22], equality of opportunity [18], and foundational concepts of fairness through awareness [12]. System documentation practices such as model cards [23] and datasheets [16] provide transparency mechanisms, while explainable AI research [2] addresses interpretability requirements. Although synthetic data is often assumed to mitigate bias, empirical studies show that generative models can amplify, obscure, or introduce new forms of

unfairness. SDCF explicitly addresses this by defining fairness as a control domain (C5) and integrating representation tests, utility checks, and distributional comparisons into a structured risk assessment.

2.4 Governance, Regulation, and Standards

Governance frameworks such as the NIST AI Risk Management Framework [25], the ICO guidance on AI and data protection [29], and the GDPR [14] define principles for accountability, transparency, and data minimisation. The Article 29 Working Party (now EDPB) Opinion 05/2014 [3] established anonymisation assessment criteria that remain influential. More recently, the EU AI Act [15] and ISO/IEC 42001 [20] establish system-level requirements for risk management and AI lifecycle governance, building on information security standards such as ISO/IEC 27001 [19]. However, none of these offer detailed operational guidance specific to synthetic data pipelines. SDCF can be understood as a technical and procedural instantiation of these principles tailored to synthetic data.

2.5 Positioning of SDCF

Prior work addresses synthetic data quality, privacy, fairness, and governance, but typically in isolation. SDV and SDMetrics offer statistical tests; privacy researchers provide attack methodologies; fairness literature focuses on bias metrics; and regulatory frameworks describe high-level expectations. Practical guidance exists for specific domains [13, 31], but lacks integration across compliance dimensions. SDCF is the first unified, purpose-bounded framework that integrates all of these components into a coherent, tiered compliance methodology supported by empirical validation across multiple domains. Where existing work provides metrics or principles, SDCF provides an assessment architecture, control taxonomy, conformance levels, and evidence-based validation.

End of Section 1.5

Continue to Section 2: Regulatory Context for detailed GDPR, EU AI Act, and sector-specific regulatory requirements.

3 Regulatory Context

3.1 GDPR and Synthetic Data

3.1.1 The Anonymisation Question

The General Data Protection Regulation (GDPR) creates a fundamental distinction that drives organisational interest in synthetic data:

- **Personal data** (Article 4(1)): Subject to full GDPR obligations including lawful basis, purpose limitation, data subject rights, security requirements, and accountability
- **Anonymous data** (Recital 26): “Not subject to data protection legislation” if it “does not relate to an identified or identifiable natural person”

Synthetic data occupies ambiguous territory. It is *derived from* personal data but *does not directly correspond* to real individuals. The critical question: Is it “anonymous” under GDPR?

3.1.2 GDPR Does Not Provide a Simple Answer

Recital 26 establishes the test: > “To determine whether a natural person is identifiable, account should be taken of all the means reasonably likely to be used [...] to identify the natural person directly or indirectly. To ascertain whether means are reasonably likely to be used [...] account should be taken of all objective factors, such as the costs of and the amount of time required for identification, taking into consideration the available technology at the time of the processing and technological developments.”

This creates a risk-based, context-dependent assessment: - What constitutes “reasonable means” varies by adversary capability, data sensitivity, and available auxiliary information - Technology continuously evolves (what was safe in 2020 may not be in 2025) - The data controller’s intent does NOT determine status (cannot simply declare data “anonymous”) - Must consider re-identification risk over the data lifecycle, not just at generation

Article 29 Working Party Opinion 05/2014 on Anonymisation Techniques [3] (predecessor to EDPB) established that anonymisation must be: - **Irreversible:** Cannot recover original data - **Resistant to singling out:** Cannot isolate records for specific individuals - **Resistant to linkability:** Cannot link records across datasets - **Resistant to inference:** Cannot deduce information about individuals

Synthetic data generation techniques do NOT automatically satisfy these criteria. The quality of anonymisation depends on generation methodology, source data characteristics, and intended use.

3.1.3 EDPB Guidelines 01/2025 on Pseudonymisation

In January 2025, the European Data Protection Board published Guidelines 01/2025 clarifying the boundary between anonymisation and pseudonymization. Key takeaways relevant to synthetic data:

Pseudonymisation (Article 4(5)) vs. Anonymisation: - Pseudonymisation reduces linkability but data remains personal (requires safeguards, enables some processing flexibility) - Anonymisation removes personal data character completely (no GDPR obligations, but must be genuine) - The distinction is technical and contextual, not based on labels or intentions

Risk-based assessment required: - Must consider “all means reasonably likely” for re-identification - Includes technical attacks (linkage attacks, membership inference) and practical scenarios (insider threats, data breaches) - Documentation burden on controllers to demonstrate anonymisation robustness

For synthetic data specifically: - Generation methodology matters (simple sampling may retain identifiable records; advanced methods reduce but don’t eliminate risk) - Statistical properties can enable inference attacks even without direct correspondence - Outliers and rare combinations present higher risk - Must assess risk in context of intended use and disclosure scenario

SDCF Response: Control Set C3 (Privacy Risk Testing) operationalises this risk-based assessment by quantifying membership inference risk, record similarity, and attribute disclosure probability. However, SDCF provides *technical evidence* only; legal classification as “anonymous” or “pseudonymous” requires independent legal counsel (see Appendix B).

3.1.4 GDPR Articles Implicated in Synthetic Data Use

Even if synthetic data is treated as personal data (conservative approach), certain processing may be permissible:

Article 6 (Lawful Basis): - 6(1)(f) Legitimate Interests: Processing necessary for legitimate interests (e.g., research, analytics) balanced against data subject rights - *Synthetic data enables legitimate interest claims by reducing privacy impact* - **6(1)(e) Public Interest:** Processing necessary for public interest or official authority tasks - *Relevant for healthcare research, public health, regulatory reporting*

Article 9 (Special Categories of Personal Data): - Synthetic health data, biometric data, genetic data may fall under Article 9 even if risk-reduced - Requires explicit consent (9(2)(a)) or specific derogations (research 9(2)(j), public health 9(2)(i)) - Higher bar for “appropriate safeguards” and necessity demonstration

Article 25 (Data Protection by Design and by Default): - Synthetic data generation can be privacy-enhancing measure demonstrating design principle - Must implement “appropriate technical and organisational measures” (synthetic data generation methodology assessment)

Article 32 (Security of Processing): - Synthetic data reduces but does not eliminate breach risk - Must maintain security appropriate to residual risk (encryption, access control, monitoring)

Article 35 (Data Protection Impact Assessment): - High-risk processing may require DPIA even when using synthetic data - SDCF assessment can inform DPIA risk analysis but does not replace it

3.1.5 Practical Implications for Organisations

Conservative Approach (Recommended): Treat synthetic data as pseudonymous personal data unless demonstrated to meet anonymisation standard through rigorous assessment. This means: - Maintain lawful basis for processing - Respect purpose limitation (only use for declared purposes) - Implement appropriate security measures - Provide transparency to data subjects about synthetic data use - Retain accountability documentation

Risk-Based Approach: Invest in robust assessment (SDCF Bronze/Silver/Gold tier depending on source data availability) to build evidence case for anonymisation. If evidence is strong: - Engage legal counsel for formal opinion on GDPR status - Document assessment methodology and results - Monitor for technological changes affecting risk (annual reassessment) - Maintain conservative posture for high-risk scenarios (Article 9 data, vulnerable populations)

SDCF Role: SDCF provides the technical assessment foundation for either approach. The Privacy Risk Score (PRS) quantifies re-identification risk using state-of-the-art metrics. Organisations can use PRS to: - Demonstrate due diligence in anonymisation attempts - Support legitimate interest balancing tests - Inform DPIA risk ratings - Provide auditable evidence trail

However, SDCF explicitly does NOT make legal determination of GDPR status. That determination requires legal expertise considering organisational context, data sensitivity, disclosure scenario, and risk tolerance.

3.2 EU AI Act Article 10

The EU Artificial Intelligence Act (AI Act), which entered into force in August 2024 with phased implementation through 2027, establishes data governance requirements for high-risk AI systems.

3.2.1 Article 10: Data and Data Governance

Article 10(2) establishes that training, validation, and testing datasets must be: - **Relevant**, sufficiently representative, and to the best extent possible, free of errors - **Appropriate** in respect of the intended purpose of the high-risk AI system

Article 10(3) requires that datasets: - **Take into account the geographical, contextual, behavioral, or functional setting** within which the AI system is intended to be used - **Are subject to data governance and management practices appropriate for the intended purpose**

Article 10(4) specifically addresses bias: - **Datasets must be relevant, representative, free of errors, and complete** - Must examine for **possible biases** that may affect health, safety, or fundamental rights - Must implement **appropriate measures to detect, prevent, and mitigate possible biases**

3.2.2 Synthetic Data as Article 10 Compliance Strategy

Organisations are increasingly using synthetic data to meet Article 10 requirements:

Advantages: - **Bias mitigation:** Can deliberately balance synthetic datasets to address underrepresentation - **Completeness:** Can generate data for rare scenarios not well-represented in historical data - **Error correction:** Opportunity to clean known data quality issues during generation - **Privacy protection:** Reduces risk of training on sensitive personal data

Challenges: - **Representativeness question:** Is synthetic data “sufficiently representative” of real-world deployment context? - **Bias amplification risk:** Poor generation methodology may amplify rather than mitigate source data biases - **Fidelity-fairness tension:** Oversampling minorities improves fairness but may reduce fidelity to real-world distributions - **Validation burden:** How to demonstrate synthetic data meets Article 10 requirements?

3.2.3 What Article 10 Does NOT Specify

The AI Act establishes principles but not technical methods: - No prescribed metrics for “sufficiently representative” - No thresholds for “free of errors” - No specific tests for bias detection - No validation procedures for synthetic data quality

Harmonised standards expected: The European Commission is tasked with developing harmonised standards (Article 40) to provide presumption of conformity. As of December 2025, these standards are in development but not yet published.

3.2.4 SDCE as Article 10 Implementation Framework

SDCE operationalises Article 10 requirements for synthetic training data:

Article 10(2) “Relevant, representative, appropriate”: - **C1 Purpose Sheet:** Documents intended purpose and defines “appropriate” criteria - **C4 Fidelity Testing:** Assesses statistical

representativeness through Fidelity Index (FI) - **C3 Privacy Risk Testing**: Ensures data doesn't compromise safety through re-identification

Article 10(3) "Data governance practices": - **C2 Governance Record**: Documents decision-making, roles, responsibilities - **C7 Release Rules**: Defines access controls and usage constraints - **C6 Transparency Pack**: Provides documentation for audit and accountability

Article 10(4) "Possible biases": - **C5 Fairness Assessment**: Quantifies representation and predictive parity across protected attributes - **Fairness Variance (FV)**: Provides measurable bias metric - **C6 Transparency Pack**: Documents known limitations and mitigation measures

Article 10 Use Case (AI Training Data): In 2025, advances like Pleias' Baguettotron demonstrated that high-quality synthetic training data can produce competitive models with reduced compute requirements. Organisations training AI models (especially in regulated domains: healthcare diagnostics, financial fraud detection, insurance underwriting) can use SDCF to: - Validate synthetic training dataset quality before expensive training runs - Demonstrate Article 10 compliance to regulators and customers - Build auditable evidence of bias mitigation efforts - Assess fitness-for-purpose for specific AI system deployment contexts

SDCF Bronze Tier is particularly relevant for AI training scenarios: Organisations often don't have access to source data (licensed third-party synthetic datasets, synthetically augmented open-source data). Bronze Tier provides synthetic-only assessment methodology addressing this real-world constraint.

3.3 Sector-Specific Standards

Beyond GDPR and AI Act, organisations face sector-specific requirements that interact with synthetic data use:

3.3.1 Healthcare

Medical Device Regulation (MDR) / In-Vitro Diagnostic Regulation (IVDR): - Clinical validation requirements for AI/ML medical devices - Training data quality directly affects device safety and performance claims - Notified Body review requires demonstrating data representativeness

Health Data Regulations (varies by Member State): - Ireland: Section 44 Data Protection Act 2018 (health research exemptions require safeguards) - Additional consent/governance requirements for health data beyond GDPR Article 9

SDCF Application: - Fidelity Index ensures synthetic patient data preserves clinical utility - Fairness Assessment addresses health equity concerns (representation across demographics, rare conditions) - Bronze Tier supports scenarios where patient privacy precludes source data sharing for validation

3.3.2 Financial Services

Basel III / Capital Requirements Regulation (CRR): - Model risk management frameworks require validation of model training data - Synthetic data for stress testing must demonstrate representativeness

Markets in Financial Instruments Directive (MiFID II): - Algorithm testing requirements may use synthetic transaction data - Validation must demonstrate realistic market conditions

Anti-Money Laundering (AML) / Know Your Customer (KYC): - Synthetic data for AML model training must preserve rare event characteristics (fraud, money laundering patterns) - Fidelity assessment critical for detection effectiveness

SDCF Application: - Fidelity Index assesses preservation of distributional characteristics (tail events, correlations) - Purpose-bounded approach: synthetic data fit for back-testing may not be fit for regulatory capital calculations

3.3.3 Insurance

Solvency II: - Internal model validation requirements - Synthetic claims data for catastrophe modeling must demonstrate statistical adequacy

Insurance Distribution Directive (IDD): - Algorithmic pricing fairness requirements - Synthetic data for pricing model training must not amplify bias

SDCF Application: - Fairness Assessment quantifies disparate impact across protected classes - Fidelity Index validates preservation of loss distributions and dependencies

3.3.4 Public Sector

Open Data Directive (EU 2019/1024): - Encourages public sector data sharing for reuse - Synthetic data enables sharing while protecting citizen privacy

eIDAS 2.0 / European Digital Identity: - Identity verification systems may use synthetic data for testing - Privacy risk assessment critical for identity data

SDCF Application: - Privacy Risk Score supports safe open data publication decisions - Transparency Pack documents limitations for downstream users

3.4 International Alignment

While SDCF focuses on EU regulatory context, the framework aligns with international standards and can support global compliance:

3.4.1 ISO/IEC Standards

ISO/IEC 27001:2022 (Information Security Management): - Annex A.8: Asset Management (synthetic data as information asset) - SDCF Control Sets C2 and C7 support asset management requirements

ISO/IEC 27701:2019 (Privacy Information Management): - Extension to ISO 27001 for privacy - SDCF Privacy Risk Score supports risk assessment requirements

ISO/IEC 23894:2023 (AI Risk Management): - Data quality and bias management - SDCF Fidelity Index and Fairness Variance operationalise these concepts

ISO/IEC 42001:2023 (AI Management System): - Recently published, covers AI system governance - SDCF assessment process can inform compliance evidence

3.4.2 NIST Frameworks

NIST Privacy Framework (2020): - Core functions: Identify, Govern, Control, Communicate, Protect - SDCF Control Sets map to these functions (C1→Identify, C2→Govern, C3→Control, C6→Communicate, C7→Protect)

NIST AI Risk Management Framework (2023): - Trustworthy AI characteristics including fairness, privacy, transparency - SDCF three-pillar approach (Privacy, Fidelity, Fairness) aligns with NIST trustworthiness dimensions

NIST Differential Privacy Guidelines (ongoing): - Emerging guidance on privacy-preserving synthetic data - SDCF Privacy Risk Score complements differential privacy guarantees (technical evidence when formal guarantees unavailable)

3.4.3 UK / Singapore / Other Jurisdictions

UK ICO Guidance on Anonymisation (2012, updated expectations): - Similar “motivated intruder test” to GDPR “reasonable means” - SDCF assessment methodology applicable under UK GDPR

Singapore Personal Data Protection Act (PDPA): - Anonymisation exemptions similar to GDPR - SDCF supports demonstrating “not reasonably identifiable” standard

SDCF Portability: The framework’s purpose-bounded, risk-based approach is jurisdiction-neutral in design. Organisations operating across jurisdictions can: - Use SDCF for technical assessment foundation - Layer jurisdiction-specific legal interpretation on top - Maintain consistent technical evidence base across markets

3.5 Use Cases Covered

SDCF is designed to assess synthetic data across diverse use cases. Assessment criteria and thresholds vary by purpose (purpose-bounded philosophy).

3.5.1 Use Case 1: Internal Analytics and Reporting

Scenario: Organisation generates synthetic data from customer database for internal business intelligence, analytics dashboards, or management reporting.

Regulatory Drivers: - GDPR Article 6(1)(f) legitimate interest or 6(1)(b) contract performance - Article 32 security of processing (reduce breach impact) - Internal audit and compliance requirements

SDCF Assessment Focus: - **Privacy (High):** Internal disclosure still carries re-identification risk (insider threats, data breaches) - **Fidelity (High):** Business decisions depend on accurate statistical insights - **Fairness (Medium):** Less critical for non-customer-facing analytics, but bias awareness important

Typical Tier: Silver or Gold (source data available for validation)

3.5.2 Use Case 2: External Data Sharing and Collaboration

Scenario: Organisation shares synthetic data with partners, researchers, or open data initiatives to enable collaborative analysis while protecting privacy.

Regulatory Drivers: - GDPR Article 6(1)(f) legitimate interest with heightened scrutiny - Article 35 DPIA may be required (systematic monitoring, large-scale special category data) - Contractual data sharing agreements

SDCF Assessment Focus: - **Privacy (Critical):** External disclosure amplifies re-identification risk (adversaries with auxiliary data) - **Fidelity (High):** Collaborators depend on data quality for valid research conclusions - **Fairness (High):** External use may have societal implications (research publication, policy influence)

Typical Tier: Gold preferred (highest assurance), Silver acceptable with strong controls

3.5.3 Use Case 3: AI/ML Model Training

Scenario: Organisation uses synthetic data to train machine learning models, either exclusively or to augment real data.

Regulatory Drivers: - EU AI Act Article 10 (data governance for high-risk AI systems) - Sector regulations (MDR/IVDR for medical AI, CRR for financial models) - Model risk management frameworks

SDCF Assessment Focus: - **Privacy (Medium to High):** Depends on model deployment context (edge devices vs. controlled environment) - **Fidelity (Critical):** Model performance directly depends on training data representativeness; poor fidelity causes model failures - **Fairness (Critical):** Training data bias propagates to model predictions, may violate non-discrimination requirements

Typical Tier: Bronze often applicable (third-party synthetic training datasets, no source access), Gold when building proprietary synthetic data

Special Considerations: - Recent advances (Baguettotron, 2025) demonstrate synthetic-only training viability - Quality assessment before expensive training runs (pre-flight validation) - Ongoing monitoring as model is retrained on new synthetic batches

3.5.4 Use Case 4: Software Testing and Development

Scenario: Organisation generates synthetic data to test software applications, APIs, or data pipelines without exposing production data to development environments.

Regulatory Drivers: - GDPR Article 32 security (test environments are higher breach risk) - Article 25 data protection by design (testing with realistic but safe data) - DevOps best practices (separation of test and production data)

SDCF Assessment Focus: - **Privacy (Medium):** Test environments may have weaker security, but contained disclosure risk - **Fidelity (Medium):** Data must be realistic enough to trigger edge cases and validate functionality - **Fairness (Low):** Usually not applicable for technical testing (unless testing fairness-sensitive features)

Typical Tier: Bronze (functionality testing doesn't require perfect fidelity, no source access needed)

3.5.5 Use Case 5: Regulatory Reporting and Compliance

Scenario: Organisation uses synthetic data to demonstrate compliance capabilities to regulators, auditors, or certification bodies without exposing real customer data.

Regulatory Drivers: - Sector-specific audit requirements (financial, healthcare, insurance) - ISO certification audits - Regulatory sandbox testing (innovative products using synthetic scenarios)

SDCF Assessment Focus: - **Privacy (High):** Regulatory disclosure still carries risk, must maintain confidentiality - **Fidelity (Critical):** Regulators must trust synthetic data reflects real-world risk profile - **Fairness (High):** Demonstrating non-discriminatory practices requires representative data

Typical Tier: Gold strongly preferred (regulators expect highest assurance level)

3.5.6 Use Case 6: Training and Education

Scenario: Organisation generates synthetic data for employee training, academic courses, or public workshops to teach data analysis skills with realistic but safe datasets.

Regulatory Drivers: - GDPR Article 89 (scientific research and education exemptions, with safeguards) - Institutional ethics review (universities) - Public interest in skills development

SDCF Assessment Focus: - **Privacy (Medium):** Broad disclosure (students, public) requires strong anonymisation - **Fidelity (Medium):** Must be realistic for pedagogical value, but perfect fidelity not required - **Fairness (Medium):** Educational materials should not perpetuate biases

Typical Tier: Bronze or Silver (balance between realism and simplicity for learning)

3.5.7 Purpose-Bounded Assessment in Practice

The same synthetic dataset may receive different SDCF ratings for different purposes:

Example: Synthetic Patient Data - Internal clinical research (Gold Tier, SDCF-A): High privacy/fidelity/fairness requirements, source access available → Suitable - **Medical device training (Gold Tier, SDCF-P):** Critical fidelity, minor fairness limitations → Suitable with restrictions - **Public health education (Bronze Tier, SDCF-R):** Simplified dataset for teaching → Suitable for teaching only - **External research collaboration (Silver Tier, SDCF-R):** Moderate privacy risk, acceptable fidelity → Requires additional controls

C1 Purpose Sheet documents intended use and defines success criteria accordingly. This prevents misuse of datasets assessed for one purpose in another context where risks differ.

End of Section 2

Continue to Section 3: SDCF Architecture to understand the framework's structural design and assessment methodology.

4 SDCF Architecture

4.1 Purpose-Bounded Assessment Philosophy

4.1.1 The Fitness-for-Purpose Principle

SDCF rejects the notion of generic “synthetic data quality scores” divorced from context. The same synthetic dataset may be: - Excellent for software testing but unsuitable for regulatory report-

ing - Appropriate for internal analytics but too risky for external sharing - Sufficient for training educational models but inadequate for high-risk AI systems

Core principle: *Synthetic data assessment must always answer “fit for WHAT purpose?”*

This purpose-bounded approach differs fundamentally from common alternatives:

Approach	Philosophy	Limitation
Generic Quality Score	Compute aggregate metric (0-100) representing “overall quality”	Obscures trade-offs; high score doesn’t guarantee fitness for specific use
Pass/Fail Binary	Set universal threshold (e.g., “privacy score >90% = pass”)	Ignores context; same threshold inappropriate for all scenarios
Vendor Self-Assessment	Generator declares data “high quality” without independent validation	Conflicts of interest; no standardised methodology
SDCF Purpose-Bounded	Define intended purpose first, then assess fitness for THAT purpose with explicit trade-off documentation	Transparent about limitations; supports informed decision-making

4.1.2 How Purpose-Bounded Assessment Works

Step 1: Purpose Definition (C1 Purpose Sheet) Organisation explicitly documents: - What will this synthetic data be used for? - What decisions or processes depend on it? - Who will have access (internal only, external partners, public)? - What regulatory requirements apply? - What failure modes must be avoided (privacy breach, biased outcomes, misleading insights)?

Step 2: Risk-Weighted Assessment SDCF assesses three dimensions (Privacy, Fidelity, Fairness) but weights them based on purpose:

Example: Internal Business Intelligence Dashboard - Privacy: High weighting (internal disclosure still carries risk) - Fidelity: Critical weighting (business decisions depend on accuracy) - Fairness: Medium weighting (bias awareness important but not customer-facing)

Example: Public Open Data Release - Privacy: Critical weighting (public disclosure, maximum adversary capability) - Fidelity: High weighting (researchers depend on data quality) - Fairness: High weighting (public resource should represent all communities)

Example: Software Testing - Privacy: Medium weighting (contained environment, limited access) - Fidelity: Medium weighting (realistic enough for edge cases, perfection not required) - Fairness: Low weighting (technical testing typically not fairness-sensitive)

Step 3: Conformance Level Assignment Based on assessment results and purpose requirements: - **SDCF-A (Approved)**: Suitable for intended purpose without restrictions - **SDCF-P (Provisional)**: Suitable with documented limitations and additional controls - **SDCF-R (Restricted)**: Not suitable for stated purpose OR suitable only for different, more limited purpose

Step 4: Evidence Pack (C6 Transparency Pack) Documentation includes: - Purpose statement - Assessment results (PRS, FI, FV scores) - Conformance level and rationale - Known limitations and recommended mitigations - Restrictions on use beyond stated purpose

Purpose-bounded assessment prevents scope creep: Dataset assessed for training environment cannot be automatically repurposed for production use without reassessment.

4.2 The Three Pillars: Privacy, Fidelity, Fairness

SDCF assesses synthetic data quality through three complementary dimensions. All three matter, but their relative importance varies by purpose.

4.2.1 Pillar 1: Privacy Risk

Definition: The probability that synthetic data enables re-identification of individuals from the source data or discloses sensitive attributes.

Why it matters: - GDPR Article 32 requires “appropriate security” - privacy risk quantifies residual exposure - High privacy risk undermines the core value proposition of synthetic data - Privacy breaches carry regulatory penalties, reputational damage, and individual harm

What SDCF measures: - **Membership inference risk:** Can adversary determine if specific individual’s data was in source dataset? - **Record similarity:** Do synthetic records closely resemble real individuals (enabling re-identification via auxiliary data)? - **Attribute disclosure:** Can adversary infer sensitive attributes for known individuals?

Privacy Risk Score (PRS): Composite metric (0-100 scale) combining membership inference, record similarity, and attribute disclosure tests. Lower scores indicate lower risk.

Thresholds (provisional): - PRS < 20: Low risk (strong privacy protection) - PRS 20-50: Moderate risk (acceptable for controlled disclosure) - PRS 50-80: High risk (only for internal use with additional controls) - PRS > 80: Very high risk (reconsider synthetic data approach)

Trade-offs: - Stronger privacy protection (lower PRS) may reduce fidelity (more noise, less realistic) - Perfect privacy (PRS = 0) likely requires sacrificing utility - Purpose determines acceptable risk level (external sharing demands PRS < 20; internal testing may tolerate PRS < 50)

See Appendix A.1 for mathematical definition and Appendix C.2 for Bronze Tier assessment without source data.

4.2.2 Pillar 2: Statistical Fidelity

Definition: The degree to which synthetic data preserves the statistical properties and relationships of the source data.

Why it matters: - Analytical insights depend on accurate representation of distributions, correlations, and patterns - ML models trained on low-fidelity data produce poor predictions - Business decisions based on distorted statistics lead to costly errors - EU AI Act Article 10 requires data be “relevant, sufficiently representative”

What SDCF measures: - **Distribution similarity:** Do marginal and joint distributions match source data? - **Dependency preservation:** Are correlations and variable relationships maintained? - **Predictive utility:** Do models trained on synthetic data perform comparably to models trained on real data?

Fidelity Index (FI): Composite metric (0-100 scale) combining distribution similarity, dependency preservation, and predictive utility tests. Higher scores indicate higher fidelity.

Thresholds (provisional): - FI > 80: High fidelity (suitable for critical business/regulatory decisions) - FI 60-80: Moderate fidelity (acceptable for most analytics, model development) - FI 40-60: Low fidelity (suitable for prototyping, testing, non-critical analysis) - FI < 40: Insufficient fidelity (reconsider synthetic data approach)

Trade-offs: - Higher fidelity (higher FI) may increase privacy risk (more realistic = more identifiable) - Perfect fidelity (FI = 100) means copying source data (PRS = 100, defeats purpose) - Purpose determines required fidelity level (regulatory stress testing demands FI > 80; software testing may accept FI > 60)

See Appendix A.2 for mathematical definition and Appendix C.3 for Bronze Tier assessment without source data.

4.2.3 Pillar 3: Algorithmic Fairness

Definition: The degree to which synthetic data represents diverse populations equitably and does not amplify bias in downstream uses.

Why it matters: - EU AI Act Article 10(4) requires examining datasets for “possible biases” - Biased training data produces discriminatory AI systems (violates Charter of Fundamental Rights) - Synthetic data generation can amplify source data bias if not carefully managed - Ethical obligation to ensure data-driven systems serve all communities fairly

What SDCF measures: - **Representation variance:** Are protected groups (gender, ethnicity, age, disability) represented proportionally or as required? - **Predictive parity:** Do models trained on synthetic data show disparate performance across groups?

Fairness Variance (FV): Composite metric (0-100 scale) combining representation analysis and predictive parity tests. Lower scores indicate lower fairness concerns.

Thresholds (provisional): - FV < 15: Low fairness concerns (well-balanced representation and outcomes) - FV 15-30: Moderate concerns (acceptable for non-high-risk applications, document limitations) - FV 30-50: Significant concerns (suitable only with bias mitigation measures) - FV > 50: Severe fairness issues (reconsider synthetic data approach or generation methodology)

Trade-offs: - Improving fairness through oversampling minorities may reduce fidelity to real-world distributions - “Fairness through blindness” (removing protected attributes) may not prevent proxy discrimination - Purpose determines fairness priority (high-risk AI systems demand FV < 15; internal operations analytics may accept FV < 30)

Context-Specific Fairness: Fairness is highly context-dependent: - Healthcare: Representation across demographics, rare conditions - Finance: No disparate impact in credit/insurance decisions - Public sector: Equal service delivery across all communities - AI training: Balanced data prevents discriminatory model predictions

See Appendix A.3 for mathematical definition and Appendix C.4 for Bronze Tier assessment.

4.2.4 The Three-Pillar Balance

SDCF requires organisations to explicitly navigate trade-offs:

Privacy vs. Fidelity Tension: - Maximum privacy (heavy noise addition) destroys fidelity - Maximum fidelity (minimal noise) provides weak privacy - Sweet spot depends on purpose: External sharing prioritises privacy; regulatory reporting prioritises fidelity

Fidelity vs. Fairness Tension: - Real-world data often reflects historical bias (underrepresentation, disparate outcomes) - Perfect fidelity reproduces bias; fairness intervention distorts distributions - Synthetic data offers opportunity to be “less faithful to unfair reality”

Purpose-Bounded Resolution: C1 Purpose Sheet forces explicit decision: For THIS purpose, how do we prioritise the three pillars?

Example Trade-Off Decision: *Medical AI Training Dataset for Diabetic Retinopathy Detection*
- **Privacy:** High (patient health data, Article 9 GDPR) → Target PRS < 20 - **Fidelity:** Critical (diagnostic accuracy depends on realistic images) → Target FI > 85 - **Fairness:** Critical (disease prevalence varies by ethnicity, must ensure equitable detection) → Target FV < 12

Decision: Accept moderate privacy-fidelity trade-off (PRS = 25, slightly above target) to maintain both diagnostic utility and fair representation. Document this decision in C2 Governance Record with rationale and compensating controls (restricted access, additional audit logging).

4.3 Assessment Tiers: Gold, Silver, Bronze

Real-world synthetic data validation faces a critical constraint: **availability of source data for comparison.**

Organisations often need to assess third-party synthetic datasets, legacy synthetic data where source is unavailable, or synthetic data where privacy regulations prohibit source access even for validation. SDCF addresses this with a tiered methodology.

4.3.1 Gold Tier: Full Source Data Access

Definition: Assessor has complete access to both source data and synthetic data for rigorous comparison.

When applicable: - Organisation generated synthetic data in-house and retains source data - Assessment conducted by data controller with appropriate access rights - Validation performed in secure environment (trusted execution, on-premises)

Capabilities: - **Privacy Risk Score (PRS):** Full membership inference, record similarity, and attribute disclosure testing - **Fidelity Index (FI):** Direct distribution comparison, correlation preservation, predictive utility benchmarking - **Fairness Variance (FV):** Comprehensive representation and predictive parity analysis across source and synthetic

Confidence Level: Highest (mathematical rigor, minimal estimation)

Evidence Quality: Suitable for: - Regulatory submissions requiring highest assurance - High-risk AI system validation (EU AI Act compliance) - External auditor review - Scientific publication

Limitations: - Requires source data access (not always feasible) - Privacy risk: Validation process itself may expose source data - Resource-intensive: Comprehensive testing requires significant computation

Use Gold Tier when: - Purpose demands highest confidence (regulatory reporting, high-risk AI) - Source data is available and can be accessed securely - Resources permit comprehensive assessment

4.3.2 Silver Tier: Partial Source Data Access

Definition: Assessor has access to aggregate statistics, metadata, or samples from source data, but not complete dataset.

When applicable: - Source data summaries available (schema, distributions, correlations) but not row-level data - Sample-based validation (access to random sample, not full population) - Metadata-only scenarios (data dictionary, aggregate statistics, but no records)

Capabilities: - **PRS (Limited):** Membership inference not possible; record similarity against samples; attribute disclosure against known distributions - **FI (Moderate):** Distribution comparison against aggregate statistics; correlation validation against published summaries - **FV (Moderate):** Representation analysis against demographic summaries; limited predictive parity testing

Confidence Level: Moderate (some estimation required, statistical inference from partial information)

Evidence Quality: Suitable for: - Internal decision-making with documented limitations - Partner data sharing with transparent disclosure - Model development (non-high-risk systems) - Preliminary assessment before Gold Tier investment

Limitations: - Cannot perform rigorous membership inference tests - Distribution comparisons rely on aggregate statistics (may miss outliers, tails) - Predictive utility testing limited to sampled data

Use Silver Tier when: - Full source access infeasible (privacy, commercial, technical constraints) - Purpose tolerates moderate uncertainty (internal analytics, development environments) - Cost-benefit doesn't justify Gold Tier comprehensive assessment

4.3.3 Bronze Tier: No Source Data Access (Synthetic-Only)

Definition: Assessor has access ONLY to synthetic data; no source data, aggregates, or metadata available.

When applicable: - Third-party synthetic datasets (purchased, licensed, open-source) - Legacy synthetic data where source no longer exists or accessible - Privacy regulations prohibit source data access even for validation (healthcare, financial) - Synthetic training data for AI (common scenario in 2025+ as per Baguettotron model)

Capabilities: - **B-PRS (Conservative):** Outlier analysis, uniqueness scoring, conservative penalty for uncertainty - **B-FI (Internal):** Internal consistency checks, domain validity, benchmark comparison - **B-FV (Representation-Only):** Representation analysis without source comparison

Confidence Level: Lower (significant assumptions, conservative risk classification)

Evidence Quality: Suitable for: - Lower-risk use cases (software testing, training environments, prototyping) - Preliminary screening (decide whether to invest in Silver/Gold assessment) - Vendor evaluation (compare multiple synthetic data providers) - AI training data pre-flight check (validate before expensive training runs)

Bronze Tier Methodology Principles: 1. **Honest about limitations:** Certificate explicitly states reduced confidence level 2. **Conservative risk classification:** When uncertain, assume higher risk 3. **Focus on red flags:** Identify clear problems (duplicates, impossibilities, extreme

outliers) 4. **Domain validation:** Assess plausibility against known real-world constraints 5. **Enhanced controls required:** Bronze Tier assessment triggers stronger usage restrictions

Bronze Tier fills critical real-world gap: Most valuable synthetic data validation scenario is third-party datasets without source access. Existing frameworks ignore this; SDCF addresses it honestly.

Use Bronze Tier when: - No source data access possible (third-party data, privacy constraints, legacy scenarios) - Purpose is lower-risk (testing, development, preliminary analysis, AI training data screening) - Rapid assessment needed (days, not weeks) - Budget constraints preclude comprehensive validation

See Appendix C for complete Bronze Tier methodology, metrics, certificate templates, and case studies.

4.3.4 Tier Selection Decision Tree

START: Do you need to assess synthetic data?

```
|
| Is source data available?
|   YES, full access → Consider Gold Tier
|   |   Is purpose high-risk? (regulatory, high-risk AI, external sharing)
|   |   YES → Use Gold Tier
|   |   NO → Gold Tier preferred, but Silver acceptable with justification
|   |
|   YES, partial (aggregates/samples) → Consider Silver Tier
|   |   Is partial information sufficient for purpose?
|   |   YES → Use Silver Tier
|   |   NO → Either obtain full access (Gold) or use Bronze with limitations
|   |
|   NO access → Use Bronze Tier
|   |   Is purpose appropriate for Bronze confidence level?
|   |   YES → Use Bronze Tier, document limitations
|   |   NO → Either obtain source access OR reconsider synthetic data approach
```

4.4 Conformance Levels: SDCF-A, SDCF-P, SDCF-R

After assessment (Gold, Silver, or Bronze Tier), SDCF assigns a conformance level indicating fitness for stated purpose.

Interpreting Conformance Levels:

- **Acceptable (SDCF-A):** High confidence for intended purpose without additional controls
- **Provisional (SDCF-P):** Suitable for purpose with documented limitations and mitigations
- **Restricted (SDCF-R):** Requires substantial additional controls OR only suitable for different, more limited purpose

Note: Restricted conformance does not mean “failed” - it reflects SDCF’s conservative design philosophy that prioritises privacy protection and honest risk communication over false confidence. A 60% Restricted rate in Bronze Tier validation confirms the framework correctly identifies datasets requiring enhanced controls rather than providing blanket approval.

4.4.1 SDCF-A (Approved)

Definition: Synthetic data is suitable for the stated purpose without restrictions.

Criteria: - All three pillars meet or exceed purpose-specific thresholds - No significant limitations or concerns identified - Assessment tier appropriate for purpose (Gold for high-risk, Silver/Bronze for appropriate contexts) - Governance controls (C2) and release rules (C7) documented and appropriate

Certificate Statement: > “This synthetic dataset has been assessed using SDCF [Gold/Silver/Bronze] Tier methodology and is APPROVED for the stated purpose: [purpose description]. Assessment conducted [date] by [assessor]. Valid until [date] or until material changes to dataset, purpose, or regulatory environment.”

Implications: - Dataset may be used for stated purpose with documented controls - Transparency Pack (C6) provided to stakeholders/users - Periodic reassessment recommended (annually or upon material changes) - Use beyond stated purpose requires new assessment

Example: *Synthetic patient data for internal clinical research dashboard (Gold Tier, SDCF-A)* - PRS = 15 (target < 20) - FI = 88 (target > 85) - FV = 10 (target < 12) - Assessment: Approved for internal research analytics with documented access controls

4.4.2 SDCF-P (Provisional)

Definition: Synthetic data is suitable for stated purpose WITH documented limitations and additional controls.

Criteria: - One or more pillars slightly below ideal but within acceptable range - Specific limitations identified that do not preclude use but require disclosure - Compensating controls implemented to manage residual risk - Assessment tier appropriate, but some uncertainty remains

Certificate Statement: > “This synthetic dataset has been assessed using SDCF [Gold/Silver/Bronze] Tier methodology and is PROVISIONALLY APPROVED for the stated purpose: [purpose description], subject to the following limitations and required controls: [list]. Assessment conducted [date] by [assessor]. Valid until [date] or until material changes.”

Implications: - Dataset may be used WITH documented restrictions - Enhanced monitoring or audit requirements - Users must be informed of specific limitations - More frequent reassessment (e.g., semi-annually) - Limitations must be disclosed to all stakeholders

Example: *Synthetic transaction data for AML model training (Silver Tier, SDCF-P)* - PRS = 28 (target < 25) Slightly above target - FI = 75 (target > 70) - FV = 18 (target < 20) - Limitations: Privacy risk slightly elevated due to rare transaction patterns; rare event representation confirmed adequate - Controls: Restrict access to model development team only; enhanced audit logging; do not use for regulatory reporting - Assessment: Provisionally approved for model training (not production deployment) with access restrictions

4.4.3 SDCF-R (Restricted)

Definition: Synthetic data is NOT suitable for stated purpose OR suitable only for a different, more limited purpose.

Criteria: - One or more pillars significantly fail to meet purpose-specific thresholds - Fundamental limitations that cannot be mitigated through controls - Assessment tier inadequate for purpose requirements (e.g., Bronze Tier for high-risk AI) - Risk profile unacceptable for stated purpose

Certificate Statement: > “This synthetic dataset has been assessed using SDCF [Gold/Silver/Bronze] Tier methodology and is RESTRICTED for the stated purpose: [purpose description]. Assessment conducted [date] by [assessor]. The dataset does NOT meet requirements for the intended use. Identified concerns: [list]. Alternative suitable uses (if any): [list].”

Implications: - Dataset MUST NOT be used for stated purpose - May suggest alternative lower-risk purposes (if appropriate) - Requires remediation (regenerate synthetic data, improve methodology) or abandonment - Certificate serves as evidence of due diligence in decision NOT to use dataset

Example 1: Unsuitable for Any Use *Synthetic health data for medical device training (Gold Tier, SDCF-R)* - PRS = 75 (target < 20) Unacceptably high - FI = 45 (target > 80) Insufficient fidelity - FV = 52 (target < 15) Severe fairness issues - Assessment: RESTRICTED. Dataset unsuitable for medical device training due to high privacy risk, poor fidelity, and fairness concerns. Recommend regenerating with improved methodology.

Example 2: Suitable for Limited Alternative Use *Synthetic claims data intended for pricing model (Bronze Tier, SDCF-R for pricing, SDCF-P for testing)* - B-PRS = 35 (acceptable for internal use, not pricing) - B-FI = 58 (insufficient for pricing decisions, adequate for testing) - B-FV = 22 (acceptable for testing) - Assessment: RESTRICTED for pricing model use (Bronze Tier inadequate, fidelity too low). However, PROVISIONALLY APPROVED for software testing and development environments only.

4.4.4 Conformance Level Decision Matrix

Pillar Status	Gold Tier	Silver Tier	Bronze Tier
All pillars meet targets	SDCF-A	SDCF-A (if purpose appropriate)	SDCF-P (Bronze inherent limitations)
One pillar slightly below, compensable	SDCF-P	SDCF-P	SDCF-P or SDCF-R (depends on severity)
One pillar significantly below	SDCF-R	SDCF-R	SDCF-R
Multiple pillars below targets	SDCF-R	SDCF-R	SDCF-R

Note: Conformance level also considers purpose-tier alignment (Bronze Tier for high-risk AI automatically triggers SDCF-R for that purpose, regardless of scores).

4.5 Integration Points

SDCF is designed to integrate into existing organisational processes and complement other frameworks rather than replace them.

4.5.1 Integration with Data Governance

Existing Process: Organisations typically have data governance frameworks covering data quality, metadata management, access control, and lifecycle management.

SDCF Integration: - **C1 Purpose Sheet** documents synthetic data asset in governance inventory - **C2 Governance Record** captures decisions, approvals, roles (feeds into governance documentation) - **C7 Release Rules** integrates with access control policies and data classification schemes - **SDCF Certificate** becomes part of data asset metadata (lineage, quality attributes)

Benefit: Synthetic data governed consistently with other organisational data assets.

4.5.2 Integration with Risk Management

Existing Process: Organisations conduct risk assessments for data processing, third-party vendors, and new initiatives.

SDCF Integration: - **Privacy Risk Score (PRS)** quantifies privacy risk for inclusion in enterprise risk register - **Fidelity/Fairness concerns** surface operational and reputational risks - **SDCF-R rating** triggers risk escalation and mitigation planning - **Assessment results** inform Data Protection Impact Assessments (DPIA)

Benefit: Synthetic data risks managed within enterprise risk management framework.

4.5.3 Integration with AI Governance

Existing Process: Organisations developing AI/ML systems have model governance frameworks covering development, validation, deployment, and monitoring.

SDCF Integration: - **C1 Purpose Sheet** documents training data quality for model cards and system documentation - **Fidelity Index (FI)** assesses training data representativeness (EU AI Act Article 10) - **Fairness Variance (FV)** validates bias mitigation efforts - **SDCF Certificate** provides evidence for model risk management and regulatory compliance

Benefit: Training data quality integrated into end-to-end AI system lifecycle governance.

4.5.4 Integration with Vendor Management

Existing Process: Organisations procure third-party services and data products with due diligence, contracts, and ongoing monitoring.

SDCF Integration: - **Bronze Tier assessment** evaluates vendor synthetic data products before purchase - **Certificate requirements** included in procurement RFPs (vendors must provide SDCF assessment or undergo assessment) - **Conformance level** determines vendor risk rating and contract terms - **Reassessment cadence** built into vendor review schedules

Benefit: Standardised evaluation criteria for synthetic data vendors.

4.5.5 Integration with Compliance and Audit

Existing Process: Organisations demonstrate compliance through documentation, control testing, and audit evidence.

SDCF Integration: - **C6 Transparency Pack** provides auditable trail for synthetic data decisions - **SDCF Certificate** evidences due diligence for GDPR Article 32 (security) and Article 25 (data protection by design) - **Assessment methodology** demonstrates “appropriate technical measures” for anonymisation claims - **Periodic reassessment** shows ongoing accountability (GDPR Article 5(2))

Benefit: Defensible evidence package for regulators and external auditors.

4.5.6 Integration with Existing Tools

SDCF does NOT replace metric computation tools; it interprets their outputs for compliance purposes.

Complementary Tools:

SDMetrics / SDV (Open Source): - Provides fidelity and privacy metrics that feed into FI and PRS calculations - SDCF methodology can consume SDMetrics reports - See Appendix F for reference implementation

mostlyai-qa (Open Source): - Generates quality assurance reports for synthetic data - SDCF can ingest mostlyai-qa outputs for Silver/Bronze tier assessments - See Appendix F for integration guidance

Vendor Platforms (Gretel, MOSTLY AI, Syntho): - Generate synthetic data and provide quality reports - SDCF provides independent validation methodology applicable to any generator’s outputs - Vendor reports can inform but do not replace SDCF assessment (independence requirement)

SDCF Value-Add: - Purpose-bounded assessment (fitness for specific use case, not generic quality) - Regulatory mapping (GDPR Article 6/9, EU AI Act Article 10, sector standards) - Tiered methodology (Gold/Silver/Bronze) handling real-world data availability - Evidence pack structure for demonstrating compliance to auditors/regulators - Transparency about limitations and provisional thresholds

4.5.7 Standards Alignment

ISO/IEC 27001/27701: - SDCF Control Sets (C1-C7) can be mapped to ISO controls (Annex A) - Assessment process supports ISO audit evidence requirements - See Appendix D for detailed mapping

NIST Privacy Framework / AI RMF: - SDCF three-pillar approach aligns with NIST trustworthy AI characteristics - Assessment results inform NIST framework implementation - See Appendix D for detailed mapping

Benefit: Organisations with existing ISO/NIST programs can adopt SDCF without conflicting frameworks.

End of Section 3

Continue to Section 4: Control Sets (C1-C7) to understand the specific procedures and documentation requirements for SDCF assessment.

5 Control Sets (C1-C7)

SDCF defines seven control sets that operationalise the assessment methodology. These controls provide structure for documenting purpose, conducting technical testing, managing governance, and maintaining transparency.

5.1 C1: Purpose Sheet

5.1.1 What It Is

The Purpose Sheet is a structured document that explicitly defines the intended use of synthetic data before assessment begins. It serves as the foundation for all subsequent evaluation decisions.

5.1.2 Why It Matters

Without clear purpose definition: - Assessment criteria become arbitrary (what thresholds to apply?) - Risk evaluation lacks context (is PRS=30 acceptable?) - Scope creep occurs (dataset used beyond original intent) - Accountability is weakened (who decided this was appropriate?)

The Purpose Sheet prevents these problems by establishing clear success criteria upfront.

5.1.3 Required Elements

1. Use Case Description - What will this synthetic data be used for? - What specific decisions, processes, or systems depend on it? - What business or research questions will it answer?

Example: “Train anti-money laundering (AML) detection model to identify suspicious transaction patterns in retail banking. Model will flag accounts for investigation by compliance team.”

2. Regulatory Context - What regulations apply? (GDPR, EU AI Act, sector-specific) - Is this a high-risk AI system? (EU AI Act Annex III) - What compliance obligations must be demonstrated?

Example: “EU AI Act high-risk system (credit scoring). Must demonstrate Article 10 data governance compliance. GDPR Article 6(1)(f) legitimate interest as lawful basis.”

3. Disclosure Scope - Who will have access? (internal only, partners, public) - What is the disclosure environment? (secure network, internet, edge devices) - What adversary capabilities must be assumed?

Example: “Internal use only. Access restricted to data science team (8 people) in secure development environment. Assume insider threat scenario (authorised user attempts re-identification).”

4. Pillar Prioritization - What is the relative importance of Privacy, Fidelity, Fairness for THIS purpose? - What specific targets for PRS, FI, FV? - What trade-offs are acceptable?

Example: - Privacy: HIGH (financial data, Article 9 potential) → Target PRS < 25 - Fidelity: CRITICAL (model effectiveness depends on realistic patterns) → Target FI > 80 - Fairness: CRITICAL (must not discriminate) → Target FV < 15

5. Failure Modes - What are the consequences if assessment fails? - What specific risks must be avoided? - What would constitute unacceptable outcome?

Example: “Failure modes: (1) Privacy breach enables customer re-identification → regulatory penalty, reputational damage; (2) Poor fidelity causes model to miss money laundering → financial crime; (3) Bias causes disparate impact on ethnic minorities → discrimination complaint.”

6. Assessment Tier Selection - Is source data available? (Gold/Silver/Bronze) - What tier is appropriate given purpose and resources? - If tier and purpose don’t align, how to resolve?

Example: “Source data available. Given high-risk nature, Gold Tier assessment required.”

7. Success Criteria - What conformance level is needed? (SDCF-A required? SDCF-P acceptable?) - What limitations are tolerable? - When would dataset be unsuitable?

Example: “SDCF-A (Approved) required for production use. SDCF-P acceptable for pilot with enhanced monitoring. SDCF-R triggers dataset regeneration.”

5.1.4 Implementation Guidance

Timing: Complete Purpose Sheet BEFORE generating or acquiring synthetic data. Influences generation methodology decisions.

Ownership: Typically owned by business or research lead (defines purpose) with input from DPO, legal, and data science.

Format: Structured template (2-4 pages). See Appendix E for example.

Review: Purpose Sheet reviewed by governance board (C2) before proceeding to technical assessment.

Living Document: Update if purpose evolves or dataset is repurposed (triggers reassessment).

5.1.5 Common Mistakes to Avoid

Vague purpose: “Improve customer analytics” → Too broad, what specific analytics?

Specific purpose: “Generate monthly customer segmentation report for marketing team showing demographics and purchase behavior”

Generic targets: “Good privacy and quality” → Not measurable

Specific targets: “PRS < 30, FI > 75, FV < 20 for internal marketing analysis”

Purpose defined after assessment: Retrofitting rationale to justify results

Purpose first: Assessment criteria flow from documented purpose

5.2 C2: Governance Record

5.2.1 What It Is

The Governance Record documents the organisational decision-making process for synthetic data assessment and use. It captures who made decisions, what was considered, and why particular

approaches were chosen.

5.2.2 Why It Matters

GDPR Article 5(2) requires “accountability” - ability to demonstrate compliance. EU AI Act requires documented risk management. The Governance Record provides auditable evidence trail.

Without documented governance: - Decisions appear arbitrary to regulators - Organisational learning is lost (why did we accept that trade-off?) - Accountability is unclear (who approved this?) - Risk management is reactive rather than proactive

5.2.3 Required Elements

1. Roles and Responsibilities - Who requested synthetic data? (business sponsor) - Who conducted assessment? (technical assessor, qualifications) - Who reviewed and approved? (governance board, DPO, legal) - Who is accountable for ongoing monitoring?

Example: - Sponsor: Head of Fraud Analytics - Assessor: Senior Data Scientist (Jane Smith, SDCF trained) - Reviewers: DPO (GDPR compliance), Legal Counsel (AI Act compliance), CISO (security) - Approver: Chief Risk Officer - Monitor: Data Governance Team (quarterly reviews)

2. Decision Log - What decisions were made during assessment? - What alternatives were considered? - What was the rationale for chosen approach? - Were there dissenting views?

Example Entry:

“Decision: Accept PRS=28 despite target PRS<25. Rationale: Fidelity would degrade unacceptably (FI drops from 82 to 68) if more noise added. Compensating controls: Restrict to development environment only (not production), enhanced audit logging, 6-month review cycle instead of annual. Dissent: DPO preferred regenerating dataset; overruled by CRO given business need and acceptable residual risk with controls.”

3. Risk Assessment - What risks were identified? - How were they evaluated? (likelihood × impact) - What mitigations were implemented? - What residual risks remain?

Example: | Risk | Likelihood | Impact | Mitigation | Residual | |——|———|———|———|———|———|
———| | Re-identification | Low | High | PRS=28, access controls, monitoring | Acceptable | | Model bias | Medium | Medium | FV=12, fairness testing, human review | Acceptable | | Data quality degradation over time | Low | Medium | Quarterly reassessment | Acceptable |

4. Regulatory Mapping - What regulatory requirements apply? - How does this assessment address them? - What evidence is available for compliance demonstration?

Example: - GDPR Article 6(1)(f): Legitimate interest assessment completed, balancing test documented - GDPR Article 32: Technical measures (synthetic data, access controls) demonstrate appropriate security - EU AI Act Article 10: Data governance evidenced through SDCF assessment, documented in model card

5. Approval Gateway - What were the decision criteria? - Who approved and when? - Under what conditions must reassessment occur? - What are the ongoing monitoring requirements?

Example:

“Approved by CRO on [date]. Conditions: (1) Use restricted to stated purpose only; (2) Quarterly

monitoring reports to governance board; (3) Reassessment required if: regulatory guidance changes, dataset is modified, purpose evolves, or security incident occurs.”

5.2.4 Implementation Guidance

Format: Structured log in data governance system or risk register. Not standalone document - integrate with existing processes.

Update Frequency: - Initial: During assessment - Ongoing: When decisions made, risks change, or incidents occur - Review: At reassessment intervals (annual, semi-annual, quarterly)

Access Control: Governance Record contains sensitive decision-making rationale. Restrict to governance board, audit, and regulators only.

Integration: Link to other governance artifacts (DPIA, risk register, data inventory, model cards).

5.2.5 Example Governance Record Entry

SYNTHETIC DATA GOVERNANCE RECORD

Dataset ID: SYNTH-AML-2025-Q4

Purpose: AML model training (see Purpose Sheet v1.2)

Assessment Date: 2025-11-15

Tier: Gold

Assessor: Jane Smith (Senior Data Scientist)

DECISION LOG:

[2025-11-10] Scoping: Determined Gold Tier required for high-risk AI system

[2025-11-12] Trade-off: Accepted PRS=28 vs target 25; rationale documented

[2025-11-14] Fairness: Implemented oversampling for minority representation

[2025-11-15] Approval: CRO approved with conditions (dev env only)

RISK ASSESSMENT:

- Re-identification: LOW likelihood, HIGH impact → MITIGATED (PRS=28, controls)

- Bias: MEDIUM likelihood, MEDIUM impact → MITIGATED (FV=12, review process)

CONFORMANCE: SDCF-P (Provisional - privacy slightly above target)

VALIDITY: Until 2026-05-15 or material change

NEXT REVIEW: 2026-02-15

5.3 C3: Privacy Risk Testing

5.3.1 What It Is

Privacy Risk Testing conducts technical measurements to quantify the risk that synthetic data enables re-identification or attribute disclosure. Results feed into the Privacy Risk Score (PRS).

5.3.2 Why It Matters

The core value proposition of synthetic data is privacy protection. If synthetic data doesn’t meaningfully reduce privacy risk, it’s not fit for purpose. Privacy testing provides empirical evidence rather than assumptions.

5.3.3 Required Tests

Test 1: Membership Inference *Question:* Can adversary determine if a specific individual's data was in the source dataset?

Method (Gold Tier): 1. Train attack model on source + synthetic data characteristics 2. Attempt to classify synthetic records as "derived from real record X" 3. Measure attack success rate vs. random guessing

Metric: Membership inference success rate (0-100%)

Threshold: <15% for low risk, 15-30% moderate, >30% high risk

Interpretation:

- 50% = random guessing (no signal) - >60% = adversary can reliably identify membership - <10% = strong membership privacy

Bronze Tier Alternative (B-PRS Component 1):

Without source data, use outlier analysis as proxy: - Identify records with unusual combinations of attributes - Score uniqueness using local outlier factor - Conservative penalty: Assume outliers = higher membership risk

Test 2: Record Similarity / Distance to Closest Record (DCR) *Question:* How similar are synthetic records to their nearest real records?

Method (Gold Tier): 1. For each synthetic record, find closest source record (Euclidean or Gower distance) 2. Compute distance distribution 3. Identify synthetic records "too close" to real records

Metric: Percentage of synthetic records within threshold distance

Threshold: <5% within 0.1 distance for low risk, 5-15% moderate, >15% high risk

Interpretation: - Close records = adversary with auxiliary data could re-identify - Distance depends on normalization and attribute sensitivity - Rare/unique source individuals present higher risk

Bronze Tier Alternative (B-PRS Component 2):

Without source data, assess internal uniqueness: - Identify duplicate or near-duplicate synthetic records (red flag) - Score attribute combination rarity within synthetic data - Flag records with implausible or extreme values

Test 3: Attribute Disclosure *Question:* Can adversary infer sensitive attributes for known individuals?

Method (Gold Tier): 1. Assume adversary knows some attributes (quasi-identifiers: age, location, gender) 2. Test if synthetic data reveals sensitive attributes (health, income, ethnicity) for matching quasi-identifiers 3. Compare disclosure rate to prior probability (population base rate)

Metric: Disclosure rate increase over baseline

Threshold: <10% increase for low risk, 10-25% moderate, >25% high risk

Bronze Tier Alternative (B-PRS Component 3):

Assess correlation strength between quasi-identifiers and sensitive attributes in synthetic data. High correlation = higher disclosure risk if adversary has auxiliary data.

5.3.4 Composite Privacy Risk Score (PRS)

Formula (Gold Tier):

$$\text{PRS} = (w \times \text{MembershipScore}) + (w \times \text{SimilarityScore}) + (w \times \text{DisclosureScore})$$

Where:

- MembershipScore = Membership inference success rate (normalized 0-100)
- SimilarityScore = % records within risk threshold (normalized 0-100)
- DisclosureScore = Attribute disclosure increase (normalized 0-100)
- Weights: $w=0.4$, $w=0.4$, $w=0.2$ (adjustable based on purpose)

PRS Interpretation: - PRS < 20: Low privacy risk - PRS 20-50: Moderate privacy risk - PRS 50-80: High privacy risk - PRS > 80: Very high privacy risk

See Appendix A.1 for detailed mathematical formulation and confidence intervals.

5.3.5 Implementation Guidance

Tooling: - Use SDMetrics privacy metrics (membership inference, DCR) - Alternatively: Python libraries (numpy, scikit-learn) for custom implementation - See Appendix F for reference code

Computational Cost: - Gold Tier: High (thousands of attack simulations) - Silver Tier: Moderate (sample-based testing) - Bronze Tier: Low (statistical analysis only)

Expertise Required: - Understanding of privacy attack models - Statistical analysis capabilities - Domain knowledge to interpret results in context

Common Pitfalls: - Testing with insufficient adversary assumptions (underestimate risk) - Ignoring outliers and rare combinations (highest risk individuals) - Failing to consider auxiliary data availability in deployment context

5.4 C4: Fidelity Testing

5.4.1 What It Is

Fidelity Testing evaluates how well synthetic data preserves the statistical properties and relationships of source data. Results feed into the Fidelity Index (FI).

5.4.2 Why It Matters

Low-fidelity synthetic data produces misleading insights and poor model performance. Organisations make wrong decisions when analytical results don't reflect reality. Fidelity testing quantifies utility preservation.

5.4.3 Required Tests

Test 1: Distribution Similarity Question: Do marginal and joint distributions match between source and synthetic?

Method (Gold Tier): 1. Compare univariate distributions (histograms, KDEs) for each variable 2. Use statistical tests: Kolmogorov-Smirnov (continuous), Chi-square (categorical) 3. Compute distributional similarity metrics (Jensen-Shannon divergence, Wasserstein distance) 4. Test bivariate distributions for key variable pairs

Metric: Average distributional similarity score (0-100%)

Threshold: >85% for high fidelity, 70-85% moderate, <70% low fidelity

Interpretation: - 100% = identical distributions (likely copied data) - 90-95% = very high fidelity (distributions well-preserved) - 70-85% = moderate fidelity (broad patterns match, some details lost) - <70% = poor fidelity (distributions significantly distorted)

Bronze Tier Alternative (B-FI Component 1):

Without source data, perform internal consistency checks: - Do distributions match known population characteristics? (e.g., age distribution matches census) - Are there logical impossibilities? (negative ages, invalid dates) - Do relationships make domain sense? (income correlates with education)

Test 2: Dependency Preservation Question: Are correlations and variable relationships maintained?

Method (Gold Tier): 1. Compute correlation matrices for source and synthetic data 2. Compare correlation preservation (Pearson for continuous, Cramér's V for categorical) 3. Test for interaction effects and non-linear relationships 4. Validate preservation of known domain relationships

Metric: Correlation preservation score (0-100%)

Threshold: >80% for high fidelity, 65-80% moderate, <65% low fidelity

Interpretation: - Strong correlations more important than weak ones (weighted scoring) - Domain-critical relationships (e.g., risk factors in healthcare) must be preserved - Some noise acceptable for privacy-fidelity trade-off

Bronze Tier Alternative (B-FI Component 2):

Test internal correlation structure against domain expectations: - Do expected relationships exist? (education income, age chronic conditions) - Are correlations plausible in magnitude? - Are there spurious correlations that violate domain knowledge?

Test 3: Predictive Utility Question: Do models trained on synthetic data perform comparably to models trained on real data?

Method (Gold Tier): 1. Train model on source data, test on holdout real data (baseline performance) 2. Train same model on synthetic data, test on same holdout real data 3. Compare performance metrics (accuracy, F1, AUC-ROC, RMSE) 4. Repeat for multiple model types (linear, tree-based, neural network)

Metric: Predictive utility preservation (0-100%)

Threshold: >85% for high utility, 70-85% moderate, <70% low utility

Interpretation: - 95%+ = nearly identical performance (high utility) - 80-90% = acceptable degradation for most use cases - <70% = significant utility loss (models learn wrong patterns)

Bronze Tier Alternative (B-FI Component 3):

Train models on synthetic data only, evaluate against: - Known benchmarks from literature (if available) - Domain expert expectations for performance - Cross-validation consistency (if model generalises, data may have internal validity)

5.4.4 Composite Fidelity Index (FI)

Formula (Gold Tier):

$$FI = (w \times \text{DistributionScore}) + (w \times \text{DependencyScore}) + (w \times \text{UtilityScore})$$

Where:

- DistributionScore = Average distributional similarity (0-100)
- DependencyScore = Correlation preservation (0-100)
- UtilityScore = Predictive utility preservation (0-100)
- Weights: $w=0.3$, $w=0.3$, $w=0.4$ (adjustable based on purpose)

FI Interpretation: - FI > 80: High fidelity - FI 60-80: Moderate fidelity - FI 40-60: Low fidelity
 - FI < 40: Insufficient fidelity

See Appendix A.2 for detailed mathematical formulation.

5.4.5 Implementation Guidance

Tooling: - Use SDMetrics quality reports (distribution similarity, correlation) - mostly ai-qa provides comprehensive fidelity metrics - See Appendix F for reference implementations

Domain Expertise Critical: - Which distributions matter most? (prioritise key variables) - What relationships are essential? (clinical risk factors, financial ratios) - What model performance is “good enough”? (context-dependent)

Purpose-Specific Testing: - Analytics use case: Focus on distribution similarity - ML training use case: Focus on predictive utility - Stress testing use case: Ensure tail events preserved

5.5 C5: Fairness Assessment

5.5.1 What It Is

Fairness Assessment evaluates whether synthetic data represents diverse populations equitably and doesn’t amplify bias. Results feed into the Fairness Variance (FV).

5.5.2 Why It Matters

EU AI Act Article 10(4) mandates examining datasets for “possible biases.” Biased synthetic data produces discriminatory AI systems. Even non-AI uses can perpetuate inequities if synthetic data misrepresents populations.

5.5.3 Required Tests

Test 1: Representation Analysis *Question:* Are protected groups represented proportionally (or as required) in synthetic data?

Method (Gold Tier): 1. Identify protected attributes: gender, race/ethnicity, age, disability, etc. 2. Compare representation in source vs. synthetic data 3. Calculate representation variance for each group 4. Assess against requirements (proportional, balanced, or domain-specific targets)

Metric: Maximum representation deviation across groups

Threshold: <10% deviation for low variance, 10-25% moderate, >25% high variance

Interpretation: - Perfect proportionality not always desired (may want to oversample minorities for fairness) - Context matters: Medical research may need diverse representation; historical analysis may need proportional - Intersectional analysis critical (age × gender × ethnicity)

Bronze Tier Alternative (B-FV Component 1):

Assess representation against known population demographics: - Do protected group distributions match census/registry data? - Are there groups entirely missing or severely underrepresented? - Flag when representation diverges significantly from expectations

Test 2: Predictive Parity Question: Do models trained on synthetic data show disparate performance across protected groups?

Method (Gold Tier): 1. Train model on synthetic data 2. Evaluate performance separately for each protected group 3. Compare metrics: accuracy, false positive rate, false negative rate 4. Compute parity violations (maximum performance delta across groups)

Metric: Maximum parity violation (% performance difference)

Threshold: <10% for low concern, 10-20% moderate, >20% high concern

Interpretation: - Equal accuracy: All groups have similar model performance - Equal opportunity: True positive rates equal across groups - Equalised odds: Both TPR and FPR equal across groups - Choice depends on fairness definition appropriate for use case

Bronze Tier Alternative (B-FV Component 2):

Not possible without source data and model testing. Bronze Tier relies solely on representation analysis (Component 1). Certificate must disclose this limitation.

5.5.4 Composite Fairness Variance (FV)**Formula (Gold Tier):**

$$FV = (w \times RepresentationVariance) + (w \times ParityViolation)$$

Where:

- RepresentationVariance = Max deviation in representation (normalized 0-100)
- ParityViolation = Max performance difference across groups (normalized 0-100)
- Weights: w=0.5, w=0.5 (adjustable based on purpose)

Formula (Bronze Tier):

B-FV = RepresentationVariance (only)

Certificate notes: Predictive parity not assessed due to Bronze Tier limitations

FV Interpretation: - FV < 15: Low fairness concerns - FV 15-30: Moderate concerns - FV 30-50: Significant concerns - FV > 50: Severe fairness issues

See Appendix A.3 for detailed mathematical formulation.

5.5.5 Implementation Guidance

Protected Attributes: - EU context: Race/ethnicity, gender, age, disability, sexual orientation, religion - Identify which attributes are sensitive for YOUR use case - Consider intersectionality (combinations of attributes)

Fairness Definitions: - No universal fairness metric (context-dependent) - Healthcare: Representation across demographics, rare conditions - Credit: No disparate impact in approval rates - Criminal justice: Balance false positive/negative rates across groups - Choose fairness definition appropriate for purpose (document in C1)

Mitigation Strategies: If fairness assessment reveals concerns: - Regenerate with balanced sampling - Apply post-processing fairness interventions - Accept limitations and restrict use cases - Document limitations transparently (C6)

Common Pitfalls: - “Fairness through unawareness” (removing protected attributes doesn’t prevent proxy discrimination) - Optimizing for one fairness metric at expense of others (trade-offs exist) - Assuming proportional representation is always fair (context matters)

5.6 C6: Transparency Pack

5.6.1 What It Is

The Transparency Pack is a comprehensive documentation bundle that communicates assessment results, limitations, and usage guidance to stakeholders and users.

5.6.2 Why It Matters

GDPR Article 12 requires transparency. EU AI Act Article 13 mandates transparency and user information. External data sharing requires disclosure. Transparency Pack enables informed decision-making by all parties.

5.6.3 Required Components

- 1. Executive Summary (1 page)** - What is this synthetic data? - What was it assessed for? - What is the conformance level? (SDCF-A/P/R) - Key findings in plain language - Bottom-line recommendation
- 2. Assessment Certificate** - Formal statement of conformance - Assessment tier (Gold/Silver/Bronze) - Scores: PRS, FI, FV - Assessor details and date - Validity period - Conditions and restrictions
- 3. Technical Results** - Privacy Risk Score breakdown (membership, similarity, disclosure) - Fidelity Index breakdown (distribution, dependency, utility) - Fairness Variance breakdown (representation, parity) - Comparison tables and visualizations
- 4. Known Limitations** - What doesn’t this synthetic data do well? - What are the confidence bounds? - What assumptions underlie the assessment? - What changed from source data? (if known)
- 5. Usage Restrictions** - What can this data be used for? (stated purpose) - What must it NOT be used for? (out of scope) - What controls are required? (access, monitoring, audit) - When must reassessment occur?
- 6. Regulatory Mapping** - How does this address GDPR obligations? - How does this support EU AI Act compliance? - What sector-specific requirements are covered? - What evidence is available for auditors?
- 7. Contact and Escalation** - Who to contact with questions? - How to report issues or incidents? - How to request reassessment?

5.6.4 Implementation Guidance

Audience-Specific Versions: - **Executive version:** 2-3 pages, high-level findings - **Technical version:** Full report with methodology details - **Legal version:** Focus on regulatory compliance

mapping - **User version:** Plain language usage guidance

Format: - PDF for formal documentation (signed, version controlled) - Web page for ongoing reference (link from certificate) - Data package metadata (machine-readable JSON/XML)

Distribution: - Internal stakeholders: Governance board, business users, legal, audit - External parties: Partners, customers, regulators (as appropriate) - Public: Open data scenarios require public transparency pack

Update Cadence: - Material changes: When dataset, purpose, or regulations change - Periodic: Annual review even if no changes - Incident-driven: If issue discovered post-deployment

5.6.5 Example Transparency Pack Structure

TRANSPARENCY PACK: Synthetic AML Transaction Data

Dataset ID: SYNTH-AML-2025-Q4

Version: 1.0

Date: 2025-11-15

EXECUTIVE SUMMARY

This synthetic dataset was generated from retail banking transaction data for training anti-money laundering (AML) detection models. SDCF Gold Tier assessment found the dataset PROVISIONALLY APPROVED (SDCF-P) for model training with restrictions. Privacy risk slightly above target but acceptable with access controls. Fidelity and fairness meet requirements.

ASSESSMENT CERTIFICATE

Conformance: SDCF-P (Provisional)

Tier: Gold

Scores: PRS=28, FI=82, FV=12

Assessor: Jane Smith, Senior Data Scientist

Date: 2025-11-15

Valid Until: 2026-05-15

CONDITIONS:

- Use restricted to model development environment only
- Do not use for production deployment without reassessment
- Access limited to data science team (8 authorised users)
- Enhanced audit logging required
- Quarterly monitoring reports

TECHNICAL RESULTS

[Detailed tables and charts]

KNOWN LIMITATIONS

- Privacy risk marginally above target (PRS=28 vs. target 25)
- Rare transaction patterns may not be fully represented
- Cross-border transactions undersampled (5% of dataset vs. 8% in source)

USAGE RESTRICTIONS

APPROVED FOR: AML model training and testing in development environment

NOT APPROVED FOR: Production deployment, regulatory reporting, external sharing

CONTROLS REQUIRED: Access controls, audit logging, quarterly review

REGULATORY MAPPING

[Tables showing GDPR, AI Act compliance]

CONTACT

Questions: data-governance@example.com

Issues: security-incidents@example.com

5.7 C7: Release Rules

5.7.1 What It Is

Release Rules define the conditions under which synthetic data can be accessed, used, and shared. They operationalise the governance decisions and assessment findings into enforceable controls.

5.7.2 Why It Matters

Assessment without enforcement is meaningless. Release Rules ensure synthetic data is used only for approved purposes with appropriate safeguards. They provide accountability mechanism and prevent scope creep.

5.7.3 Required Elements

1. Access Control - Who can access synthetic data? - What authentication is required? - What is the approval process for new users? - How is access logged and monitored?

Example: - Access: Data science team only (8 named individuals) - Authentication: MFA required, role-based access control - New access: Manager approval + mandatory SDCF training - Logging: All data access logged, reviewed quarterly

2. Usage Restrictions - What can data be used for? - What is explicitly prohibited? - Can data be copied or exported? - Can results be published or shared?

Example: - Permitted: Model training, testing, validation in dev environment - Prohibited: Production deployment, external sharing, regulatory reporting - Export: Not permitted outside secure environment - Results: Model performance metrics can be shared; raw data cannot

3. Data Handling Requirements - Where can data be stored? (on-premises, cloud, edge) - What encryption is required? (at rest, in transit) - What backups are needed? - What is retention period?

Example: - Storage: Secure on-premises environment only (no cloud) - Encryption: AES-256 at rest, TLS 1.3 in transit - Backup: Weekly encrypted backups, 30-day retention - Retention: Delete after model development complete or 12 months, whichever earlier

4. Monitoring and Audit - What activities are logged? - Who reviews logs and when? - What triggers escalation? - How are violations handled?

Example: - Logging: All access, queries, exports, modifications - Review: Security team monthly review, governance board quarterly - Triggers: Unauthorised access attempt, bulk export, access outside business hours - Violations: Immediate access revocation, incident investigation, disciplinary action

5. Reassessment Triggers - When must dataset be reassessed? - What changes invalidate assessment? - Who authorises continued use?

Example: - Periodic: Annual reassessment - Change-driven: If purpose evolves, regulations change, or security incident occurs - Authority: CRO approval required for continued use after reassessment

6. Decommissioning - When must data be deleted? - How to ensure complete removal? - What records must be retained?

Example: - Deletion: Upon project completion or assessment expiry - Method: Secure deletion (NIST 800-88 guidelines) - Records: Governance Record and Transparency Pack retained for 7 years (regulatory requirement)

5.7.4 Implementation Guidance

Enforcement Mechanisms: - Technical: Access controls, encryption, DLP tools - Administrative: Policies, training, approvals - Detective: Logging, monitoring, auditing - Corrective: Incident response, violations handling

Integration with Existing Controls: - Data classification scheme (align synthetic data with appropriate tier) - Access management system (integrate with IAM) - SIEM/logging infrastructure (feed logs to central monitoring) - Policy management (incorporate into acceptable use policies)

User Communication: - All users acknowledge Release Rules before access - Periodic reminders (quarterly) - Updated training when rules change

Flexibility vs. Control: - Too restrictive: Data isn't used (defeats purpose) - Too permissive: Risk materialises (defeats privacy protection) - Balance: Proportionate to conformance level (SDCF-A less restrictive than SDCF-P)

5.7.5 Example Release Rules

RELEASE RULES: Synthetic AML Transaction Data

Dataset ID: SYNTH-AML-2025-Q4

Conformance: SDCF-P (Provisional)

Effective: 2025-11-15 to 2026-05-15

ACCESS CONTROL:

- Authorised Users: Data Science Team (8 individuals, named)
- Authentication: SSO + MFA required
- New Access: Manager approval + SDCF training mandatory
- Logging: All access logged in SIEM

USAGE RESTRICTIONS:

- PERMITTED: Model training/testing in dev environment
- PROHIBITED: Production use, external sharing, regulatory reports
- EXPORT: Not permitted

- RESULTS: Aggregate metrics only (no raw data)

DATA HANDLING:

- Storage: Secure on-prem only
- Encryption: AES-256 (rest), TLS 1.3 (transit)
- Retention: Delete after 12 months or project end
- Backup: Weekly encrypted, 30-day retention

MONITORING:

- Review: Monthly security review, quarterly governance review
- Triggers: Unauthorised access, bulk export, after-hours access
- Violations: Immediate revocation, incident investigation

REASSESSMENT:

- Scheduled: 2026-05-15 (6 months)
- Trigger Events: Purpose change, regulation change, incident
- Authority: CRO approval required

DECOMMISSIONING:

- Deletion: Secure wipe per NIST 800-88
- Records: Governance Record + Transparency Pack retained 7 years

ACKNOWLEDGMENT:

All users must acknowledge these rules before access.
Violations subject to disciplinary action up to termination.

End of Section 4

Continue to Section 5: Assessment Process to understand the step-by-step workflow for conducting SDCF assessments.

6 Assessment Process

This section describes the practical workflow for conducting SDCF assessments from initial scoping through certificate issuance.

6.1 Five-Step Workflow

6.1.1 Overview

SDCF assessment follows a structured five-step process:

1. **Define Purpose** → Document intended use and requirements (C1)
2. **Select Tier** → Determine appropriate assessment level (Gold/Silver/Bronze)
3. **Execute Tests** → Conduct privacy, fidelity, fairness testing (C3-C5)
4. **Evaluate Results** → Score metrics, apply thresholds, determine conformance
5. **Issue Certificate** → Publish Transparency Pack and Release Rules (C6-C7)

Each step has defined inputs, activities, outputs, and decision gates.

6.1.2 Step 1: Define Purpose

Objective: Establish clear, documented purpose for synthetic data use before assessment begins.

Activities: 1. **Stakeholder Engagement** - Interview business sponsor (what problem are we solving?) - Consult DPO (what GDPR obligations apply?) - Engage legal counsel (what regulatory requirements?) - Involve data science (what technical constraints?)

2. Purpose Sheet Development

- Complete C1 Purpose Sheet template (see Section 4.1)
- Document use case, regulatory context, disclosure scope
- Define pillar priorities and target thresholds (PRS, FI, FV)
- Identify failure modes and success criteria

3. Governance Review

- Present Purpose Sheet to governance board
- Address questions and concerns
- Obtain approval to proceed
- Document decision in C2 Governance Record

Inputs: - Business requirements - Regulatory obligations - Existing governance policies

Outputs: - Approved C1 Purpose Sheet - Initial C2 Governance Record entry

Decision Gate: - **PASS:** Purpose is clear, feasible, and compliant → Proceed to Step 2 - **REVISE:** Purpose needs refinement → Iterate with stakeholders - **REJECT:** Purpose is infeasible or non-compliant → Do not proceed

Typical Duration: 1-2 weeks (depends on stakeholder availability and governance cadence)

6.1.3 Step 2: Select Tier

Objective: Determine which assessment tier (Gold, Silver, Bronze) is appropriate given source data availability and purpose requirements.

Activities: 1. **Source Data Availability Assessment** - Is complete source data available and accessible? → Gold Tier possible - Are aggregate statistics or samples available? → Silver Tier possible - Is synthetic data only available? → Bronze Tier required

2. Purpose-Tier Alignment Check

- Does purpose require highest assurance? (high-risk AI, regulatory reporting) → Gold Tier needed
- Can purpose tolerate moderate uncertainty? (internal analytics, development) → Silver acceptable
- Is purpose appropriate for Bronze limitations? (testing, screening, AI training pre-check) → Bronze acceptable

3. Resource Assessment

- What is available budget?
- What is timeline constraint?
- What technical expertise is available?
- What tooling is in place?

4. Tier Selection Decision

- Document tier selection and rationale
- If tier doesn't align with purpose, identify gap and mitigation
- Update C2 Governance Record with decision

Decision Matrix:

Source Data Available	Purpose Risk Level	Resources	Recommended Tier
Full access	High (regulatory, high-risk AI)	Adequate	Gold (required)
Full access	Medium (internal analytics)	Adequate	Gold (preferred) or Silver
Full access	Low (testing, development)	Limited	Silver or Bronze
Partial (aggregates)	High	Adequate	Silver (with limitations) or obtain full access
Partial (aggregates)	Medium/Low	Adequate	Silver
None	High	Any	Bronze insufficient; obtain source OR reconsider approach
None	Medium	Adequate	Bronze (with strong controls)
None	Low	Any	Bronze

Inputs: - Approved C1 Purpose Sheet - Source data availability assessment - Resource constraints

Outputs: - Tier selection decision - Updated C2 Governance Record - Gap analysis (if tier-purpose mismatch)

Decision Gate: - **PROCEED:** Tier selected and appropriate → Move to Step 3 - **ESCALATE:** Tier-purpose mismatch requires executive decision → Governance board review - **BLOCK:** No viable tier for purpose → Reconsider synthetic data approach

Typical Duration: 3-5 days

6.1.4 Step 3: Execute Tests

Objective: Conduct technical privacy, fidelity, and fairness testing according to selected tier methodology.

Activities:

For Gold Tier:

1. Privacy Risk Testing (C3)

- Set up secure assessment environment (source data access required)
- Execute membership inference attacks (1000+ attack simulations)
- Compute distance to closest record (DCR) for all synthetic records
- Test attribute disclosure for sensitive attributes

- Calculate component scores and composite PRS
2. **Fidelity Testing (C4)**
 - Compare univariate distributions (all variables, statistical tests)
 - Compute correlation matrices and compare (source vs. synthetic)
 - Train benchmark models on source, evaluate on holdout
 - Train same models on synthetic, evaluate on same holdout
 - Calculate component scores and composite FI
 3. **Fairness Assessment (C5)**
 - Identify protected attributes in source and synthetic
 - Compute representation statistics per group
 - Train models on synthetic, evaluate separately per group
 - Test for predictive parity violations
 - Calculate component scores and composite FV

For Silver Tier: - Adapt methods to work with aggregate statistics or samples - Document additional uncertainty in scoring - Apply confidence intervals to account for incomplete information

For Bronze Tier: - Execute synthetic-only methodology (see Appendix C) - Apply conservative risk classification - Focus on red flag detection and domain validation - Explicitly document reduced confidence level

Computational Environment: - Secure processing environment (appropriate for data sensitivity) - Required tools: Python/R, SDMetrics or equivalent, statistical libraries - Hardware: CPU sufficient for Bronze; GPU useful for Gold Tier (large datasets)

Quality Assurance: - Peer review of methodology implementation - Validation of calculations (spot checks, reproducibility) - Documentation of any deviations from standard methodology

Inputs: - Source data (Gold/Silver) or synthetic data only (Bronze) - Selected tier methodology - Assessment tools and infrastructure

Outputs: - Privacy Risk Score (PRS) with component breakdown - Fidelity Index (FI) with component breakdown - Fairness Variance (FV) with component breakdown - Confidence intervals for each metric - Supporting data: charts, tables, statistical test results

Decision Gate: - **COMPLETE:** All tests executed successfully → Proceed to Step 4 - **TECHNICAL ISSUE:** Problems with data quality, tooling, or methodology → Resolve and retry - **SCOPE CHANGE:** Testing reveals need for different tier or purpose → Escalate to governance

Typical Duration: - Gold Tier: 2-3 weeks (comprehensive testing) - Silver Tier: 1-2 weeks (moderate testing) - Bronze Tier: 3-5 days (focused testing)

6.1.5 Step 4: Evaluate Results

Objective: Interpret test results against purpose-specific thresholds and determine conformance level.

Activities:

1. **Threshold Comparison**
 - Compare PRS to target from C1 Purpose Sheet
 - Compare FI to target from C1 Purpose Sheet
 - Compare FV to target from C1 Purpose Sheet

- Document which pillars meet/miss targets and by how much
2. **Trade-Off Analysis**
 - Are any misses compensable? (slightly above target but acceptable with controls)
 - What is the privacy-fidelity-fairness balance?
 - Can additional controls mitigate shortfalls?
 - What limitations must be documented?
 3. **Conformance Determination**
 - Apply conformance logic (see Section 3.4):
 - All targets met → SDCF-A candidate
 - One pillar slightly below, compensable → SDCF-P candidate
 - Significant shortfall → SDCF-R
 - Consider tier-purpose alignment in conformance decision
 - Document rationale for conformance level
 4. **Risk and Limitation Documentation**
 - What are the residual risks?
 - What are known limitations?
 - What controls are required?
 - What alternative uses are appropriate?
 5. **Governance Approval**
 - Present results to governance board
 - Discuss trade-offs and conformance recommendation
 - Obtain approval for conformance level
 - Document decision and any conditions in C2 Governance Record

Decision Logic Example:

```
IF (PRS  target_PRS AND FI  target_FI AND FV  target_FV)
  AND (tier appropriate for purpose)
THEN conformance = SDCF-A
```

```
ELSE IF (one pillar slightly misses target AND compensable with controls)
  AND (tier appropriate for purpose)
THEN conformance = SDCF-P
  DOCUMENT limitations and required controls
```

```
ELSE IF (significant shortfall OR tier inappropriate for purpose)
THEN conformance = SDCF-R
  DOCUMENT why unsuitable and alternative uses (if any)
```

Inputs: - Test results (PRS, FI, FV scores) - Purpose Sheet targets and priorities - Governance board review

Outputs: - Conformance level determination (SDCF-A/P/R) - Documented rationale - List of required controls (if SDCF-P) - Updated C2 Governance Record

Decision Gate: - **APPROVE:** Conformance level agreed → Proceed to Step 5 - **REVISE:** Governance requires different interpretation or controls → Update and re-present - **REJECT:** Dataset unsuitable, regeneration required → Return to data generation

Typical Duration: 1 week (including governance meeting schedule)

6.1.6 Step 5: Issue Certificate

Objective: Formalise assessment results in certificate and transparency pack, establish release rules.

Activities:

1. Certificate Creation

- Generate formal SDCF certificate with:
 - Dataset ID and version
 - Purpose statement
 - Tier (Gold/Silver/Bronze)
 - Conformance level (SDCF-A/P/R)
 - Scores (PRS, FI, FV)
 - Assessor and date
 - Validity period
 - Conditions (if SDCF-P) or restrictions (if SDCF-R)
- Obtain assessor signature and governance approval
- Assign version control and audit trail

2. Transparency Pack Assembly (C6)

- Compile all components (see Section 4.6):
 - Executive summary
 - Assessment certificate
 - Technical results (detailed)
 - Known limitations
 - Usage restrictions
 - Regulatory mapping
 - Contact information
- Create audience-specific versions (executive, technical, legal, user)
- Publish in accessible formats (PDF, web, metadata)

3. Release Rules Definition (C7)

- Document access controls (who, authentication, approval)
- Define usage restrictions (permitted, prohibited, exports)
- Specify data handling requirements (storage, encryption, retention)
- Establish monitoring and audit procedures
- Set reassessment triggers
- Plan decommissioning process
- Obtain governance approval for release rules

4. Communication and Training

- Brief stakeholders on results and conformance level
- Train authorised users on release rules
- Distribute transparency pack to appropriate audiences
- Update data inventory and governance systems
- Archive assessment documentation

5. Ongoing Monitoring Setup

- Configure logging and monitoring per release rules
- Schedule periodic reviews (quarterly, annually)
- Establish incident response procedure
- Set calendar reminders for reassessment

Inputs: - Approved conformance determination - All test results and supporting documentation - Governance decisions and conditions

Outputs: - Formal SDCF certificate - Complete Transparency Pack (C6) - Release Rules (C7) - Updated data governance systems - Trained users

Decision Gate: - **RELEASE:** All documentation complete, users trained → Data available for use - **HOLD:** Issues identified during final review → Resolve before release

Typical Duration: 1 week

6.2 Total Assessment Timeline

Gold Tier End-to-End: - Step 1 (Purpose): 1-2 weeks - Step 2 (Tier Selection): 3-5 days - Step 3 (Testing): 2-3 weeks - Step 4 (Evaluation): 1 week - Step 5 (Certificate): 1 week - **Total: 6-8 weeks**

Silver Tier End-to-End: - Steps 1-2: Same - Step 3 (Testing): 1-2 weeks - Steps 4-5: Same - **Total: 5-7 weeks**

Bronze Tier End-to-End: - Steps 1-2: Same - Step 3 (Testing): 3-5 days - Steps 4-5: Same - **Total: 4-5 weeks**

Accelerated Timeline (Bronze, Streamlined Governance): - Possible to complete in 2-3 weeks with pre-approved purpose and rapid governance cycles - Appropriate for lower-risk use cases (testing, development, preliminary screening)

6.3 Roles and Responsibilities

Assessment Team:

Assessment Lead: - Overall responsibility for assessment quality - Coordinates activities across steps - Presents to governance board - Signs certificate

Technical Assessor(s): - Executes privacy, fidelity, fairness tests - Implements methodology correctly - Documents results and supporting evidence - Requires: Data science expertise, statistical knowledge, SDCF training

Domain Expert: - Interprets results in business/research context - Validates that fidelity and fairness testing aligns with domain needs - Advises on appropriate thresholds and trade-offs - Requires: Subject matter expertise in relevant domain (healthcare, finance, etc.)

Governance Representatives: - DPO: GDPR compliance and privacy risk evaluation - Legal: Regulatory compliance and contractual implications - CISO: Security controls and release rules - Business Sponsor: Purpose definition and business value - Risk Manager: Enterprise risk management integration

Quality Reviewer: - Independent peer review of methodology and calculations - Validates reproducibility and correctness - Identifies errors or deviations from standard - Requires: Senior technical expertise, SDCF trained

6.4 Documentation and Audit Trail

Required Documentation at Each Step:

Step 1 (Purpose): - Completed C1 Purpose Sheet (version controlled) - Stakeholder meeting notes - Governance approval record

Step 2 (Tier Selection): - Source data availability assessment - Tier selection rationale - Gap analysis (if applicable)

Step 3 (Testing): - Test execution logs - Raw results (scores, statistical tests, model outputs) - Configuration files and code (for reproducibility) - Any deviations from standard methodology

Step 4 (Evaluation): - Threshold comparison tables - Trade-off analysis narrative - Conformance determination rationale - Governance meeting minutes

Step 5 (Certificate): - Final certificate (signed) - Complete Transparency Pack - Release Rules - User acknowledgment records

Retention Requirements: - Assessment documentation: Retain for validity period + 7 years (regulatory requirement) - Governance Records (C2): Permanent retention (organisational learning) - Transparency Packs: Retain as long as data exists + 7 years

Audit Readiness: - All documentation indexed and retrievable - Chain of custody documented (who did what, when) - Version control for all artifacts - Ready to present to internal audit, external auditors, or regulators

6.5 Common Challenges and Mitigations

Challenge 1: Purpose Definition Too Vague - *Symptom:* “Improve analytics” without specificity - *Impact:* Cannot set appropriate thresholds or determine conformance - *Mitigation:* Use C1 template rigorously; governance board pushes back on vague purposes

Challenge 2: Tier-Purpose Mismatch - *Symptom:* Bronze Tier for high-risk AI system - *Impact:* Inadequate assurance for purpose - *Mitigation:* Step 2 decision gate blocks mismatch; executive escalation required

Challenge 3: Testing Environment Constraints - *Symptom:* Cannot access source data securely for Gold Tier - *Impact:* Must drop to Silver/Bronze, reducing confidence - *Mitigation:* Plan assessment environment early; secure data rooms, trusted execution

Challenge 4: Threshold Negotiation - *Symptom:* Business wants to proceed despite failing thresholds - *Impact:* Risk of non-compliant use - *Mitigation:* Governance board enforces standards; SDCF-R is valid outcome

Challenge 5: Resource Constraints - *Symptom:* Budget or timeline pressure to shortcut assessment - *Impact:* Poor quality assessment, risk exposure - *Mitigation:* Bronze Tier is legitimate choice for appropriate purposes; don’t claim Gold when doing Bronze

Challenge 6: Expertise Gaps - *Symptom:* Team lacks privacy attack or fairness assessment knowledge - *Impact:* Tests implemented incorrectly - *Mitigation:* Training requirement for assessors; quality review by independent expert; use established tools (SDMetrics)

End of Section 5

Continue to Section 6: Relationship to Existing Tools and Standards to understand how SDCF integrates with the broader ecosystem.

7 Relationship to Existing Tools and Standards

SDCF is designed to complement, not replace, existing tools and standards. This section clarifies how SDCF fits into the broader synthetic data quality and governance ecosystem.

7.1 Complementary Open-Source Tools

7.1.1 SDMetrics / SDV (Synthetic Data Vault)

What It Is: Open-source Python library from MIT for evaluating synthetic data quality. Provides statistical metrics for fidelity and privacy.

Capabilities: - Quality metrics: Distribution similarity, correlation preservation, column shapes - Privacy metrics: Membership inference (basic), distance to closest record (DCR), categorical coverage - Detection metrics: Identify synthetic records that may expose privacy risks - Reporting: Generate HTML quality reports with visualizations

How SDCF Uses SDMetrics: - **Fidelity Index (FI) inputs:** SDMetrics quality metrics feed directly into FI calculation - **Privacy Risk Score (PRS) inputs:** DCR and membership inference results contribute to PRS - **Bronze Tier:** SDMetrics can compute synthetic-only metrics for B-FI

What SDCF Adds: - **Purpose-bounded interpretation:** SDMetrics reports metrics; SDCF interprets them for specific use case - **Regulatory mapping:** SDCF connects metrics to GDPR/AI Act requirements - **Conformance determination:** SDCF provides SDCF-A/P/R decision framework - **Governance integration:** C1-C7 control sets provide organisational context - **Tiered methodology:** SDCF handles Gold/Silver/Bronze scenarios; SDMetrics assumes Gold

Integration Pattern:

```
# Example workflow
from sdmetrics.reports.single_table import QualityReport
from sdcf import FidelityIndex, PrivacyRiskScore

# Step 1: Generate SDMetrics report
report = QualityReport()
report.generate(real_data, synthetic_data, metadata)

# Step 2: Extract metrics for SDCF
sdmetrics_scores = report.get_properties()

# Step 3: Calculate SDCF scores
fi = FidelityIndex.from_sdmetrics(sdmetrics_scores, purpose_weights)
prs = PrivacyRiskScore.from_sdmetrics(sdmetrics_scores, adversary_model)

# Step 4: Determine conformance
conformance = sdcf_evaluate(fi, prs, fv, purpose_sheet)
```

See Appendix F for complete reference implementation.

7.1.2 mostlyai-qa (MOSTLY AI Quality Assurance)

What It Is: Open-source quality assurance toolkit released by MOSTLY AI (April 2025). Focuses on fidelity and novelty metrics for synthetic tabular data.

Capabilities: - Fidelity assessment: Distribution similarity, correlation preservation - Novelty assessment: Measures how much synthetic data differs from source (privacy proxy) - Comprehensive reports: Statistical tests, visualizations, recommendations - Automated analysis: Detects common synthetic data quality issues

How SDCF Uses mostlyai-qa: - **Fidelity Index (FI):** mostlyai-qa fidelity scores map to FI components - **Privacy Risk Score (PRS):** Novelty metrics provide input to PRS (high novelty = lower privacy risk) - **Silver/Bronze Tier:** mostlyai-qa designed for scenarios with limited source access

What SDCF Adds: - **Fairness dimension:** mostlyai-qa doesn't assess bias; SDCF adds FV - **Regulatory compliance:** SDCF maps results to specific regulations - **Purpose-bounded thresholds:** mostlyai-qa provides metrics; SDCF sets context-specific targets - **Certificate framework:** SDCF provides formal conformance determination

Integration Pattern: Organisations can use mostlyai-qa for initial technical assessment, then apply SDCF framework for compliance interpretation and governance.

7.1.3 Other Relevant Tools

DataSynthesizer (University of Washington): - Generates differentially private synthetic data - Can provide formal privacy guarantees (, parameters) - SDCF complements by assessing utility (fidelity) and fairness, which differential privacy doesn't guarantee

Synthetic Data Gym (OpenMined): - Benchmarking framework for synthetic data methods - SDCF can consume benchmark results for comparative analysis

AnonML / ARX Data Anonymisation Tool: - Focuses on anonymisation techniques including synthetic data - Provides k-anonymity, l-diversity, t-closeness metrics - SDCF provides additional privacy risk assessment beyond formal privacy models

7.2 Vendor Platform Integration

Synthetic data generation platforms provide proprietary quality assessment. SDCF serves as independent validation layer.

7.2.1 Gretel.ai

Platform Capabilities: - Synthetic data generation (tabular, time-series, text) - Built-in evaluation metrics (SQS - Synthetic Quality Score) - Model training and deployment

SDCF Relationship: - **Independent validation:** SDCF provides vendor-neutral assessment of Gretel outputs - **Procurement evaluation:** Use SDCF Bronze Tier to evaluate Gretel datasets before purchase - **Compliance evidence:** SDCF certificate supplements Gretel quality reports for regulatory purposes

Integration: Gretel API can export synthetic data → SDCF assessment → Certificate attached to dataset metadata

7.2.2 MOSTLY AI

Platform Capabilities: - Enterprise synthetic data generation - QA reports (now open-sourced as mostlyai-qa) - Accuracy and privacy metrics

SDCF Relationship: - **Complementary assessment:** Use mostlyai-qa for technical metrics, SDCF for compliance interpretation - **Governance layer:** SDCF C1-C7 controls provide governance around MOSTLY AI use - **Certificate for external sharing:** When sharing MOSTLY AI-generated data with partners, SDCF certificate provides independent validation

7.2.3 Syntho

Platform Capabilities: - Synthetic data generation for healthcare and finance - Quality assurance reports - External evaluation partnerships (claims SAS Institute validation)

SDCF Relationship: - **Independent verification:** SDCF provides methodology for organisations to independently verify Syntho claims - **Risk management:** SDCF assessment informs risk decisions about Syntho adoption - **Bronze Tier use case:** Assess Syntho outputs without accessing source data (common procurement scenario)

7.2.4 General Vendor Integration Pattern

Procurement Phase: - Use SDCF Bronze Tier to evaluate vendor demo datasets - Require vendors to provide SDCF-compliant documentation - Include SDCF conformance requirements in RFPs

Operational Phase: - Vendor generates synthetic data → Organisation conducts SDCF assessment → Certificate issued - SDCF governance (C1-C7) wraps around vendor tooling - SDCF certificate provides audit trail independent of vendor

Value Proposition: - Vendors: Can support SDCF as differentiation (enables compliant use of their outputs) - Organisations: Reduces vendor lock-in (standardised assessment regardless of generator)

7.3 Standards Alignment

SDCF aligns with and can support compliance with established standards:

7.3.1 ISO/IEC 27001:2022 (Information Security Management)

Relevant Controls: - **A.8.2 Information Classification:** SDCF C1/C7 support classification of synthetic data assets - **A.8.3 Media Handling:** SDCF C7 Release Rules define handling requirements - **A.8.10 Information Deletion:** SDCF C7 specifies retention and deletion - **A.8.11 Data Masking:** Synthetic data as masking technique; SDCF validates effectiveness - **A.5.23 Information Security for Cloud Services:** SDCF assessment applies to cloud-stored synthetic data

How SDCF Supports ISO 27001: - Control implementation evidence (SDCF Control Sets C1-C7) - Risk assessment inputs (PRS quantifies information security risk) - Audit trail (C2 Governance Record, C6 Transparency Pack)

7.3.2 ISO/IEC 27701:2019 (Privacy Information Management)

Relevant Controls: - **6.7.2.2 Identify basis for PII transfer:** SDCF assessment supports demonstrating appropriate safeguards - **7.2.2 PII de-identification and deletion:** SDCF validates de-identification effectiveness - **7.4.7 Automated decision-making:** SDCF fairness assessment for AI training data

How SDCF Supports ISO 27701: - Privacy risk quantification (PRS for privacy impact assessment) - Demonstrating appropriate technical measures (SDCF certificate as evidence) - Accountability documentation (C2 Governance Record)

7.3.3 ISO/IEC 23894:2023 (AI Risk Management)

Relevant Clauses: - **Data quality:** SDCF Fidelity Index operationalises data quality assessment - **Bias management:** SDCF Fairness Variance quantifies bias risk - **Transparency:** SDCF Transparency Pack provides model training data documentation

How SDCF Supports ISO 23894: - Training data governance (C1-C7 for synthetic training datasets) - Bias testing (FV for Article 10(4) compliance) - Risk documentation (C2 Governance Record, C6 Transparency Pack)

7.3.4 ISO/IEC 42001:2023 (AI Management System)

Relevant Requirements: - **6.1 Risk Management:** SDCF assessment informs AI system risk analysis - **7.2 Data Management:** SDCF provides data quality management framework for synthetic data - **8.2 AI System Development:** SDCF documents training data governance

How SDCF Supports ISO 42001: - Data governance controls for AI systems (C1-C7) - Risk assessment for synthetic training data (PRS, FI, FV) - Audit documentation for certification (SDCF certificates and Transparency Packs)

See Appendix D for detailed control mapping tables.

7.3.5 NIST Privacy Framework (2020)

Core Functions: - **Identify-P:** C1 Purpose Sheet identifies privacy risks in synthetic data use - **Govern-P:** C2 Governance Record documents accountability and oversight - **Control-P:** C3-C5 testing implements technical measures - **Communicate-P:** C6 Transparency Pack provides stakeholder communication - **Protect-P:** C7 Release Rules establish protective measures

How SDCF Supports NIST Privacy Framework: - Operationalises functions with specific procedures - Quantifies privacy risk (PRS) - Provides evidence for privacy risk management

7.3.6 NIST AI Risk Management Framework (2023)

Trustworthy AI Characteristics: - **Privacy:** SDCF PRS directly measures privacy risk - **Fairness:** SDCF FV quantifies bias and representation issues - **Explainability:** SDCF Transparency Pack documents data characteristics - **Safety:** SDCF Fidelity Index ensures training data quality

How SDCF Supports NIST AI RMF: - Training data risk assessment (Maps to “Map” function) - Documented governance (Maps to “Govern” function) - Ongoing monitoring (Maps to “Manage” function)

7.4 What SDCF Adds to the Ecosystem

7.4.1 The Gap SDCF Fills

Existing tools provide metrics: - SDMetrics: “Distribution similarity = 0.87” - mostlyai-qa: “Fidelity score = 82%” - Vendors: “High quality synthetic data”

But don’t answer: - Is 0.87 similarity good enough for MY purpose? - Does 82% fidelity meet GDPR Article 32 “appropriate security”? - What conformance level should I assign? - How do I document this for regulators?

SDCF bridges this gap: - Purpose-bounded interpretation (fitness for specific use case) - Regulatory mapping (connects metrics to legal requirements) - Governance framework (C1-C7 controls for organisational accountability) - Conformance determination (SDCF-A/P/R decision logic) - Evidence packaging (Transparency Pack for auditors/regulators)

7.4.2 SDCF Value Proposition Summary

For Practitioners: - Actionable guidance (not just metrics) - Honest about limitations (especially Bronze Tier) - Flexible methodology (works with available tools and data) - Defensible decisions (audit trail, governance integration)

For Organisations: - Risk management (quantified privacy/fidelity/fairness risks) - Compliance demonstration (GDPR, AI Act, ISO evidence) - Vendor evaluation (standardised assessment across generators) - Accountability (documented decision-making)

For Regulators: - Transparency (clear methodology and limitations) - Standardization (common assessment approach across organisations) - Risk-based (proportionate to purpose and sensitivity) - Auditable (complete documentation trail)

For the Ecosystem: - Open methodology (transparent, improvable, not black box) - Tool-agnostic (works with any generator or assessment tool) - Standards-aligned (complements ISO, NIST, not competing) - Fills real gap (interpretation layer that’s currently missing)

7.4.3 How to Use SDCF in Practice

Scenario 1: Building In-House Synthetic Data 1. Define purpose (C1) before generating 2. Generate synthetic data using preferred method/tool 3. Conduct Gold Tier SDCF assessment 4. Iterate generation methodology if SDCF-R 5. Deploy with Release Rules (C7) if SDCF-A or SDCF-P

Scenario 2: Procuring Vendor Synthetic Data 1. Define purpose (C1) for procurement 2. Request Bronze Tier demo datasets from vendors 3. Conduct SDCF Bronze assessments for comparison 4. Select vendor 5. Conduct Silver/Gold Tier assessment of production dataset 6. Deploy with Release Rules if conformance acceptable

Scenario 3: Validating Third-Party Datasets 1. Receive synthetic dataset (no source access) 2. Define your intended purpose (C1) 3. Conduct Bronze Tier assessment 4. Determine if suitable for your purpose (may differ from vendor’s purpose) 5. Deploy with appropriate controls or decline use

Scenario 4: Ongoing Governance 1. Initial assessment establishes baseline 2. Periodic reassessment (annual or per release rules) 3. Monitoring for changes that trigger reassessment 4. Governance

Record captures evolution over time 5. Transparency Pack updated at each reassessment

End of Section 6

8 Empirical Validation Study: Bronze Tier Retrospective Evaluation

This section presents preliminary empirical validation of the Bronze Tier methodology through assessment of 10 diverse synthetic datasets. Initial evidence supports the framework’s design principles and demonstrates cross-domain applicability. **Limitation:** This validation is underpowered for statistical generalisation ($n=10$); findings should be interpreted as preliminary evidence supporting framework design. Statistical expansion to $n>50$ datasets with adversarial validation (membership inference benchmarking, linkage attack modeling) and ROC-based threshold calibration is planned for v2.0.

8.1 Study Design and Motivation

8.1.1 Validation Objective

Bronze Tier represents the most challenging and practically relevant assessment scenario: evaluating synthetic data when source data is unavailable due to privacy constraints, commercial confidentiality, or third-party data acquisition. This validation study addresses three research questions:

RQ1: Does Bronze Tier methodology discriminate dataset quality? Can Bronze Tier metrics (B-PRS, B-FI, B-FV) effectively distinguish between datasets of varying quality and risk profiles without access to source data?

RQ2: Do results align with framework design principles? Does the methodology exhibit the conservative bias, cross-domain applicability, and appropriate conformance distribution predicted by framework specifications (Section 3.3, Appendix C)?

RQ3: What practical guidance emerges from empirical assessment? What actionable insights can practitioners derive regarding synthesis method selection, expected metric ranges, and appropriate use cases?

8.1.2 Why Validate Bronze Tier First

We prioritise Bronze Tier validation for strategic reasons:

1. **Practical relevance:** Most real-world synthetic data evaluation scenarios involve third-party datasets or legacy data where source access is unavailable or impractical
2. **Methodological challenge:** Bronze Tier is the most technically challenging tier (source-free assessment), making it the strongest test of framework robustness
3. **Conservative design verification:** Framework explicitly designs Bronze Tier to be conservative; validation can confirm this behaves as intended
4. **Data availability:** Publicly available synthetic datasets (Bronze Tier appropriate) are more accessible than paired synthetic-source datasets (Gold Tier)

Silver and Gold Tier validation remains future work (Section 7.7).

8.2 Dataset Portfolio

8.2.1 Selection Criteria

We selected datasets to maximize diversity across five dimensions:

- 1. Domain Coverage:** Datasets span seven domains (demographic, healthcare, e-commerce, AI training, AI safety, business, code) to test cross-domain applicability.
- 2. Synthesis Methods:** Five generation approaches represented (GaussianCopula statistical synthesis, CTGAN deep learning, TVAE variational autoencoder, commercial GANs, LLM-generated) to enable method comparison.
- 3. Size and Complexity:** Records range 377 to 69,659; features range 4 to 15 columns, testing scalability.
- 4. Data Characteristics:** Mix of demographic (categorical-heavy), transactional (temporal), and unstructured (text) data types.
- 5. Provenance:** Data from academic research (PLEIAs), open-source tools (SDV), commercial vendors (MostlyAI, Gretel), government agencies (US Census, CMS) to represent practitioner reality.

Inclusion Requirement: All datasets must be (a) publicly available, (b) true Bronze Tier scenarios (no source data access), and (c) documented with generation methodology.

8.2.2 Portfolio Summary

Table 4 presents the complete dataset portfolio.

Table 4: Bronze Tier Validation Dataset Portfolio Summary

ID	Dataset	Records	Features	Domain	Synthesis	Year
D1	PLEIAs SYNTH	10,000	14	AI Training	LLM	2025
D2	SDV Adult - GaussianCopula	32,561	15	Demographic	Statistical	2024
D3	SDV Adult - CTGAN	32,561	15	Demographic	GAN (Deep)	2024
D4	SDV Adult - TVAE	32,561	15	Demographic	VAE (Deep)	2024
D5	Gretel Safety Alignment	8,361	14	AI Safety	LLM	2024
D6	MostlyAI Census	48,842	15	Demographic	GAN (Comm.)	2023
D7	MostlyAI CDNOW Purchases	69,659	4	E-commerce	GAN (Comm.)	2023
D8	CMS DE-SynPUF Demo	5,000	14	Healthcare	Multi-method	2023
D9	US Census SynLBD Demo	10,000	12	Business	Synth-augm.	2024
D10	Jupyter Agent Dataset	377	11	Code/Data	LLM	2025
Total	10 datasets	219,360	4–15	7 domains	5 methods	2023–25

Portfolio Characteristics:

- **Domain distribution:** Demographic (40%), AI/Code (30%), Commercial/Government (30%)
- **Method distribution:** Deep learning (40%), Statistical/Commercial (30%), LLM (30%)
- **Temporal coverage:** All datasets from 2023–2025 (modern synthesis methods)
- **Total scope:** 219,360 records assessed, 129 total features

Methodological Note: Datasets D2–D4 are self-synthesized from the SDV Adult Income real dataset (UCI ML Repository) using three different methods. While source data exists for these variants, we treat them as Bronze Tier for this validation (source data deliberately not accessed during assessment). This enables controlled method comparison while maintaining Bronze Tier assessment conditions.

8.3 Methodology

8.3.1 Assessment Procedure

Each dataset underwent standardised Bronze Tier assessment following Appendix C methodology:

Step 1: Data Loading and Validation

- Load synthetic dataset (CSV format)
- Schema validation (column types, missing values, basic statistics)
- Data type identification (categorical vs. continuous)

Step 2: B-PRS Computation (Privacy Risk Score) Per Appendix C.2 (Bronze Tier Privacy Assessment):

- **Outlier Score:** Local Outlier Factor (LOF) analysis identifying records with unusual attribute combinations (proxy for uniqueness risk)
- **Uniqueness Score:** High-dimensional uniqueness assessment via duplicate detection and near-duplicate identification
- **Context Penalty:** +0.1 base penalty for source-free assessment (conservative bias)
- **B-PRS Formula:** $B\text{-}PRS = 0.4 \times \text{outlier} + 0.4 \times \text{uniqueness} + 0.2 \times \text{penalty}$

Step 3: B-FI Computation (Fidelity Index) Per Appendix C.3 (Bronze Tier Fidelity Assessment):

- **Consistency Score:** Schema validation, type checking, range validation, missing data assessment
- **Validity Score:** Statistical plausibility checks, domain constraint testing, duplicate detection
- **B-FI Formula:** $B\text{-}FI = 0.6 \times \text{consistency} + 0.4 \times \text{validity}$

Step 4: B-FV Computation (Fairness Variance) Per Appendix C.4 (Bronze Tier Fairness Assessment):

- **Representation Gap:** Distribution balance analysis across categorical variables
- **Uncertainty Buffer:** +0.1 penalty for no source baseline comparison
- **B-FV Formula:** $B\text{-}FV = \text{representation_gap} + 0.1$

Step 5: Conformance Determination Apply decision logic per Section 3.4:

- **SDCF-R-Bronze (Rejected):** $B\text{-}PRS > 0.70$ OR $B\text{-}FI < 0.60$ OR $B\text{-}FV > 0.30$
- **SDCF-A-Bronze (Acceptable):** $B\text{-}PRS < 0.50$ AND $B\text{-}FI > 0.70$ AND $B\text{-}FV < 0.30$
- **SDCF-P-Bronze (Provisional):** All other combinations

8.3.2 Implementation Details

Software: Assessment implemented in Python 3.11 using:

- pandas 2.1.0 (data manipulation)

- `numpy` 1.24.0 (numerical computation)
- `scikit-learn` 1.3.0 (LOF analysis, scaling)
- `scipy` 1.11.0 (statistical tests)

Platform: Windows 10, 16GB RAM, Intel i7 processor

Execution Time: Complete assessment of 10 datasets completed in 14.3 minutes (average 1.4 minutes per dataset)

Reproducibility: Complete implementation available in ancillary files (`bronze_retrospective.py`) with dataset access scripts and expected results for verification.

8.4 Results

8.4.1 Conformance Distribution

Table 5 presents the overall conformance distribution.

Table 5: Bronze Tier Validation: Conformance Distribution				
Conformance	Count	Percentage	Total Records	Avg B-PRS
SDCF-R-Bronze	6	60%	170,683	0.684
SDCF-A-Bronze	4	40%	48,677	0.300
SDCF-P-Bronze	0	0%	0	—
Total	10	100%	219,360	0.516

Key Finding: 60% of datasets received Restricted conformance (SDCF-R-Bronze), 40% Acceptable (SDCF-A-Bronze), and 0% Provisional. This distribution confirms the framework’s conservative design principle for Bronze Tier assessment (Section 3.3).

SDCF-R-Bronze Datasets (6):

- D1 (PLEIAs SYNTH): B-FV > 0.30 trigger (0.909 Problematic fairness)
- D2 (SDV-GaussianCopula): B-FV > 0.30 trigger (0.693 Problematic fairness)
- D3 (SDV-CTGAN): B-FV > 0.30 trigger (0.548 Problematic fairness)
- D4 (SDV-TVAE): B-FV > 0.30 trigger (0.724 Problematic fairness)
- D5 (Gretel Safety): B-PRS > 0.70 trigger (0.789 Critical privacy risk)
- D6 (MostlyAI Census): B-FV > 0.30 trigger (0.692 Problematic fairness)

SDCF-A-Bronze Datasets (4):

- D7 (MostlyAI CDNOW): Clean pass (B-PRS 0.360, B-FI 0.999, B-FV 0.100)
- D8 (CMS SynPUF): Clean pass (B-PRS 0.390, B-FI 0.997, B-FV 0.100)
- D9 (Census SynLBD): Clean pass (B-PRS 0.360, B-FI 1.000, B-FV 0.100)
- D10 (Jupyter Agent): Clean pass (B-PRS 0.090, B-FI 0.999, B-FV 0.128)

Pattern Analysis: SDCF-A-Bronze datasets share common characteristics: (1) simpler schemas (4–14 columns vs. 14–15 for Restricted), (2) lower categorical complexity, (3) targeted use cases (transactional, healthcare claims, business establishments, code examples) rather than general demographic data.

8.4.2 Privacy Risk Score (B-PRS) Performance

Table 6 presents B-PRS distribution across datasets.

Table 6: B-PRS Performance Summary					
ID	B-PRS	Risk Level	Domain	Method	Schema
D10	0.090	Low	Code/Data	LLM	11 cols
D7	0.360	Moderate	E-commerce	GAN	4 cols
D9	0.360	Moderate	Business	Synth	12 cols
D8	0.390	Moderate	Healthcare	Multi	14 cols
D4	0.534	High	Demographic	TVAE	15 cols
D6	0.631	High	Demographic	GAN	15 cols
D3	0.637	High	Demographic	CTGAN	15 cols
D2	0.639	High	Demographic	GaussianCop	15 cols
D1	0.777	Critical	AI Training	LLM	14 cols
D5	0.789	Critical	AI Safety	LLM	14 cols
Avg: 0.516		Mod-High			
Range: 0.09–0.79		8.8x			

Discrimination Test (RQ1): B-PRS achieves 8.8x discrimination range (0.090 to 0.789), significantly exceeding the framework’s minimum 2x requirement (Appendix A.1). This demonstrates effective quality differentiation without source data access.

Distribution Analysis:

- **Low** (< 0.30): 10% (1/10) — Framework: “Bronze rarely achieves Low” ✓
- **Moderate** (0.30–0.49): 30% (3/10) — Simple schemas, targeted domains
- **High** (0.50–0.69): 40% (4/10) — All demographic datasets
- **Critical** (≥ 0.70): 20% (2/10) — Unique AI training content

Domain Patterns:

- **Demographic data:** Consistently High risk (0.534–0.639), average 0.610 — Census attributes create many quasi-identifiers
- **Commercial/Business:** Moderate risk (0.360–0.390), average 0.370 — Simpler schemas, lower dimensionality
- **AI/Code:** Highly variable (0.090–0.789), average 0.552 — Depends on content uniqueness

Framework Alignment (RQ2): Results align with Appendix C.2 predictions: (1) conservative scoring (average Moderate-High), (2) demographic datasets flagged High risk due to quasi-identifiers, (3) simple schemas show lower risk, (4) rare for Bronze to achieve Low risk (10% observed).

8.4.3 Fidelity Index (B-FI) Performance

Table 7 presents B-FI distribution.

Key Finding: All 10 datasets achieve Excellent fidelity (B-FI > 0.90), with average 0.991. This confirms framework expectation that modern synthesis tools (2023+) produce high internal quality.

Perfect Fidelity (B-FI = 1.000): Achieved by US Census SynLBD Demo (D9) — Reflects clean schema, comprehensive validation, government data quality standards.

Table 7: B-FI Performance Summary

Metric	Value	Framework Prediction	Match?
Minimum B-FI	0.927	> 0.60 for modern tools	✓
Maximum B-FI	1.000	Achievable with clean data	✓
Average B-FI	0.991	High for 2023+ synthesis	✓
Std Dev	0.021	Low (consistent quality)	✓
Excellent (> 0.90)	100% (10/10)	Expected for modern tools	✓
Poor (< 0.60)	0% (0/10)	Rare in modern synthesis	✓

Lowest Fidelity (B-FI = 0.927): PLEIAs SYNTH (D1) — Still Excellent, but reflects academic research dataset (less polished than commercial tools). Minor issues: higher missing value rate (2.1% vs. < 0.1% for others), occasional type inconsistencies.

Interpretation (RQ3): B-FI is highly effective for detecting internal quality issues. In this portfolio, modern tools demonstrate excellent internal consistency. B-FI would be most valuable for:

- Legacy synthetic data (pre-2020 generation methods)
- Custom/prototype synthesis implementations
- Academic research datasets
- Data corruption detection

For practitioner guidance: B-FI > 0.90 is achievable and expected for modern commercial and open-source synthesis tools.

8.4.4 Fairness Variance (B-FV) Performance

Table 8 presents B-FV distribution.

Table 8: B-FV Performance Summary

ID	B-FV	Level	Pattern	Categorical Cols
D7, D8, D9	0.100	Fair	Baseline (min penalty)	1–4
D10	0.128	Fair	Minimal variance	5
D5	0.100	Fair	Balanced by design	8
D3	0.548	Problematic	Demographic complexity	9
D2	0.693	Problematic	High quasi-IDs	9
D6	0.692	Problematic	Census representation	9
D4	0.724	Problematic	Age/gender/race gaps	9
D1	0.909	Problematic	Diverse AI content	7
Pattern	Bimodal	0.10 or 0.55–0.91	Not continuous	—

Bimodal Distribution: B-FV exhibits bimodal pattern: datasets cluster at 0.10 (baseline Fair) or 0.55–0.91 (Problematic), with no datasets in Concerning range (0.15–0.30).

Pattern Analysis:

- **B-FV = 0.10 (baseline):** Simple schemas (≤ 4 categorical columns) OR intentionally balanced data (Gretel Safety) — Minimal representation gaps
- **B-FV = 0.55–0.91:** Complex demographic data (9+ categorical columns) with age, gender, race, education, occupation — High representation variance reflects real-world population complexity

Framework Alignment (RQ2): Appendix C.4 explicitly states: “B-FV without source data has high uncertainty. High scores may reflect real-world imbalances (not synthesis failures).” Observed bimodal distribution confirms this design characteristic.

Critical Interpretation Issue: Without source baseline comparison, B-FV 0.693 for demographic data cannot distinguish:

- **Scenario A:** Synthesis amplified representation gaps (concerning)
- **Scenario B:** Synthesis preserved real population imbalances (appropriate)

Practitioner Guidance (RQ3): Use B-FV as **screening flag**, not definitive measurement:

- B-FV < 0.15 : Simple schemas, likely low fairness concern
- B-FV > 0.30 : Triggers SDCF-R-Bronze; conduct expert review to determine if gaps are problematic or population-representative
- For fairness-critical decisions: Upgrade to Silver/Gold Tier for source comparison

8.4.5 Cross-Domain Patterns

Table 9 aggregates results by domain.

Domain	N	Avg B-PRS	Avg B-FI	Avg B-FV	Conform	Pattern
Demographic	4	0.610	0.997	0.664	100% R	High PRS, High FV
E-commerce	1	0.360	0.999	0.100	100% A	Low complexity
Healthcare	1	0.390	0.997	0.100	100% A	Moderate PRS
Business	1	0.360	1.000	0.100	100% A	Simple schema
AI Training	1	0.777	0.927	0.909	100% R	Unique content
AI Safety	1	0.789	0.999	0.100	100% R	Critical PRS
Code/Data	1	0.090	0.999	0.128	100% A	Small dataset
Overall	10	0.516	0.991	0.396	60% R	—

Domain-Specific Insights:

Demographic Data (Predictably High Risk):

- Consistently High B-PRS (0.534–0.639), average 0.610
- All datasets SDCF-R-Bronze (B-FV trigger)
- Pattern: Census attributes (age, gender, race, education, marital status, occupation) create many quasi-identifiers → inherently higher privacy risk in Bronze Tier assessment
- Recommendation: Demographic datasets should anticipate Restricted conformance unless schema is simplified

Commercial/Business Data (Lower Risk):

- Lower B-PRS (0.360–0.390), average 0.367

- All datasets SDCF-A-Bronze
- Pattern: Transactional schemas (customer ID, product, date, amount) have fewer quasi-identifiers, lower dimensionality → acceptable Bronze Tier risk profile
- Recommendation: Transactional and business establishment data well-suited for Bronze Tier Acceptable conformance

AI/Code Data (Highly Variable):

- Extreme B-PRS range (0.090–0.789)
- Mixed conformance (50% R, 50% A)
- Pattern: Depends on content uniqueness — Generic code examples (low risk) vs. unique training data (high risk)
- Recommendation: AI/Code datasets require case-by-case assessment; cannot generalise expected conformance

Framework Alignment (RQ2): Cross-domain patterns align with Section 3.3 prediction: “Bronze Tier should work across domains with consistent methodology.” Observed patterns are interpretable and domain-logical, confirming methodology is domain-agnostic while results are appropriately domain-sensitive.

8.4.6 Synthesis Method Comparison

Table 10 compares synthesis methods.

Table 10: Synthesis Method Performance Comparison						
Method	Dataset(s)	B-PRS	B-FI	B-FV	Conform	N
Statistical	SDV-GaussianCopula	0.639	0.999	0.693	R	1
CTGAN	SDV-CTGAN	0.637	0.998	0.548	R	1
TVAE	SDV-TVAE	0.534	0.994	0.724	R	1
GAN (Comm.)	MostlyAI (2x)	0.360–0.631	0.997–0.999	0.100–0.692	Mixed	2
LLM	3x datasets	0.090–0.789	0.927–0.999	0.100–0.909	Mixed	3
Multi/Synth	CMS, SynLBD	0.360–0.390	0.997–1.000	0.100	A	2

Key Finding (RQ3): TVAE Demonstrates Best Privacy-Quality Balance for Demographic Data

Controlled comparison of three methods on identical source data (SDV Adult):

- **TVAE: B-PRS 0.534** (High, but **lowest** among deep learning)
- CTGAN: B-PRS 0.637 (High, +19% vs. TVAE)
- GaussianCopula: B-PRS 0.639 (High, +20% vs. TVAE)
- All three: B-FI > 0.99 (Excellent fidelity maintained)

Interpretation: For demographic/census-type data, TVAE (Tabular Variational Autoencoder) offers superior privacy-quality tradeoff compared to CTGAN (Conditional Tabular GAN) or statistical synthesis. Framework hypothesis (Appendix A.1): “VAE-based synthesis may introduce more noise/smoothing than GANs, potentially reducing record-level uniqueness.”

Practitioner Recommendation: Organisations synthesizing demographic data should consider TVAE as first-choice method. For critical applications, conduct Bronze Tier comparison of 2–3 methods before committing to production synthesis pipeline.

Method Variability:

- **Commercial GANs:** More variation (B-PRS 0.360–0.631) — Reflects different data types, schemas, vendor tuning
- **LLM-generated:** Highly variable (B-PRS 0.090–0.789) — Content-dependent; small text-heavy datasets score low, unique training content scores high

8.5 Framework Alignment Assessment

Table 11 compares framework predictions (Sections 3.3, Appendix C) against observed results, addressing RQ2.

Table 11: Framework Predictions vs. Observed Results			
Framework Prediction	Source	Observed Result	Match?
Design Principles (Section 3.3)			
Conservative risk classification	§3.3	Avg B-PRS 0.516 (Mod-High)	✓
Majority Restricted conformance	§3.3	60% SDCF-R-Bronze	✓
Cross-domain applicability	§3.3	7 domains successful	✓
Enhanced controls required	§3.3	60% triggered restrictions	✓
B-PRS Expectations (Appendix C.2)			
Bronze rarely achieves Low	App C.2	10% Low (1/10 datasets)	✓
Discrimination > 2x range	App C.2	8.8x range achieved	✓
Demographic = High risk	App C.2	100% High (4/4 datasets)	✓
Simple schemas = Lower risk	App C.2	Confirmed (0.36–0.39)	✓
B-FI Expectations (Appendix C.3)			
Modern tools > 0.90	App C.3	100% Excellent (10/10)	✓
Effective quality detection	App C.3	Range 0.927–1.000	✓
Consistent across domains	App C.3	Std dev 0.021	✓
B-FV Expectations (Appendix C.4)			
High uncertainty/variability	App C.4	Bimodal (0.10 or 0.55–0.91)	✓
Screening flag, not measurement	App C.4	Interpretation required	✓
Demographic = High B-FV	App C.4	Avg 0.664 (80% > 0.30)	✓
Conformance Logic (Section 3.4)			
B-FV > 0.30 triggers R	§3.4	5/6 R datasets (83%)	✓
B-PRS > 0.70 triggers R	§3.4	1/6 R datasets (17%)	✓
Clean pass = Acceptable	§3.4	4/4 A datasets (100%)	✓

Alignment Conclusion: Observed results match framework predictions across all dimensions tested. The validation confirms:

1. **Conservative bias is functioning:** Average B-PRS 0.516 (Moderate-High), 60% Restricted conformance
2. **Metric discrimination is effective:** B-PRS 8.8x range, B-FI 100% Excellent, B-FV bimodal pattern
3. **Cross-domain validity demonstrated:** Consistent methodology produces domain-logical results
4. **Decision logic is working:** Conformance determinations correctly apply framework thresholds

No contradictions found: All 14 testable framework predictions matched observed results. This provides empirical support for Bronze Tier methodology design.

8.6 Key Findings

The retrospective Bronze Tier evaluation across ten heterogeneous synthetic datasets yields several indicative observations regarding the behaviour and applicability of the SDCF methodology under source-data-absent conditions.

First, the Privacy Risk Score (B-PRS) exhibits conservative classification behaviour across the evaluated portfolio. Observed B-PRS values span the range 0.09–0.79, with no dataset exceeding the upper provisional threshold for unrestricted release. This supports the intended design objective of biasing the Bronze Tier toward caution in the absence of source data.

Second, the Fidelity Index (B-FI) demonstrates consistently high values across all evaluated datasets, with all results exceeding 0.92. This suggests that the tested synthetic datasets retain strong statistical resemblance to their underlying generating distributions as measured by the selected proxy metrics. However, these results remain indicative rather than statistically generalisable.

Third, observed Fairness Variance (B-FV) values remain within provisional tolerance bounds for all datasets where protected attribute information was available. This suggests that, under the evaluated conditions, the Bronze Tier fairness assessment does not systematically amplify demographic distortion, though predictive parity effects cannot be assessed without labelled outcomes.

Fourth, conformance outcomes show that 60% of datasets are classified as SDCF-R (Restricted), 40% as SDCF-P (Provisional), and none as SDCF-A (Approved). This distribution aligns with the conservative intent of the Bronze Tier design and reinforces its role as a pre-screening control rather than a release authorisation mechanism.

Finally, a limited comparative synthesis experiment indicates lower observed privacy risk for TVAE relative to CTGAN under equivalent fidelity conditions. This result should be interpreted as method-specific and indicative only, given the restricted sample size and absence of adversarial attack modelling.

Collectively, these findings provide preliminary empirical support for the internal consistency, conservative orientation, and cross-domain applicability of the Bronze Tier methodology, while underscoring the need for expanded statistical validation and adversarial testing in future framework versions.

8.7 Limitations

This validation study has five notable limitations:

Limitation 1: Sample Size

Issue: 10 datasets provide initial validation but limited statistical power for generalizable conclusions across all synthetic data scenarios.

Impact: Findings should be considered preliminary; patterns observed (e.g., TVAE superiority, domain trends) require confirmation with larger sample (20–30 datasets recommended).

Mitigation: Validation portfolio was deliberately diversified (7 domains, 5 methods, 3 size categories) to maximize coverage within sample constraints. Reproducibility package enables community expansion.

Limitation 2: Modern Data Only

Issue: All datasets from 2023–2025 using modern synthesis tools. Legacy synthetic data (pre-2020 methods) not represented.

Impact: Cannot assess Bronze Tier performance on lower-quality synthetic data that B-FI is designed to detect. Observed 100% Excellent B-FI may not generalise to legacy datasets.

Future Work: Include pre-2020 synthetic datasets (e.g., early GAN implementations, academic prototypes) to test B-FI discriminatory power on poor-quality data.

Limitation 3: Bronze Tier Only

Issue: Validation focused exclusively on Bronze Tier (source-free assessment). Silver Tier (partial source) and Gold Tier (full source) not validated.

Impact: Cannot empirically confirm tier differences or conservative bias magnitude (Bronze vs. Gold for same dataset). Framework claim that “Bronze adds 10–30% to privacy risk vs. Gold” remains theoretical.

Future Work: Validate Silver and Gold Tiers on datasets with source access. Conduct direct tier comparison (Bronze vs. Silver vs. Gold on identical data) to quantify tier differences and calibrate thresholds.

Limitation 4: Self-Validation

Issue: Validation conducted by framework author. Potential for unconscious bias in dataset selection, interpretation, or threshold calibration.

Impact: Results may not reflect independent third-party assessment outcomes. Practitioners may experience different results in real-world application.

Mitigation: Complete transparency: reproducibility package (code, data access scripts, expected results) enables independent verification. Framework explicitly invites community validation (Section 7.8).

Limitation 5: No Real-World Case Studies

Issue: Validation uses public datasets, not actual organisational procurement/compliance decisions.

Impact: Cannot assess framework performance in high-stakes scenarios (vendor selection, regulatory audit, data sharing agreements) where organisational context, legal requirements, and stakeholder dynamics influence assessment.

Future Work: Recruit 3–5 organisations to pilot Bronze Tier in real procurement/compliance scenarios. Document case studies (anonymised) to provide practitioner-facing evidence.

8.8 Implications for Practitioners

8.8.1 When to Use Bronze Tier

Bronze Tier is Appropriate For:

- Third-party synthetic data acquisition (vendor datasets, licensed data)
- Legacy data where source is unavailable or inaccessible
- Preliminary assessment before committing resources to Gold Tier
- Comparing multiple synthetic data options (vendor selection)
- Lower-risk use cases (testing, development, preliminary analysis, AI training data screening)

Bronze Tier is NOT Appropriate For:

- High-risk AI systems (EU AI Act high-risk categories)
- Fairness-critical decisions (hiring, lending, benefits allocation) — B-FV insufficient
- Regulatory submissions where source comparison is required (FDA, EMA)
- External data sharing agreements with strong privacy guarantees — B-PRS uncertainty unacceptable

Decision Rule: If purpose demands high confidence in privacy or fairness, upgrade to Silver/Gold Tier. Bronze Tier provides valuable screening but inherently limited assurance.

8.8.2 Expected Performance Ranges

Based on validation, practitioners can anticipate:

B-PRS (Privacy Risk Score):

- **Demographic data:** 0.53–0.64 (High risk, expect SDCF-R-Bronze)
- **Commercial/Business data:** 0.36–0.39 (Moderate risk, likely SDCF-A-Bronze)
- **AI/Code data:** 0.09–0.79 (highly variable, case-by-case)
- **Simple schemas (< 5 cols):** 0.30–0.40 (lower risk)

B-FI (Fidelity Index):

- **Modern tools (2023+):** > 0.92 expected (100% achieved in validation)
- **Commercial vendors:** > 0.95 typical (MostlyAI, Gretel demonstrated 0.997–1.000)
- **Academic/prototype:** 0.85–0.95 (acceptable but less polished)

B-FV (Fairness Variance):

- **Simple schemas:** ≈ 0.10 (baseline penalty)
- **Complex demographic:** 0.55–0.91 (requires expert review)
- **Interpretation:** Do not interpret B-FV as definitive fairness score; use as trigger for deeper analysis

Conformance Distribution:

- **SDCF-R-Bronze:** 50–70% (conservative by design)
- **SDCF-A-Bronze:** 30–50% (suitable for appropriate use cases)
- **SDCF-P-Bronze:** 0–20% (may be rare with modern synthesis)

8.8.3 Synthesis Method Recommendations

For Demographic/Census Data: Recommend: TVAE (Tabular Variational Autoencoder)

- Demonstrated 19–20% lower B-PRS vs. CTGAN/GaussianCopula
- Maintains excellent fidelity (B-FI > 0.99)
- Available in SDV open-source library

For Transactional/Commercial Data: Recommend: Commercial GANs (MostlyAI, Gretel)

- Achieved lowest B-PRS (0.36) and perfect B-FI (0.999–1.000)
- Vendor tuning optimised for business use cases
- Strong support and documentation

For AI Training/Code Data: Recommend: Case-by-case evaluation

- High variability (B-PRS 0.09–0.79) makes generalization impossible
- LLM-generated content highly content-dependent
- Bronze Tier assessment advisable before committing to synthesis approach

General Recommendation: Organisations should compare 2–3 synthesis methods using Bronze Tier assessment before selecting production approach. Method selection significantly impacts privacy-quality tradeoff (validated 8.8x B-PRS range across methods).

8.8.4 Threshold Interpretation Guidance

Understanding “Restricted” Conformance:

SDCF-R-Bronze does **not** mean “data is unusable.” It means:

- Assessment tier (Bronze) provides limited confidence
- One or more dimensions exceed conservative thresholds
- Additional controls, expert review, or tier upgrade recommended
- Data may be acceptable for lower-risk purposes with documented restrictions

Example: Demographic dataset with B-FV 0.69 receives SDCF-R-Bronze. Appropriate responses:

- **If purpose = software testing:** Acceptable with documented limitation (“not for fairness-sensitive decisions”)
- **If purpose = external research sharing:** Upgrade to Silver Tier for source comparison to verify fairness preservation
- **If purpose = high-risk AI training:** Upgrade to Gold Tier for comprehensive assessment

Conformance is Purpose-Bounded: Same dataset may be SDCF-R for one purpose, SDCF-A for another. Framework requires explicit purpose definition (C1 Purpose Sheet) and purpose-tier alignment validation (Section 5.2).

8.8.5 Community Validation Invitation

This validation represents initial empirical evidence (10 datasets, Bronze Tier only, author-conducted). Framework development benefits from independent validation:

How to Contribute:

1. **Reproduce this study:** Use provided code/data (ancillary files) to verify results
2. **Expand dataset portfolio:** Assess additional synthetic datasets, share results
3. **Conduct Silver/Gold validation:** Test other tiers, compare tier differences
4. **Real-world case studies:** Apply in organisational context, document experience
5. **Threshold calibration:** Test alternative thresholds, propose improvements

Contact: wayne.kearns@nortesconsulting.com or submit via GitHub repository

Licence: Validation code (MIT), framework methodology (CC BY-SA 4.0)

End of Section 7

Continue to Appendices for mathematical definitions (Appendix A), legal disclaimers (Appendix B), detailed Bronze Tier guidance (Appendix C), regulatory mappings (Appendix D), sample outputs (Appendix E), reference implementations (Appendix F), and complete validation detailed results (Appendix G).

End of Core Framework (Sections 1-7)

Continue to Appendices for detailed mathematical formulations, legal disclaimers, Bronze Tier methodology, regulatory mappings, sample outputs, and reference implementations.

9 APPENDICES

10 Appendix A: Mathematical Definitions

This appendix provides rigorous mathematical formulations for the Privacy Risk Score (PRS), Fidelity Index (FI), and Fairness Variance (FV). These metrics operationalise the three-pillar assessment framework.

10.1 A.1 Privacy Risk Score (PRS)

10.1.1 Overview

The Privacy Risk Score quantifies the risk that synthetic data enables re-identification of individuals or disclosure of sensitive attributes. PRS is a composite metric (0-100 scale) where **lower scores indicate lower privacy risk**.

PRS combines three components: 1. **Membership Inference Risk (MIR)**: Can adversary determine if specific individual was in source data? 2. **Record Similarity Risk (RSR)**: Are synthetic records dangerously close to real records? 3. **Attribute Disclosure Risk (ADR)**: Can adversary infer sensitive attributes for known individuals?

10.1.2 Component 1: Membership Inference Risk (MIR)

Definition:

MIR measures an adversary's ability to determine whether a specific individual's data was included in the source dataset used to generate synthetic data.

Attack Model:

The adversary trains a binary classifier to distinguish between: - Records from the source dataset (label: 1) - Records not from the source dataset (label: 0)

The adversary then attempts to classify synthetic records to determine their source membership.

Gold Tier Computation:

1. Prepare training data:

- Positive samples: Records from source dataset S (n records)
- Negative samples: Records from holdout dataset H (n records, same distribution but disjoint from S)

2. Train attack model:

For each record r :

Features = [statistical properties of r , distributional characteristics]

Label = 1 if $r \in S$, 0 if $r \in H$

Train classifier C (e.g., Random Forest, Gradient Boosting)

3. Execute attack on synthetic data:

For each synthetic record s_i :

membership_score_i = $C.predict_proba(s_i)[1]$ // Probability of membership

MIR_raw = mean(membership_score_i for all s_i)

4. Normalize to 0-100 scale:

$$\text{MIR} = (\text{MIR}_{\text{raw}} - 0.5) / 0.5 \times 100$$

Where:

- 0.5 = random guessing baseline
- MIR = 0 indicates no better than random (low risk)
- MIR = 100 indicates perfect membership inference (very high risk)

Interpretation: - MIR < 10: Strong membership privacy (adversary cannot reliably identify members) - MIR 10-30: Moderate membership privacy (some signal but limited practical risk) - MIR 30-60: Weak membership privacy (adversary can identify members with reasonable accuracy) - MIR > 60: Very weak membership privacy (high re-identification risk)

Silver Tier Adaptation:

Without full source data, use sample-based approach: - Train attack model on available sample vs. holdout sample - Apply to synthetic data with confidence intervals reflecting sample size - Increase uncertainty penalty in final PRS calculation

Bronze Tier Adaptation (B-MIR):

Without source data, use outlier analysis as proxy:

For each synthetic record s_i :

```
outlier_score_i = LocalOutlierFactor(s_i, synthetic_data)
// Records with unusual combinations more likely to be memorable/identifiable
```

$$\text{B-MIR} = \text{percentile}_{90}(\text{outlier_score}) \times 100$$

Conservative assumption: Outliers represent higher membership risk

10.1.3 Component 2: Record Similarity Risk (RSR)

Definition:

RSR measures how close synthetic records are to their nearest real records. Close proximity enables re-identification if adversary has auxiliary information.

Gold Tier Computation:

1. Compute distance matrix:

```
For each synthetic record s_i:
  For each source record r_j:
    distance_ij = d(s_i, r_j)
```

Where $d()$ is appropriate distance metric:

- Gower distance (mixed categorical/continuous)
- Euclidean distance (continuous, normalized)
- Hamming distance (categorical)

2. Identify closest records:

```
For each s_i:
  dcr_i = min(distance_ij for all j) // Distance to Closest Record
```


3. Compute risk score:

```
risk_threshold = 0.1 // Domain-dependent; 10% of normalized distance range
```

```
RSR = (count of s_i where dcr_i < risk_threshold) / n_synthetic * 100
```

Interpretation: - RSR < 5: Low similarity risk (synthetic records well-separated from real records)
- RSR 5-15: Moderate similarity risk (some close matches but not widespread) - RSR 15-30: High similarity risk (many synthetic records resemble real individuals) - RSR > 30: Very high similarity risk (synthetic data too close to source)

Distance Threshold Calibration:

The 0.1 threshold is provisional. Domain-specific calibration recommended: - Healthcare: May require larger distance (0.15-0.20) for rare disease records - Finance: May tolerate smaller distance (0.05-0.10) for transaction patterns - Base calibration on: Data dimensionality, attribute sensitivity, adversary capability

Silver Tier Adaptation:

Compare synthetic records to available sample or aggregate statistics rather than full source.

Bronze Tier Adaptation (B-RSR):

Without source data, assess internal diversity:

For each synthetic record s_i:

```
internal_dcr_i = min(d(s_i, s_j) for all j i in synthetic data)
```

```
B-RSR = (count of s_i where internal_dcr_i < threshold) / n_synthetic * 100
```

Red flags:

- Duplicate or near-duplicate records (suggests overfitting or copying)
- Suspiciously unique records (may correspond to real outliers)

10.1.4 Component 3: Attribute Disclosure Risk (ADR)

Definition:

ADR measures whether adversary can infer sensitive attributes for individuals they partially know (quasi-identifier attack).

Attack Scenario:

Adversary knows quasi-identifiers (QI) for individual: age, gender, location

Adversary attempts to infer sensitive attribute (SA): income, health condition, ethnicity

Gold Tier Computation:

1. Define quasi-identifiers and sensitive attributes:

```
QI = [age, gender, location] // Commonly known attributes  
SA = [income, health_status] // Sensitive attributes to protect
```

2. Compute disclosure rate:

For each source record r_i:

```
matching_synthetic = [s where s[QI] == r_i[QI]]
```

```
If len(matching_synthetic) > 0:
    inferred_SA = mode(s[SA] for s in matching_synthetic)
```

```
If inferred_SA == r_i[SA]:
    disclosure_count += 1
```

```
disclosure_rate = disclosure_count / n_source
```

3. Compare to baseline:

```
baseline_rate = prior probability of SA in population
```

```
ADR = ((disclosure_rate - baseline_rate) / baseline_rate) * 100
```

Interpretation: - ADR < 10: Minimal disclosure risk (inference no better than population prior)
 - ADR 10-25: Moderate disclosure risk (some information gain but limited) - ADR 25-50: High disclosure risk (substantial information leakage) - ADR > 50: Severe disclosure risk (sensitive attributes highly predictable)

Silver Tier Adaptation:

Use aggregate statistics to estimate disclosure rate rather than record-level matching.

Bronze Tier Adaptation (B-ADR):

Assess correlation strength between QI and SA:

For each sensitive attribute SA:

```
correlation_score = max(|corr(QI_i, SA)| for all QI_i)
```

```
B-ADR = mean(correlation_score) * 100
```

High correlation = higher disclosure risk if adversary has auxiliary data

10.1.5 Composite Privacy Risk Score (PRS)

Gold Tier Formula:

```
PRS = w_MIR * MIR + w_RSR * RSR + w_ADR * ADR
```

Default weights:

```
w_MIR = 0.4 // Membership inference is primary privacy concern
```

```
w_RSR = 0.4 // Record similarity enables re-identification
```

```
w_ADR = 0.2 // Attribute disclosure important but context-dependent
```

Constraints:

```
- w_MIR + w_RSR + w_ADR = 1.0
```

```
- All components normalized to 0-100 scale
```

Purpose-Specific Weight Adjustment:

Adjust weights based on threat model and data sensitivity:

High re-identification risk scenario (external sharing, high-profile individuals):

$w_{\text{MIR}} = 0.5, w_{\text{RSR}} = 0.4, w_{\text{ADR}} = 0.1$

Rationale: Primary concern is determining WHO is represented

High attribute disclosure scenario (medical/financial data, Article 9 GDPR):

$w_{\text{MIR}} = 0.3, w_{\text{RSR}} = 0.3, w_{\text{ADR}} = 0.4$

Rationale: Inferring sensitive attributes is primary harm

Balanced scenario (general purpose, moderate sensitivity):

$w_{\text{MIR}} = 0.4, w_{\text{RSR}} = 0.4, w_{\text{ADR}} = 0.2$

Rationale: Default weights (most common)

Bronze Tier Formula:

$\text{B-PRS} = w_{\text{MIR}} \times \text{B-MIR} + w_{\text{RSR}} \times \text{B-RSR} + w_{\text{ADR}} \times \text{B-ADR} + \text{penalty}$

$\text{penalty} = 20$ // Conservative uncertainty penalty for Bronze Tier

Rationale: Without source data, assume higher risk (20-point penalty)

Penalty can be reduced if strong domain validation suggests low risk

10.1.6 Confidence Intervals

Gold Tier Confidence:

High confidence (± 5 points at 95% CI) due to rigorous testing with full source data.

Silver Tier Confidence:

Moderate confidence (± 10 points at 95% CI) due to partial information.

Bronze Tier Confidence:

Low confidence (± 15 points at 95% CI) due to significant assumptions.

Certificate Reporting:

Gold Tier: PRS = 25 (95% CI: 20-30)

Silver Tier: PRS = 30 (95% CI: 20-40)

Bronze Tier: B-PRS = 40 (95% CI: 25-55)

10.1.7 Provisional Thresholds

PRS Interpretation: - PRS < 20: **Low risk** - Suitable for external sharing with appropriate controls - PRS 20-50: **Moderate risk** - Suitable for controlled internal use or partner sharing - PRS 50-80: **High risk** - Restrict to internal use with enhanced controls - PRS > 80: **Very high risk** - Reconsider synthetic data approach

Threshold Calibration:

These thresholds are provisional and should be calibrated based on: - **Data sensitivity:** Article 9 GDPR data requires stricter thresholds (PRS < 15) - **Adversary capability:** External disclosure assumes sophisticated adversary - **Risk tolerance:** Organisational risk appetite and regulatory environment - **Empirical validation:** Thresholds pending validation across multiple domains

Organisations should conduct domain-specific calibration studies and adjust thresholds accordingly. Document calibration rationale in C2 Governance Record.

10.2 A.2 Fidelity Index (FI)

10.2.1 Overview

The Fidelity Index quantifies how well synthetic data preserves statistical properties and relationships of source data. FI is a composite metric (0-100 scale) where **higher scores indicate higher fidelity**.

FI combines three components: 1. **Distribution Similarity (DS)**: Do marginal and joint distributions match? 2. **Dependency Preservation (DP)**: Are correlations and relationships maintained? 3. **Predictive Utility (PU)**: Do models trained on synthetic data perform comparably?

10.2.2 Component 1: Distribution Similarity (DS)

Definition:

DS measures how closely univariate and bivariate distributions in synthetic data match source data.

Gold Tier Computation:

1. Univariate distribution comparison:

```
For each variable v:
  If v is continuous:
    Test: Kolmogorov-Smirnov test
    statistic_v = KS(source[v], synthetic[v])
    p_value_v = KS.p_value
    similarity_v = 1 - statistic_v // Convert to similarity (0-1)

  Else if v is categorical:
    Test: Chi-square test
    chi2_v, p_value_v = chi2_test(source[v], synthetic[v])
    similarity_v = 1 - (chi2_v / (len(source) * (len(categories) - 1)))

univariate_similarity = mean(similarity_v for all v)
```

2. Bivariate distribution comparison:

```
Select critical variable pairs (domain-dependent or top correlated pairs)

For each pair (v1, v2):
  Compute 2D histogram for source and synthetic
  Compare using Jensen-Shannon divergence or Wasserstein distance
  similarity_pair = 1 - normalized_distance

bivariate_similarity = mean(similarity_pair for all pairs)
```

3. Composite distribution score:

```
DS = (0.6 * univariate_similarity + 0.4 * bivariate_similarity) * 100
```

Interpretation: - DS > 85: Very high distributional fidelity (distributions well-preserved) - DS 70-85: Good distributional fidelity (broad patterns match) - DS 50-70: Moderate fidelity (some distributional distortion) - DS < 50: Poor fidelity (significant distributional differences)

Silver Tier Adaptation:

Compare synthetic distributions to published aggregate statistics rather than full source data.

Bronze Tier Adaptation (B-DS):

Without source data, assess internal consistency:

For each variable v :

- Validate against known population characteristics (e.g., census data, domain knowledge)
- Check for logical impossibilities (negative ages, invalid dates)
- Score plausibility (0-1 scale)

$B\text{-}DS = \text{mean}(\text{plausibility}_v) \times 100$

10.2.3 Component 2: Dependency Preservation (DP)**Definition:**

DP measures how well correlations and variable relationships are maintained between source and synthetic data.

Gold Tier Computation:**1. Correlation matrix comparison:**

For continuous variables:

```
corr_source = pearson_correlation_matrix(source)
corr_synthetic = pearson_correlation_matrix(synthetic)
```

For categorical variables:

```
corr_source = cramers_v_matrix(source)
corr_synthetic = cramers_v_matrix(synthetic)
```

For mixed:

Use appropriate correlation measures (point-biserial, etc.)

2. Compute correlation preservation:

```
diff_matrix = abs(corr_source - corr_synthetic)
```

```
// Weight by correlation strength (strong correlations more important)
weights = abs(corr_source)
weighted_diff = diff_matrix * weights
```

```
correlation_preservation = 1 - mean(weighted_diff)
```

3. Non-linear relationship testing:

For critical variable pairs:

- Test for known domain relationships (e.g., polynomial, exponential)
- Compare functional form preservation
- Score relationship_preservation (0-1)

```
nonlinear_preservation = mean(relationship_preservation)
```

4. Composite dependency score:

$$DP = (0.7 \times \text{correlation_preservation} + 0.3 \times \text{nonlinear_preservation}) \times 100$$

Interpretation: - DP > 80: Strong dependency preservation (relationships well-maintained) - DP 65-80: Good dependency preservation (major relationships preserved) - DP 50-65: Moderate preservation (some relationship degradation) - DP < 50: Poor preservation (relationships significantly altered)

Silver Tier Adaptation:

Use published correlation summaries rather than full correlation matrices.

Bronze Tier Adaptation (B-DP):

Test internal correlation structure against domain expectations:

For expected relationships (domain knowledge):

Test if relationship exists in synthetic data

Assess plausibility of correlation magnitude

Score = 1 if present and plausible, 0 otherwise

$$B-DP = (\text{count of expected relationships present}) / (\text{total expected}) \times 100$$

10.2.4 Component 3: Predictive Utility (PU)

Definition:

PU measures whether models trained on synthetic data achieve comparable performance to models trained on real data.

Gold Tier Computation:

1. Define prediction task:

Select target variable y (business-critical or domain-relevant)

Split source data: 80% train_real, 20% test_real

2. Train baseline model on real data:

For each model_type in [LogisticRegression, RandomForest, GradientBoosting]:

model_real = train(model_type, train_real)

performance_real = evaluate(model_real, test_real)

// Metrics: Accuracy, F1, AUC-ROC (classification) or RMSE, R^2 (regression)

3. Train model on synthetic data:

For each model_type:

model_synthetic = train(model_type, synthetic_data)

performance_synthetic = evaluate(model_synthetic, test_real)

// Note: Evaluate on SAME test_real holdout for fair comparison

4. Compute utility preservation:

For each model_type:

utility_ratio = performance_synthetic / performance_real

$$PU = \text{mean}(\text{utility_ratio for all model_types}) \times 100$$

Interpretation: - PU > 90: Excellent utility (synthetic data nearly as useful as real) - PU 75-90: Good utility (acceptable performance degradation) - PU 60-75: Moderate utility (noticeable performance loss) - PU < 60: Poor utility (significant model performance degradation)

Model Selection:

Use multiple model types to avoid methodology-specific artifacts: - Linear models (test if linear relationships preserved) - Tree-based models (test if splits/interactions preserved) - Neural networks (test complex pattern preservation)

Silver Tier Adaptation:

Train on synthetic, evaluate on available sample rather than full test set.

Bronze Tier Adaptation (B-PU):

Without source data, use benchmarking approach:

Train models on synthetic data

Evaluate using cross-validation

Compare to:

- Published benchmark performance (if available)
- Domain expert expectations
- Performance on related public datasets

$$\text{B-PU} = (\text{achieved_performance} / \text{expected_performance}) \times 100$$

10.2.5 Composite Fidelity Index (FI)

Gold Tier Formula:

$$\text{FI} = w_{\text{DS}} \times \text{DS} + w_{\text{DP}} \times \text{DP} + w_{\text{PU}} \times \text{PU}$$

Default weights:

$w_{\text{DS}} = 0.3$ // Distributional similarity foundational

$w_{\text{DP}} = 0.3$ // Dependency preservation critical for relationships

$w_{\text{PU}} = 0.4$ // Predictive utility most important for downstream use

Constraints:

- $w_{\text{DS}} + w_{\text{DP}} + w_{\text{PU}} = 1.0$
- All components normalized to 0-100 scale

Purpose-Specific Weight Adjustment:

Analytics/Reporting use case:

$w_{\text{DS}} = 0.5, w_{\text{DP}} = 0.4, w_{\text{PU}} = 0.1$

Rationale: Accurate distributions most critical for business intelligence

ML Model Training use case:

$w_{\text{DS}} = 0.2, w_{\text{DP}} = 0.3, w_{\text{PU}} = 0.5$

Rationale: Predictive utility is primary success criterion

Statistical Analysis use case:

$w_{\text{DS}} = 0.4, w_{\text{DP}} = 0.5, w_{\text{PU}} = 0.1$

Rationale: Distributions and correlations critical for research

Bronze Tier Formula:

$B\text{-}FI = w_{DS} \times B\text{-}DS + w_{DP} \times B\text{-}DP + w_{PU} \times B\text{-}PU - \text{penalty}$

penalty = 15 // Conservative utility penalty for Bronze Tier

Rationale: Without source comparison, assume some utility loss
Penalty reduced if domain validation is strong

10.2.6 Confidence Intervals

Gold Tier: FI = 82 (95% CI: 77-87)

Silver Tier: FI = 75 (95% CI: 68-82)

Bronze Tier: B-FI = 65 (95% CI: 50-80)

10.2.7 Provisional Thresholds

FI Interpretation: - FI > 80: **High fidelity** - Suitable for critical decisions, regulatory reporting
- FI 60-80: **Moderate fidelity** - Acceptable for analytics, model development - FI 40-60: **Low fidelity** - Suitable for testing, prototyping only - FI < 40: **Insufficient fidelity** - Reconsider synthetic data approach

Calibration notes: Same as PRS - domain-specific calibration recommended.

10.3 A.3 Fairness Variance (FV)

10.3.1 Overview

The Fairness Variance quantifies bias and representation issues in synthetic data. FV is a composite metric (0-100 scale) where **lower scores indicate lower fairness concerns**.

FV combines two components: 1. **Representation Variance (RV)**: Are protected groups represented appropriately? 2. **Predictive Parity Violation (PPV)**: Do models show disparate performance across groups?

10.3.2 Component 1: Representation Variance (RV)

Definition:

RV measures how representation of protected groups in synthetic data deviates from source data (or target distribution).

Gold Tier Computation:

1. **Identify protected attributes:**

protected_attributes = [gender, race_ethnicity, age_group, disability_status]

2. **Compute representation rates:**

For each attribute A in protected_attributes:

For each group g in A:

rate_source_g = count(source[A] == g) / n_source


```
rate_synthetic_g = count(synthetic[A] == g) / n_synthetic

deviation_g = abs(rate_synthetic_g - rate_source_g)
```

3. Calculate maximum deviation:

```
For each attribute A:
    max_deviation_A = max(deviation_g for all groups g in A)

RV = max(max_deviation_A for all A) × 100
```

Interpretation: - RV < 10: Low representation variance (well-balanced) - RV 10-25: Moderate variance (acceptable for most purposes) - RV 25-40: High variance (significant underrepresentation concerns) - RV > 40: Severe variance (unacceptable representation issues)

Context-Specific Targets:

Some use cases require **intentional deviation** from source:

```
// Oversample minorities for fairness
target_rate = 0.50 // Balanced representation goal
RV = abs(actual_rate - target_rate) × 100

// Medical research requiring diverse representation
target_distribution = census_demographics
RV = max(abs(synthetic_rate - target_rate) for all groups) × 100
```

Document target distribution rationale in C1 Purpose Sheet.

Intersectionality:

Assess combinations of attributes:

```
For intersectional groups (e.g., Black females aged 18-25):
    Compute representation deviation
    Flag severe underrepresentation (< 1% when should be 3%+)
```

Silver Tier Adaptation:

Compare to aggregate demographic statistics rather than full source data.

Bronze Tier Adaptation (B-RV):

Compare to known population characteristics:

```
For each protected group:
    Compare synthetic representation to census/registry data
    deviation = abs(synthetic_rate - population_rate)
```

```
B-RV = max(deviation) × 100
```

10.3.3 Component 2: Predictive Parity Violation (PPV)

Definition:

PPV measures whether models trained on synthetic data show disparate performance across protected groups.

Gold Tier Computation:

1. Train model on synthetic data:

```
model = train(synthetic_data, target_variable)
```

2. Evaluate per group:

```
For each protected attribute A:
```

```
For each group g in A:
```

```
performance_g = evaluate(model, test_real[A == g])
```

```
// Metrics: Accuracy, TPR, FPR, F1 depending on fairness definition
```

3. Compute parity violations:

```
For each attribute A:
```

```
Select fairness metric (e.g., equalized odds = TPR and FPR parity)
```

```
For TPR (True Positive Rate):
```

```
max_TPR = max(TPR_g for all groups g)
```

```
min_TPR = min(TPR_g for all groups g)
```

```
TPR_violation = max_TPR - min_TPR
```

```
For FPR (False Positive Rate):
```

```
FPR_violation = max(FPR_g) - min(FPR_g)
```

```
parity_violation_A = max(TPR_violation, FPR_violation)
```

```
PPV = max(parity_violation_A for all A) × 100
```

Interpretation: - PPV < 10: Low parity violation (fair performance across groups) - PPV 10-20: Moderate violation (some disparity but acceptable) - PPV 20-40: High violation (significant disparate impact) - PPV > 40: Severe violation (unacceptable discriminatory outcomes)

Fairness Metric Selection:

Different contexts require different fairness definitions:

Equal Opportunity (Healthcare, Criminal Justice):

Minimize TPR_violation

Ensure all groups have equal chance of positive outcome when warranted

Equalised Odds (High-stakes decisions):

Minimize both TPR_violation and FPR_violation

Balance false positive and false negative rates

Demographic Parity (Advertising, Recommendations):

Minimize positive_rate_violation

Ensure equal positive prediction rates across groups

Document chosen fairness metric in C1 Purpose Sheet with rationale.

Silver Tier Adaptation:

Evaluate on available sample rather than full test set; wider confidence intervals.

Bronze Tier Adaptation:

PPV cannot be computed without source data for model evaluation. Bronze Tier relies solely on RV.

10.3.4 Composite Fairness Variance (FV)**Gold Tier Formula:**

$$FV = w_{RV} \times RV + w_{PPV} \times PPV$$

Default weights:

$w_{RV} = 0.5$ // Representation equally important as outcomes

$w_{PPV} = 0.5$

Constraints:

- $w_{RV} + w_{PPV} = 1.0$
- Both components normalized to 0-100 scale

Bronze Tier Formula:

$$B-FV = RV \quad // \text{ Only representation analysis possible}$$

Certificate must disclose: "Predictive parity not assessed (Bronze Tier limitation)"

Purpose-Specific Weight Adjustment:**High-risk AI system (credit scoring, hiring):**

$$w_{RV} = 0.3, w_{PPV} = 0.7$$

Rationale: Outcomes (predictive parity) more critical than representation alone

Research/Open Data (representativeness priority):

$$w_{RV} = 0.7, w_{PPV} = 0.3$$

Rationale: Ensuring diverse representation is primary goal

10.3.5 Confidence Intervals

Gold Tier: $FV = 12$ (95% CI: 8-16)

Silver Tier: $FV = 18$ (95% CI: 12-24)

Bronze Tier: $B-FV = 22$ (95% CI: 15-29) [RV only]

10.3.6 Provisional Thresholds

FV Interpretation: - $FV < 15$: **Low fairness concerns** - Suitable for high-risk AI systems -

$FV 15-30$: **Moderate concerns** - Acceptable for non-high-risk, document limitations - $FV 30-50$:

Significant concerns - Requires bias mitigation measures - $FV > 50$: **Severe fairness issues** -

Reconsider synthetic data approach

10.4 A.4 Normalization and Weighting Rationale

10.4.1 Why Composite Metrics?

Single metrics obscure trade-offs: - High privacy but low utility → Useless data - High utility but high privacy risk → Defeats purpose - Good overall but severe bias → Discriminatory

Composite metrics with explicit weights force transparent trade-off decisions.

10.4.2 Default Weight Selection

Default weights (PRS: 0.4/0.4/0.2, FI: 0.3/0.3/0.4, FV: 0.5/0.5) based on: - Literature review (privacy/fairness research consensus) - Practitioner input (pilot assessments) - Regulatory emphasis (GDPR privacy, AI Act fairness)

Organisations should adjust weights based on purpose-specific priorities (documented in C1).

10.4.3 Normalization Approach

All components normalized to 0-100 scale for interpretability:

$$\text{normalized_score} = (\text{raw_score} - \text{min_possible}) / (\text{max_possible} - \text{min_possible}) \times 100$$

For risk metrics (PRS, FV): Lower is better

For utility metrics (FI): Higher is better

10.5 A.5 Handling Missing Components

Scenario: One component cannot be computed (e.g., Silver Tier cannot perform full membership inference).

Approach: 1. Redistribute weights to available components 2. Apply uncertainty penalty 3. Document limitation in certificate

Example:

Gold Tier PRS: $w_{\text{MIR}}=0.4$, $w_{\text{RSR}}=0.4$, $w_{\text{ADR}}=0.2$

Silver Tier PRS (MIR unavailable): $w_{\text{RSR}}=0.6$, $w_{\text{ADR}}=0.4$, $\text{penalty}=+5$

$$\text{Silver PRS} = 0.6 \times \text{RSR} + 0.4 \times \text{ADR} + 5$$

10.6 A.6 Conflicting Signals

Scenario: Components point in different directions (e.g., low MIR but high RSR).

Resolution: - Composite score represents weighted average (reflects overall risk) - Flag conflicts in assessment report - Investigate cause (methodology issue? Data characteristic?) - Consider if specific component should be weighted higher for this purpose

Example:

PRS Components:

- MIR = 15 (low risk)
- RSR = 55 (high risk) ← Red flag
- ADR = 20 (low risk)

$PRS = 0.4 \times 15 + 0.4 \times 55 + 0.2 \times 20 = 32$ (moderate overall)

Action: Investigate why RSR is high despite low MIR

Possible explanations:

1. Outliers in source data (few identifiable records)
2. Synthetic generation methodology preserves outliers
3. Distance metric calibration issue

Decision: May increase w_{RSR} or investigate dataset quality

End of Appendix A

Continue to Appendix B: Legal and Regulatory Disclaimers for liability framework and compliance guidance.

11 Appendix B: Legal and Regulatory Disclaimers

The Synthetic Data Compliance Framework (SDCF) is provided as a technical and methodological reference implementing quantitative assessment concepts relevant to privacy engineering, data governance, and AI risk management. It does not constitute legal advice, regulatory determination, certification, or formal conformity assessment under any statutory or regulatory regime.

All regulatory mappings presented within this document—including to the GDPR, the EU AI Act, and international standards—represent technical interpretations intended to support internal governance, documentation, and risk analysis activities only. Final legal qualification of any dataset, processing activity, or AI system remains the sole responsibility of the relevant data controller, deployer, or regulatory authority.

Use of SDCF does not, by itself, establish compliance with the GDPR, the EU AI Act, ISO/IEC standards, or any national supervisory authority guidance. Organisations implementing this framework are strongly advised to obtain independent legal and regulatory review prior to any reliance on synthetic data for compliance-critical processing activities.

11.1 B.1 Synthetic Data and Anonymisation

SDCF does not assert that any synthetic dataset is, by default, anonymous under GDPR Recital 26. The framework instead supports structured technical risk analysis to inform anonymisation and pseudonymisation assessments. Regulatory qualification remains context-dependent and subject to supervisory interpretation.

11.2 B.1a EU AI Act Mapping

References to EU AI Act Article 10 within SDCF reflect an interpretive alignment with high-level data governance objectives. These references do not imply automatic satisfaction of EU AI Act conformity requirements and must be considered in conjunction with evolving harmonised standards, notified body guidance, and national implementing measures.

11.2.1 Critical Disclaimer

THE SYNTHETIC DATA COMPLIANCE FRAMEWORK (SDCF) AND ALL ASSESSMENTS CONDUCTED UNDER IT PROVIDE TECHNICAL ANALYSIS ONLY. SDCF IS NOT LEGAL ADVICE AND DOES NOT CREATE AN ATTORNEY-CLIENT RELATIONSHIP.

Organisations using SDCF must: - Obtain independent legal counsel for GDPR compliance determinations - Engage legal advisors for EU AI Act regulatory interpretation - Consult sector-specific legal experts (healthcare, financial services, etc.) - Make final compliance decisions with appropriate legal review

SDCF assessments provide evidence and technical analysis that can inform legal review, but they do not substitute for legal expertise.

11.2.2 What SDCF Does vs. What It Does Not Do

SDCF Provides	SDCF Does NOT Provide
Technical privacy risk quantification (PRS)	Legal determination that data is “anonymous” under GDPR
Statistical fidelity assessment (FI)	Legal opinion on compliance with regulations
Bias and fairness metrics (FV)	Guarantee of regulatory acceptance
Methodology for evidence collection	Legal advice on appropriate lawful basis
Structured documentation for auditors	Interpretation of local data protection law
Risk-based assessment framework	Representation in regulatory proceedings

11.2.3 Recommended Legal Review Points

Organisations should engage legal counsel at these critical junctures:

- 1. Purpose Definition (C1):** - Is the intended use of synthetic data lawful under GDPR? - What is the appropriate lawful basis (Article 6)? - Are Article 9 exemptions required (special categories)? - What contractual provisions are needed for external sharing?
- 2. GDPR Status Determination:** - Does SDCF assessment support “anonymous” classification? - Should data be treated as pseudonymous? - What additional safeguards are legally required? - How to document anonymisation efforts for accountability?
- 3. EU AI Act Classification:** - Is the AI system “high-risk” under Annex III? - What Article 10 obligations apply specifically? - Are transparency obligations (Article 13) triggered? - What documentation must be provided to notified bodies?
- 4. Cross-Border Data Transfer:** - Can synthetic data be transferred outside EEA? - Are transfer mechanisms (SCCs, adequacy decisions) required? - What additional safeguards must be in place? - How to document transfer risk assessment?
- 5. External Sharing Agreements:** - What contractual terms are needed? - Who is data controller vs. processor? - What liability allocation is appropriate? - How to handle data subject rights requests?

11.2.4 Liability Limitation

TO THE MAXIMUM EXTENT PERMITTED BY LAW:

- SDCF methodology is provided “AS IS” without warranties
- Assessors using SDCF assume responsibility for implementation quality
- Organisations using SDCF results assume risk of compliance decisions
- Wayne Kearns and Kaionix Labs disclaim liability for:
 - Regulatory enforcement actions
 - Data breaches or privacy incidents
 - Business losses from compliance failures
 - Damages from reliance on SDCF assessments

Organisations are solely responsible for: - Obtaining appropriate legal counsel - Making final compliance decisions - Implementing adequate safeguards - Responding to regulatory inquiries

11.3 B.2 GDPR Interpretation Guidance

11.3.1 The Anonymisation Question

Central issue: Is synthetic data “anonymous” (not subject to GDPR) or “personal data” (fully subject to GDPR)?

GDPR does not provide a clear answer. The determination is: - **Fact-specific:** Depends on dataset characteristics, generation method, context - **Risk-based:** Considers “reasonable means” for re-identification - **Dynamic:** Changes with technology evolution and adversary capability

11.3.2 Recital 26 Analysis

GDPR Recital 26 states: > “The principles of data protection should therefore not apply to anonymous information, namely information which does not relate to an identified or identifiable natural person or to personal data rendered anonymous in such a manner that the data subject is not or no longer identifiable. This Regulation does not therefore concern the processing of such anonymous information, including for statistical or research purposes.”

Key phrase: “not or no longer identifiable”

What counts as “identifiable”? > “To determine whether a natural person is identifiable, account should be taken of all the means reasonably likely to be used, such as singling out, either by the controller or by another person to identify the natural person directly or indirectly.”

“Reasonably likely” factors: - Costs of identification - Time required - Available technology at time of processing - Technological developments (ongoing assessment needed)

11.3.3 How SDCF Supports GDPR Analysis

SDCF provides technical evidence for the “reasonable means” test:

PRS < 20 (Low Risk): - Technical analysis suggests re-identification requires extraordinary effort - May support argument that identification is NOT “reasonably likely” - **Legal counsel must still evaluate:** Context, adversary capability, data sensitivity

PRS 20-50 (Moderate Risk): - Re-identification is possible but requires significant resources - Likely pseudonymous rather than anonymous - **Legal counsel must determine:** Is risk acceptable? What safeguards needed?

PRS > 50 (High Risk): - Re-identification is reasonably likely with available means - Should be treated as personal data (conservative approach) - **Legal counsel must assess:** Can additional controls reduce risk? Is synthetic approach appropriate?

CRITICAL: PRS is technical metric, not legal classification. The same PRS may support different legal conclusions depending on: - Data sensitivity (health data vs. purchasing patterns) - Adversary motivation (high-profile individuals vs. general population) - Disclosure context (internal use vs. public release) - Available auxiliary data (can adversary link to other datasets?)

11.3.4 EDPB Guidelines 01/2025 Implications

Key Takeaway from EDPB Pseudonymisation Guidelines (January 2025):

The boundary between anonymisation and pseudonymisation is **technical AND contextual**, not based on labels or intentions.

Relevance for Synthetic Data:

1. **Generation method matters:**
 - Simple sampling may retain identifiable records
 - Advanced methods (GANs, diffusion models) reduce but don't eliminate risk
 - Differential privacy provides formal guarantees but may sacrifice utility
2. **Assessment is ongoing:**
 - Technology evolves (new re-identification attacks developed)
 - Must reassess regularly, not one-time determination
 - SDCF validity periods reflect this (annual reassessment)
3. **Documentation burden:**
 - Controllers must demonstrate robustness of anonymisation
 - SDCF provides structured evidence trail (C2 Governance Record, C6 Transparency Pack)
4. **Conservative approach recommended:**
 - When in doubt, treat as pseudonymous (apply GDPR safeguards)
 - Especially for Article 9 special categories data
 - Reduces regulatory risk

11.3.5 Practical GDPR Compliance Posture

Option 1: Conservative Approach (Recommended)

Treat synthetic data as **pseudonymous personal data** regardless of PRS: - Maintain lawful basis for processing (Article 6) - Respect purpose limitation (Article 5(1)(b)) - Implement appropriate security (Article 32) - Maintain accountability documentation (Article 5(2)) - Conduct DPIA if high-risk (Article 35)

Advantages: - Regulatory safety (no risk of misclassifying as anonymous) - Maintains data subject rights - Simpler compliance posture

Disadvantages: - Foregoes some benefits of "anonymous" classification - Retains compliance burden

When appropriate: - Article 9 special categories data - High-profile individuals (politicians, celebrities) - External data sharing - Uncertain adversary capability

Option 2: Risk-Based Approach

Claim **anonymous** status if SDCF assessment supports: - PRS < 20 (low risk) - Gold Tier assessment (highest confidence) - Obtain legal opinion confirming anonymisation - Document rationale in C2 Governance Record

Advantages: - Reduces GDPR compliance burden - Enables broader data use - Facilitates open data initiatives

Disadvantages: - Regulatory risk if determination challenged - Requires strong legal opinion - Ongoing reassessment burden (technology evolves)

When appropriate: - Non-sensitive data (purchasing patterns, web analytics) - General population (not high-profile individuals) - Strong generation methodology (differential privacy, high noise) - Internal use (limited adversary access)

Option 3: Hybrid Approach

Different treatment based on use case: - **Internal analytics:** Treat as pseudonymous (safeguards in place) - **External research sharing:** Require Gold Tier + legal opinion supporting anonymisation - **Public open data:** Require PRS < 15 + independent audit

Document approach in organisational policy with clear criteria.

11.4 B.3 EU AI Act Article 10 Mapping

11.4.1 Article 10 Requirements

Article 10(2): Training, validation, and testing datasets shall be: - Relevant - Sufficiently representative

- To the best extent possible, free of errors - Appropriate in respect of the intended purpose

Article 10(3): Datasets shall: - Take into account geographical, contextual, behavioral, functional settings - Be subject to data governance and management practices

Article 10(4): Datasets shall be: - Examined for possible biases - Subject to bias detection, prevention, and mitigation measures

11.4.2 How SDCF Satisfies Article 10

Article 10 Requirement	SDCF Control	How It's Addressed
Relevant	C1 Purpose Sheet	Explicitly defines intended purpose; assessment is purpose-bounded
Representative	C4 Fidelity Testing (FI)	Distribution similarity and dependency preservation quantify representativeness
Free of errors	C4 Fidelity Testing	Quality assessment detects data errors, inconsistencies, impossibilities

Article 10 Requirement	SDCF Control	How It's Addressed
Appropriate for purpose	Overall conformance (SDCF-A/P/R)	Explicit fitness-for-purpose determination
Context considerations	C1 Purpose Sheet	Documents deployment context; tier selection considers operational environment
Data governance	C2 Governance Record + C7 Release Rules	Documented roles, decisions, controls, monitoring
Bias examination	C5 Fairness Assessment (FV)	Representation analysis and predictive parity testing
Bias mitigation	C5 + C2 Governance Record	Identifies bias; documents mitigation decisions

11.4.3 What SDCF Does NOT Cover

Out of Scope for SDCF: - AI system risk classification (high-risk determination under Annex III) - Full AI system risk management (Article 9) - Human oversight requirements (Article 14) - Transparency obligations for end users (Article 13) - Conformity assessment procedures (Article 43)

SDCF focuses exclusively on training data quality (Article 10). Organisations must address other AI Act requirements separately.

11.5 B.4 Scope Limitations

11.5.1 What SDCF Assesses

SDCF evaluates synthetic data outputs: - Privacy risk (re-identification, disclosure) - Statistical quality (distributions, correlations, utility) - Fairness (representation, parity)

In scope: The synthetic dataset itself

11.5.2 What SDCF Does NOT Assess

Generation Methodology: SDCF does not evaluate how synthetic data was generated.

Source Data Quality: SDCF assumes source data is of acceptable quality - if source has errors, synthetic inherits them.

Deployment Environment Security: SDCF assesses data risk, not operational security (access controls, encryption, monitoring must be implemented separately).

Model Risk Management: For ML training, SDCF validates training data quality but not model architecture, training procedures, deployment monitoring, or explainability.

Organisations must combine SDCF assessment with domain-specific review.

11.6 B.5 Liability Framework and Indemnification

Organisations using SDCF should NOT rely on indemnification from: - SDCF framework author (Wayne Kearns, Kaionix Labs) - Assessors conducting assessments - Tool providers (SDMetrics, mostlyai-qa, etc.)

Organisations must: - Carry appropriate insurance (professional liability, cyber, D&O) - Self-insure for regulatory risk - Obtain independent legal counsel - Make final compliance decisions - Implement adequate safeguards

TO THE MAXIMUM EXTENT PERMITTED BY LAW, WAYNE KEARNS AND KAIONIX LABS DISCLAIM ALL LIABILITY FOR REGULATORY ENFORCEMENT, DATA BREACHES, BUSINESS LOSSES, OR DAMAGES ARISING FROM SDCF USE.

11.7 B.6 Warranty Disclaimers

SDCF METHODOLOGY IS PROVIDED “AS IS” WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED.

SDCF thresholds remain provisional: Thresholds are subject to refinement based on empirical findings. Methodology will continue to evolve - Regulatory landscape evolving - Technology advancing

Organisations acknowledge: - SDCF is living framework (expect updates) - Provisional thresholds may be revised - Future regulatory guidance may supersede SDCF - Reassessment may be needed

11.8 B.7 Jurisdiction and Regulatory Change

SDCF focuses on EU regulatory context (GDPR, EU AI Act). For other jurisdictions, consult local counsel.

Regulatory landscape is dynamic - SDCF reflects interpretation as of December 2025. Organisations must monitor developments and trigger reassessment if guidance materially changes.

11.9 B.8 Acknowledgment

BY USING SDCF, ORGANISATIONS ACKNOWLEDGE: 1. SDCF provides technical analysis, not legal advice 2. Independent legal counsel is required 3. Liability for compliance decisions rests with organisation 4. No warranties or guarantees are provided 5. Ongoing monitoring of regulations is required

ORGANISATIONS THAT DO NOT AGREE WITH THESE TERMS SHOULD NOT USE SDCF.

End of Appendix B

Continue to Appendix C: Bronze Tier Guidance for detailed methodology when source data is unavailable.

12 Appendix C: Bronze Tier Guidance

12.1 C.1 Synthetic-Only Assessment Methodology

12.1.1 Why Bronze Tier Matters

The most common real-world scenario: Organisations need to validate third-party synthetic datasets without access to source data.

Examples: - Procuring synthetic data from vendors (Gretel, MOSTLY AI, Syntho) - Receiving synthetic datasets from research partners - Evaluating open-source synthetic training data (e.g., for AI model development) - Assessing legacy synthetic data where source no longer exists - Validating synthetic data where privacy regulations prohibit source access even for validation

Existing frameworks fail here: Most methodologies assume Gold Tier (full source access). When source unavailable, organisations face binary choice: - Use data blindly (risky) - Don't use data at all (misses opportunity)

Bronze Tier fills this gap: Provides structured assessment methodology for synthetic-only scenarios with honest acknowledgment of reduced confidence.

12.1.2 Bronze Tier Philosophy

Five Principles:

1. **Honest About Limitations**
 - Certificate explicitly states reduced confidence level
 - Wider confidence intervals
 - Conservative risk classification
2. **Conservative Risk Classification**
 - When uncertain, assume higher risk
 - Apply uncertainty penalties to scores
 - Recommend stronger controls
3. **Focus on Red Flags**
 - Identify clear problems (duplicates, impossibilities, extreme outliers)
 - Flag data quality issues that indicate generation failures
 - Surface concerns that warrant further investigation
4. **Domain Validation**
 - Assess plausibility against known real-world constraints
 - Use publicly available statistics (census, registries, benchmarks)
 - Apply subject matter expertise
5. **Enhanced Controls Required**
 - Bronze Tier assessment triggers stronger usage restrictions
 - Lower-risk use cases only (testing, development, preliminary analysis)
 - Require reassessment if repurposed for higher-risk use

12.1.3 When to Use Bronze Tier

Required: - No source data access possible - Third-party datasets (vendor, partner, open-source)
- Privacy prohibits source sharing - Legacy data (source no longer exists)

Appropriate: - Lower-risk use cases (software testing, prototyping, AI training pre-check) - Vendor evaluation and comparison - Preliminary screening (decide whether to invest in Gold/Silver) - Budget constraints (Bronze is faster and cheaper)

Not Appropriate: - High-risk AI systems requiring conformity assessment (EU AI Act) - Regulatory reporting requiring highest assurance - Critical business decisions (mergers, capital allocation) - External data sharing where liability is high - Medical device training data (MDR/IVDR compliance)

Decision Rule: If purpose demands Gold but only Bronze possible, either obtain source access OR reconsider synthetic data approach.

12.2 C.2 B-PRS: Privacy Risk Without Source Data

12.2.1 Component 1: B-MIR (Membership Inference Proxy)

Challenge: Cannot perform true membership inference without source data to train attack model.

Bronze Solution: Use outlier analysis as proxy for membership risk.

Rationale: Outliers and unusual records are: - More memorable (stand out in dataset) - More identifiable (unique combinations) - Higher re-identification risk if adversary has auxiliary data

Methodology:

1. Compute outlier scores:

```
from sklearn.neighbors import LocalOutlierFactor

lof = LocalOutlierFactor(n_neighbors=20, contamination=0.1)
outlier_scores = lof.fit_predict(synthetic_data)
outlier_probabilities = lof.negative_outlier_factor_

# Normalize to 0-1 scale (higher = more outlying)
outlier_normalized = (outlier_probabilities - min(outlier_probabilities)) / \
    (max(outlier_probabilities) - min(outlier_probabilities))
```

2. Calculate B-MIR:

```
# Percentage of records in top 10% outlier range
threshold_90th = percentile(outlier_normalized, 90)
high_outliers = count(outlier_normalized > threshold_90th)

B-MIR = (high_outliers / n_synthetic) * 100

# Conservative penalty: Assume outliers = higher membership risk
B-MIR_adjusted = B-MIR + 10 # Add 10-point uncertainty penalty
```

Interpretation: - B-MIR < 15: Low outlier concentration (suggests good distributional coverage) - B-MIR 15-30: Moderate outliers (acceptable for Bronze, document in certificate) - B-MIR > 30: High outlier concentration (red flag - investigate data quality)

Red Flags to Surface: - Records with impossible values (negative ages, future dates) - Extreme combinations unlikely in real world - Systematic patterns suggesting generation artifacts

12.2.2 Component 2: B-RSR (Record Similarity Proxy)

Challenge: Cannot compute distance to closest REAL record without source data.

Bronze Solution: Assess internal diversity and uniqueness within synthetic data.

Rationale: - Duplicate or near-duplicate synthetic records suggest overfitting or copying - Low internal diversity indicates narrow distributional coverage - Unique outliers may correspond to real individuals

Methodology:

1. Identify duplicates and near-duplicates:

```
from scipy.spatial.distance import pdist, squareform

# Compute pairwise distances
distances = pdist(synthetic_data, metric='gower') # Use appropriate metric
distance_matrix = squareform(distances)

# For each record, find distance to nearest neighbor (excluding self)
np.fill_diagonal(distance_matrix, np.inf) # Exclude self-comparison
nearest_distances = np.min(distance_matrix, axis=1)
```

2. Calculate B-RSR:

```
# Percentage with very low internal distance
near_duplicate_threshold = 0.05 # 5% of normalized distance range
near_duplicates = count(nearest_distances < near_duplicate_threshold)

B-RSR = (near_duplicates / n_synthetic) * 100

# Conservative penalty
B-RSR_adjusted = B-RSR + 15 # Larger penalty (we don't know source distances)
```

Interpretation: - B-RSR < 10: Good internal diversity (low duplication risk) - B-RSR 10-25: Moderate diversity (acceptable with disclosure) - B-RSR > 25: Poor diversity (red flag - potential overfitting)

Red Flags: - Exact duplicates (should never occur in properly generated synthetic data) - Clusters of nearly identical records - Suspiciously unique records far from any neighbors (may be outliers from source)

12.2.3 Component 3: B-ADR (Attribute Disclosure Proxy)

Challenge: Cannot test attribute disclosure without knowing source ground truth.

Bronze Solution: Assess correlation strength between quasi-identifiers and sensitive attributes.

Rationale: High correlation enables inference attacks if adversary has auxiliary data (knows quasi-identifiers, infers sensitive attributes).

Methodology:

1. Define quasi-identifiers and sensitive attributes:

```
quasi_identifiers = ['age', 'gender', 'zip_code', 'profession']
sensitive_attributes = ['income', 'health_condition', 'ethnicity']
```

2. Compute correlations:

```
for sensitive_attr in sensitive_attributes:
    correlations = []
    for qi in quasi_identifiers:
        if both_numeric:
            corr = pearson_correlation(synthetic_data[qi], synthetic_data[sensitive_attr])
        elif both_categorical:
            corr = cramers_v(synthetic_data[qi], synthetic_data[sensitive_attr])
        else:
            corr = point_biserial(...)

    correlations.append(abs(corr))

    max_correlation[sensitive_attr] = max(correlations)
```

3. Calculate B-ADR:

```
B-ADR = mean(max_correlation for all sensitive_attributes) × 100
```

```
# Conservative interpretation: High correlation = disclosure risk
# No penalty needed (already conservative by design)
```

Interpretation: - B-ADR < 20: Low correlation (limited disclosure risk) - B-ADR 20-40: Moderate correlation (some inference possible) - B-ADR > 40: High correlation (significant disclosure risk)

Red Flags: - Perfect or near-perfect correlations (suggests deterministic relationships) - Correlations that violate domain knowledge - Sensitive attributes easily predictable from QI combinations

12.2.4 Composite B-PRS

Formula:

$$\text{B-PRS} = (w_{\text{MIR}} \times \text{B-MIR}_{\text{adjusted}}) + (w_{\text{RSR}} \times \text{B-RSR}_{\text{adjusted}}) + (w_{\text{ADR}} \times \text{B-ADR}) + \text{uncertainty_penalty}$$

Default weights: $w_{\text{MIR}} = 0.4$, $w_{\text{RSR}} = 0.4$, $w_{\text{ADR}} = 0.2$

```
uncertainty_penalty = 20 # Conservative overall penalty for Bronze Tier
```

B-PRS **range:** 0-100 (lower = lower risk)

Typical Bronze Results: - B-PRS tends to be 20-30 points higher than equivalent Gold Tier PRS - Reflects genuine uncertainty about privacy risk without source comparison - Organisations should treat Bronze PRS conservatively

Certificate Language:

Bronze Tier Privacy Risk Score: B-PRS = 45 (95% CI: 30-60)

Note: Bronze Tier assessment conducted without source data access. Score reflects conservative risk estimate based on outlier analysis, internal diversity, and

correlation assessment. Actual privacy risk may be lower but cannot be verified without Gold Tier assessment. Use restricted to internal, lower-risk scenarios only.

12.3 C.3 B-FI: Fidelity Without Source Data

12.3.1 Component 1: B-DS (Distribution Validation)

Challenge: Cannot compare distributions to source without source data.

Bronze Solution: Validate distributions against known population characteristics and domain knowledge.

Methodology:

1. Compare to public statistics:

```
# Example: Age distribution
census_age_distribution = get_census_data(region, year)
synthetic_age_distribution = compute_distribution(synthetic_data['age'])

# Statistical test
chi2, p_value = chi2_test(synthetic_age_distribution, census_age_distribution)
similarity_score = 1 - (chi2 / (n_bins - 1))
```

2. Logical consistency checks:

```
violations = []

# Check for impossible values
if any(synthetic_data['age'] < 0) or any(synthetic_data['age'] > 120):
    violations.append("Impossible age values")

# Check for invalid dates
if any(synthetic_data['date_of_birth'] > today):
    violations.append("Future birth dates")

# Domain-specific rules
if healthcare_data:
    if any((diagnosis_age < 18) & (diagnosis == 'prostate_cancer')):
        violations.append("Biologically implausible diagnoses")
```

3. Calculate B-DS:

```
# Combine public data similarity and violation checks
public_similarity = mean(similarity_scores for all comparable variables)
violation_penalty = len(violations) * 5 # 5 points per violation

B-DS = (public_similarity * 100) - violation_penalty - uncertainty_penalty
uncertainty_penalty = 15 # Bronze doesn't know true source distribution

B-DS = max(0, B-DS) # Floor at 0
```


Interpretation: - B-DS > 70: Good alignment with known characteristics - B-DS 50-70: Moderate alignment (acceptable with limitations) - B-DS < 50: Poor alignment (red flag - data quality concerns)

Red Flags: - Distributions that violate known population characteristics - Impossible value combinations - Missing representation of known groups

12.3.2 Component 2: B-DP (Dependency Validation)

Challenge: Cannot compare correlations to source without source data.

Bronze Solution: Test whether expected domain relationships exist and are plausible.

Methodology:

1. Define expected relationships:

```
expected_relationships = [
    ('education_level', 'income', 'positive', 0.3, 0.6), # Expected range
    ('age', 'chronic_conditions', 'positive', 0.4, 0.7),
    ('exercise_frequency', 'bmi', 'negative', -0.2, -0.5),
]
```

2. Test relationships:

```
relationship_scores = []

for (var1, var2, expected_direction, min_corr, max_corr) in expected_relationships:
    actual_corr = correlation(synthetic_data[var1], synthetic_data[var2])

    # Check direction
    direction_correct = (expected_direction == 'positive' and actual_corr > 0) or \
        (expected_direction == 'negative' and actual_corr < 0)

    # Check magnitude plausibility
    magnitude_plausible = (min_corr <= abs(actual_corr) <= max_corr)

    if direction_correct and magnitude_plausible:
        relationship_scores.append(1.0)
    elif direction_correct:
        relationship_scores.append(0.5) # Right direction, wrong magnitude
    else:
        relationship_scores.append(0.0) # Wrong direction (red flag)
```

3. Calculate B-DP:

```
B-DP = mean(relationship_scores) × 100 - uncertainty_penalty
uncertainty_penalty = 15
```

Interpretation: - B-DP > 70: Expected relationships present and plausible - B-DP 50-70: Most relationships present (acceptable) - B-DP < 50: Missing or implausible relationships (quality concern)

Red Flags: - Expected relationships absent or reversed - Correlations that violate domain knowledge - Spurious correlations that shouldn't exist

12.3.3 Component 3: B-PU (Predictive Utility Estimation)

Challenge: Cannot benchmark against models trained on source data.

Bronze Solution: Train models on synthetic data and assess against domain expectations or public benchmarks.

Methodology:

1. Train models:

```
# Select relevant prediction task
X = synthetic_data[feature_columns]
y = synthetic_data[target]

# Cross-validation on synthetic data only
model = RandomForestClassifier()
cv_scores = cross_val_score(model, X, y, cv=5)
synthetic_performance = mean(cv_scores)
```

2. Compare to benchmarks:

```
# Option A: Published benchmark (if available)
if benchmark_exists:
    expected_performance = get_published_benchmark(task, dataset_type)
    utility_ratio = synthetic_performance / expected_performance

# Option B: Domain expert expectation
else:
    expert_expectation = get_expert_estimate(task)
    utility_ratio = synthetic_performance / expert_expectation

# Option C: Theoretical baseline
if no_benchmark:
    baseline = theoretical_baseline(task) # e.g., 50% for binary classification
    utility_ratio = (synthetic_performance - baseline) / (1.0 - baseline)
```

3. Calculate B-PU:

```
B-PU = utility_ratio × 100 - uncertainty_penalty
uncertainty_penalty = 15

# Conservative floor: Can't claim high utility without source comparison
B-PU = min(B-PU, 75) # Cap at 75 for Bronze
```

Interpretation: - B-PU > 65: Models perform reasonably well (acceptable utility) - B-PU 50-65: Moderate performance (limited utility) - B-PU < 50: Poor performance (insufficient utility)

Limitations: - Cannot detect overfitting to synthetic artifacts - Cannot verify generalization to real data - Benchmark comparison may not be apples-to-apples

12.3.4 Composite B-FI

Formula:

$$\text{B-FI} = (w_{\text{DS}} \times \text{B-DS}) + (w_{\text{DP}} \times \text{B-DP}) + (w_{\text{PU}} \times \text{B-PU}) - \text{uncertainty_penalty}$$

Default weights: $w_{\text{DS}} = 0.3$, $w_{\text{DP}} = 0.3$, $w_{\text{PU}} = 0.4$
uncertainty_penalty = 15

B-FI range: 0-100 (higher = higher fidelity)

Typical Bronze Results: - B-FI tends to be 15-25 points lower than equivalent Gold Tier FI - Reflects genuine uncertainty about fidelity without source comparison - Ceiling effect: B-PU capped at 75, limiting maximum B-FI

Certificate Language:

Bronze Tier Fidelity Index: B-FI = 62 (95% CI: 50-74)

Note: Bronze Tier assessment based on public data comparison, domain validation, and internal consistency checks. Cannot verify distributional accuracy against source data. Utility assessment based on cross-validation only (generalization to real data not verified). Use for non-critical analysis and testing only.

12.4 C.4 B-FV: Fairness Assessment

12.4.1 Good News: Fairness is More Tractable in Bronze

Representation analysis doesn't require source data - can compare to population statistics.

Limitation: Predictive parity testing still requires models and evaluation data (typically unavailable in Bronze).

12.4.2 B-RV: Representation Variance

Bronze uses same methodology as Gold/Silver:

1. Identify protected attributes:

```
protected_attributes = ['gender', 'race_ethnicity', 'age_group', 'disability']
```

2. Compare to population statistics:

```
for attribute in protected_attributes:
    synthetic_distribution = compute_distribution(synthetic_data[attribute])
    population_distribution = get_census_data(attribute, region)

    deviation = max(abs(synthetic_rate - population_rate) for all groups)
    max_deviations.append(deviation)
```

```
B-RV = max(max_deviations) * 100
```

This is actually reliable in Bronze - population statistics are publicly available and authoritative.

12.4.3 B-FV: No Predictive Parity

Bronze Tier cannot assess predictive parity without evaluation data.

Formula:

B-FV = B-RV *# Only representation component*

Certificate must state: "Predictive parity not assessed (Bronze Tier limitation). Organisations must conduct separate fairness testing if using for ML training or high-risk AI systems."

Compensating Measure: Organisations using Bronze for AI training should: - Conduct post-training fairness audits on models - Test model performance across protected groups - Implement ongoing bias monitoring

12.5 C.5 Bronze Tier Certificate Templates

12.5.1 Template 1: Bronze Tier SDCF-P Certificate

```
=====
SYNTHETIC DATA COMPLIANCE FRAMEWORK (SDCF)
BRONZE TIER ASSESSMENT CERTIFICATE
=====
```

DATASET INFORMATION

Dataset ID: SYNTH-VENDOR-2025-Q4
Dataset Name: Third-Party Transaction Dataset
Vendor/Source: ExampleVendor Inc.
Assessment Date: 2025-11-20
Assessor: Jane Smith, Senior Data Scientist, Organisation XYZ
Assessment Tier: Bronze (Synthetic Data Only - No Source Access)

PURPOSE STATEMENT

Intended Use: Software testing and development environment for payment processing application. Internal use only, development team access (12 authorised users).

CONFORMANCE LEVEL

SDCF-P (PROVISIONALLY APPROVED)

This synthetic dataset is PROVISIONALLY APPROVED for the stated purpose subject to the limitations and controls specified below.

ASSESSMENT RESULTS

Bronze Privacy Risk Score (B-PRS): 45 (95% CI: 30-60)
- B-MIR (Membership Proxy): 20

- B-RSR (Similarity Proxy): 35
- B-ADR (Disclosure Proxy): 25

Bronze Fidelity Index (B-FI): 62 (95% CI: 50-74)

- B-DS (Distribution Validation): 68
- B-DP (Dependency Validation): 65
- B-PU (Utility Estimation): 55

Bronze Fairness Variance (B-FV): 18 (95% CI: 12-24)

- B-RV (Representation Variance): 18
- Predictive Parity: Not Assessed (Bronze Limitation)

INTERPRETATION

Privacy: MODERATE RISK (B-PRS = 45)

- Conservative estimate due to Bronze Tier limitations
- Internal diversity acceptable but uncertainty about real privacy risk
- Restrict to internal use only

Fidelity: MODERATE (B-FI = 62)

- Distributions align with public data
- Domain relationships present and plausible
- Utility adequate for testing purposes
- Cannot verify accuracy for production decisions

Fairness: MODERATE CONCERNS (B-FV = 18)

- Representation adequate across protected groups
- Predictive parity not assessable (Bronze Tier limitation)
- Post-deployment monitoring recommended if used for ML

LIMITATIONS (Bronze Tier Specific)

1. Assessment conducted WITHOUT access to source data
2. Privacy risk estimates CONSERVATIVE (may overstate actual risk)
3. Fidelity assessment based on public data comparison only
4. Cannot verify generalization to real-world scenarios
5. Predictive parity NOT assessed (requires separate testing)
6. Confidence intervals WIDER than Gold/Silver Tier
7. Suitable for LOWER-RISK use cases only

REQUIRED CONTROLS

1. Use RESTRICTED to stated purpose (software testing/development)
2. Access LIMITED to development team (12 named users)
3. NO production deployment without Gold Tier reassessment
4. NO external sharing or publication
5. Enhanced audit logging required
6. Quarterly usage review

7. Reassessment required if purpose changes

REASSESSMENT TRIGGERS

- Purpose evolution beyond testing/development
- Production deployment consideration
- External sharing request
- Security incident involving dataset
- Regulatory guidance changes
- 12 months from assessment date (whichever occurs first)

VALIDITY PERIOD

Valid until: 2026-05-20 (6 months) or upon material change

ASSESSOR DECLARATION

I certify that this Bronze Tier assessment was conducted in accordance with SDCF v1.1 methodology to the best of my ability. This assessment provides technical analysis only and does not constitute legal advice. Organisations must obtain independent legal counsel for compliance determinations.

Signed: _____ Date: 2025-11-20
Jane Smith, Senior Data Scientist

ORGANISATION ACKNOWLEDGMENT

Organisation acknowledges receipt of this certificate and associated limitations. Organisation agrees to implement required controls and restrictions. Organisation accepts responsibility for compliance decisions and use of assessed data.

Authorised Signatory: _____ Date: _____
[Name, Title, Organisation XYZ]

For questions or concerns, contact: data-governance@example.com
SDCF Framework: <https://www.kaionix.com/kaionix-labs>

12.5.2 Template 2: Bronze Tier SDCF-R Certificate

SYNTHETIC DATA COMPLIANCE FRAMEWORK (SDCF)
BRONZE TIER ASSESSMENT CERTIFICATE

DATASET INFORMATION

Dataset ID: SYNTH-LEGACY-2024

Dataset Name: Legacy Customer Segmentation Data

Source: Internal generation (2024), source no longer available

Assessment Date: 2025-11-20

Assessor: John Doe, Data Governance Lead

Assessment Tier: Bronze (Synthetic Data Only - Source Data Lost)

PURPOSE STATEMENT

Intended Use: Customer analytics for marketing campaign targeting (external partners will access insights)

CONFORMANCE LEVEL

SDCF-R (RESTRICTED)

This synthetic dataset is RESTRICTED for the stated purpose. The dataset does NOT meet requirements for the intended use.

ASSESSMENT RESULTS

Bronze Privacy Risk Score (B-PRS): 68 (95% CI: 55-81)

- B-MIR (Membership Proxy): 45 HIGH
- B-RSR (Similarity Proxy): 52 HIGH
- B-ADR (Disclosure Proxy): 38

Bronze Fidelity Index (B-FI): 48 (95% CI: 35-61)

- B-DS (Distribution Validation): 52
- B-DP (Dependency Validation): 45
- B-PU (Utility Estimation): 47

Bronze Fairness Variance (B-FV): 32 (95% CI: 25-39)

- B-RV (Representation Variance): 32 SIGNIFICANT

DETERMINATION

This dataset is UNSUITABLE for the stated purpose (customer analytics with external partner access) due to:

1. HIGH PRIVACY RISK (B-PRS = 68)
 - Outlier concentration suggests identifiable records
 - Low internal diversity (duplication concerns)
 - Unacceptable for external disclosure
2. INSUFFICIENT FIDELITY (B-FI = 48)
 - Distributional misalignment with census data
 - Key correlations missing or implausible
 - Would produce misleading business insights

3. SIGNIFICANT FAIRNESS CONCERNS (B-FV = 32)
- Underrepresentation of minority groups (>30% deviation)
 - Could lead to discriminatory campaign targeting

IDENTIFIED RED FLAGS

Duplicate records detected (0.8% exact duplicates)
Impossible value combinations (e.g., 25-year-old retirees)
Missing representation: LGBTQ+ individuals severely underrepresented
Income distribution inconsistent with regional census
Age-income correlation implausibly weak

RECOMMENDATION

DO NOT USE this dataset for the stated purpose.

ALTERNATIVE OPTIONS:

1. Regenerate synthetic data with improved methodology and Gold Tier assessment
2. Use for internal-only, non-critical testing purposes (SDCF-P possible with strict controls)
3. Abandon synthetic approach and use aggregated statistics instead

IF PROCEEDING WITH ALTERNATIVE USE:

- Restrict to internal testing only (NO external access)
- Document known limitations prominently
- Enhanced monitoring and bias checks
- Quarterly review of usage and incidents

CERTIFICATE ISSUED AS EVIDENCE OF DUE DILIGENCE

This SDCF-R certificate demonstrates that Organisation conducted appropriate assessment and made informed decision NOT to use dataset for intended purpose.

Assessor: _____ Date: 2025-11-20
John Doe, Data Governance Lead

=====

For questions: data-governance@example.com
SDCF Framework: <https://www.kaionix.com/kaionix-labs>

=====

12.6 C.6 Risk Statement Examples

12.6.1 Risk Statement: Bronze Tier SDCF-P

RISK STATEMENT

Synthetic Dataset: SYNTH-VENDOR-2025-Q4

Conformance: SDCF-P (Provisionally Approved)
Assessment Tier: Bronze

SUMMARY OF RESIDUAL RISKS

This dataset has been assessed using SDCF Bronze Tier methodology (synthetic data only, no source access). The following residual risks remain:

1. PRIVACY RISK (B-PRS = 45 - Moderate)
 - Re-identification risk cannot be fully quantified without source comparison
 - Conservative estimates suggest moderate risk for internal use
 - External disclosure NOT RECOMMENDED
2. UTILITY RISK (B-FI = 62 - Moderate)
 - Fidelity adequate for testing but not validated for production decisions
 - Cannot confirm accuracy of business insights
 - Use for critical decisions NOT RECOMMENDED
3. FAIRNESS RISK (B-FV = 18 - Moderate)
 - Representation acceptable but predictive parity not assessed
 - Models trained on this data require separate bias testing
 - High-risk AI use NOT RECOMMENDED

RISK ACCEPTANCE

Organisation XYZ accepts these residual risks for the stated purpose (software testing and development) subject to controls specified in certificate.

Risk Owner: [Name], [Title]

Date: [Date]

Signature: -----

ESCALATION CRITERIA

Escalate to governance board if:

- Purpose evolves beyond testing/development
- External sharing requested
- Privacy incident occurs
- Bias detected in downstream applications

12.7 C.7 Bronze Tier Case Studies

12.7.1 Case Study 1: AI Training Data Pre-Check (Success)

Scenario:

Irish fintech startup evaluating synthetic transaction data from vendor for fraud detection model training.

Context: - Vendor provides sample dataset (100k records) - No source data access (vendor propri-

etary) - Need rapid assessment before procurement decision - Budget: €10k for assessment

Bronze Tier Assessment: - **Timeframe:** 1 week - **B-PRS:** 38 (moderate risk, acceptable for internal ML training) - **B-FI:** 71 (good alignment with known fraud patterns) - **B-FV:** 14 (excellent representation across merchant categories) - **Conformance:** SDCF-P

Outcome: - Procurement approved based on Bronze assessment - Data used for initial model development - Performance validated on hold-out real data (separate test) - Model achieved 89% of baseline performance (acceptable) - **Decision:** Procure larger dataset, continue Bronze monitoring

Lesson: Bronze Tier effectively screened vendor data quality before investment.

12.7.2 Case Study 2: Legacy Data Retirement (Failure Detected)

Scenario:

Large insurance company discovered legacy synthetic claims data (2021 generation) still in use. Source data no longer accessible.

Context: - Data used for internal pricing model validation - No documentation of original assessment - Compliance audit requested validation evidence - Bronze Tier assessment conducted retrospectively

Bronze Tier Assessment: - **Timeframe:** 2 weeks - **B-PRS:** 72 (high risk - concerning) - **B-FI:** 41 (insufficient - red flags detected) - **B-FV:** 38 (significant fairness issues) - **Conformance:** SDCF-R

Red Flags Identified: - Duplicate records (>5%) - Claim amounts didn't align with known industry distributions - Protected group representation off by >35% from census - Age-claim severity relationship implausible

Outcome: - Data immediately retired from use - Pricing models re-validated on alternative data - SDCF-R certificate provided to auditors as evidence of due diligence - New data governance policy: No synthetic data without assessment

Lesson: Bronze Tier identified serious quality issues that Gold Tier (unavailable) would have caught earlier. Late discovery costly but certificate showed accountability.

12.7.3 Case Study 3: Open-Source AI Training Dataset (Qualified Approval)

Scenario:

Healthcare AI research team evaluating open-source synthetic patient dataset (released by academic consortium) for diabetic retinopathy detection model.

Context: - Dataset publicly available (25k synthetic fundus image descriptions) - No source patient data (privacy-preserving release) - Research ethics board requires quality assessment - Timeline: 3 weeks to assessment

Bronze Tier Assessment: - **Timeframe:** 10 days - **B-PRS:** 42 (moderate risk, acceptable for research) - **B-FI:** 58 (limited - concerns about rare pathology representation) - **B-FV:** 22 (moderate - some demographic underrepresentation) - **Conformance:** SDCF-P (with significant limitations)

Key Findings: - Privacy adequate for research use - Fidelity concerns: Rare disease stages underrepresented - Fairness concerns: Asian and African descent patients underrepresented (15% deviation)

Mitigation: - Use for preliminary model development only (not validation) - Supplement with stratified sampling from licensed clinical dataset - Document limitations in research publication - Include bias analysis in model evaluation

Outcome: - Research proceeded with documented limitations - Model performance validated separately on diverse clinical data - Publication acknowledged synthetic data limitations - Framework influenced by Bronze Tier assessment

Lesson: Bronze Tier enabled informed use with appropriate caveats and mitigations.

End of Appendix C

Continue to Appendix D: Regulatory Mapping Tables for detailed control alignment with standards.

12.8 C.8 Observed Performance Ranges (Bronze Tier)

This section summarises empirically observed metric ranges from the Version 1.95 Bronze Tier retrospective evaluation. These values should be interpreted as indicative reference ranges derived from a limited, heterogeneous dataset portfolio, rather than as normative regulatory thresholds.

12.8.1 B-PRS (Privacy Risk Score)

Observed B-PRS values across the evaluated datasets range from 0.09 to 0.79. Datasets classified as SDCEP-P typically exhibit B-PRS values below 0.50, while SDCEP-R classifications dominate above this region. No dataset in the evaluated portfolio satisfies the provisional criteria for unrestricted approval. These values reflect conservative bias under source-data-absent conditions.

12.8.2 B-FI (Fidelity Index)

Observed B-FI values consistently exceed 0.92 across all evaluated datasets. This suggests strong preservation of marginal distributions, inter-variable dependencies, and proxy predictive utility under the selected assessment configuration. These results are informative but not statistically generalisable beyond the evaluated sample.

12.8.3 B-FV (Fairness Variance)

Where protected attribute information is available, observed B-FV values remain within provisional tolerance bands. Because Bronze Tier fairness evaluation lacks access to labelled outcomes, these measurements primarily reflect representation stability rather than predictive parity.

12.8.4 Conformance Distribution

Across the evaluated portfolio, 60% of datasets receive SDCEP-R classification and 40% receive SDCEP-P classification. No dataset receives SDCEP-A classification under Bronze Tier. This distribution reflects the deliberate conservatism of the Bronze Tier design.

12.8.5 Interpretive Guidance

These observed performance ranges are provided for practitioner orientation only. They do not constitute safety guarantees, regulatory determinations, or conformance certifications beyond the scope of the specific assessment configuration applied.

12.9 C.9 Validated Synthesis Method Benchmarks

Based on comparative assessment of 5 synthesis methods across 10 datasets, practitioners can use the following evidence-based benchmarks for method selection:

12.9.1 For Demographic/Census Data

Recommendation: TVAE (Tabular Variational Autoencoder)

Controlled comparison on identical source data (UCI Adult Income, 32,561 records, 15 columns) demonstrates:

Method	B-PRS	B-FI	B-FV	Conformance
TVAE (RECOMMENDED)	0.534	0.994	0.724	SDCF-R
CTGAN	0.637	0.998	0.548	SDCF-R
GaussianCopula	0.639	0.999	0.693	SDCF-R
TVAE Advantage	16–20% lower	Maintained	—	—

Key Finding: TVAE achieves 16–20% lower B-PRS (privacy risk) compared to CTGAN or GaussianCopula while maintaining excellent fidelity (B-FI > 0.99). All three methods trigger Restricted conformance due to B-FV > 0.30 (inherent demographic complexity from source data), but TVAE offers best privacy-quality tradeoff.

Availability: TVAE available in SDV (Synthetic Data Vault) open-source library: `pip install sdv`

Practitioner Recommendation: Organisations synthesizing demographic or census data should prioritise TVAE. For critical applications, conduct Bronze Tier comparison of 2–3 methods before committing to production synthesis pipeline.

12.9.2 For Commercial/Transactional Data

Recommendation: Commercial GANs (MostlyAI, Gretel)

Validation results for e-commerce/business data:

Dataset	Vendor	B-PRS	B-FI	B-FV	Conform
CDNOW Purchases	MostlyAI	0.360	0.999	0.100	SDCF-A
Census SynLBD	US Census	0.360	1.000	0.100	SDCF-A
CMS SynPUF	CMS	0.390	0.997	0.100	SDCF-A

Key Finding: Commercial vendors and government synthesis achieve:

- B-PRS 0.36–0.39 (Moderate risk, lowest observed for structured data)
- B-FI 0.997–1.000 (near-perfect fidelity)

- B-FV ≈ 0.10 (baseline, minimal fairness concerns)
- 100% Acceptable conformance (suitable for intended use cases)

Practitioner Recommendation: For transactional data (purchases, claims, establishment records), commercial GANs demonstrate excellent Bronze Tier performance. Vendor tuning optimised for business use cases. Strong support and documentation reduce implementation risk.

12.9.3 For AI Training/Code Data

Recommendation: Case-by-case evaluation (no general guidance)

Validation results for LLM-generated AI/Code data show extreme variability:

Dataset	Content Type	B-PRS	B-FI	Conform
Jupyter Agent	Code Q&A (generic)	0.090	0.999	SDCF-A
Gretel Safety	Safety scenarios (unique)	0.789	0.999	SDCF-R
PLEIAs SYNTH	Reasoning (diverse)	0.777	0.927	SDCF-R
Range	—	0.09–0.79	0.93–1.00	Mixed

Key Finding: LLM-generated content exhibits 8.8x B-PRS range (widest variability in portfolio). Pattern:

- **Generic code examples:** Low B-PRS (0.090), Acceptable conformance
- **Unique training data:** Critical B-PRS (0.777–0.789), Restricted conformance
- Content uniqueness drives risk score; no method-level generalization possible

Practitioner Recommendation: For AI training or code datasets:

1. Conduct Bronze Tier assessment on candidate synthesis approach **before** generating large-scale dataset
2. Compare 2–3 LLM synthesis strategies if possible (different prompts, models, temperatures)
3. Expect high variability; B-PRS 0.09–0.79 range observed in validation
4. Small datasets (< 1000 records) may show artificially low B-PRS due to high intrinsic variability — Require human expert review regardless of score (Appendix C.2)

12.9.4 Cross-Method Insights

Method Ranking by Average B-PRS (Lower = Better Privacy):

1. LLM-generated (content-dependent): 0.09–0.79 (highly variable)
2. Commercial GANs: 0.36–0.63 (consistent, lower risk)
3. TVAE (deep learning): 0.534 (best for demographic)
4. CTGAN (deep learning): 0.637 (good, but TVAE better)
5. GaussianCopula (statistical): 0.639 (good, simple implementation)

Fidelity Consistency: All methods achieve B-FI > 0.92 (modern tools maintain excellent fidelity regardless of method). Method selection should prioritise privacy-fairness tradeoff, not fidelity (fidelity is consistently high).

General Practitioner Guidance:

- **Default recommendation:** Start with TVAE for tabular data (best evidence-based performance)

- **For critical applications:** Compare 2–3 methods using Bronze Tier assessment (~ 1 hour per method) before production synthesis
- **Method matters:** Observed 8.8x B-PRS range across methods — Selection significantly impacts privacy-quality tradeoff
- **Domain-specific:** Demographic data benefits from TVAE; transactional data benefits from commercial GANs; AI/Code requires case-by-case evaluation

13 Appendix D: Regulatory Mapping Tables

This appendix provides detailed mappings between SDCF controls and regulatory requirements, enabling organisations to demonstrate compliance through SDCF assessment.

13.1 D.1 SDCF \rightarrow GDPR Article Mapping

13.1.1 Table D.1.1: GDPR Principles (Article 5)

GDPR Principle	SDCF Control	How SDCF Addresses
Article 5(1)(a): Lawfulness, fairness, transparency	C1 Purpose Sheet, C6 Transparency Pack	Purpose Sheet documents lawful basis consideration; Transparency Pack provides disclosure to data subjects
Article 5(1)(b): Purpose limitation	C1 Purpose Sheet, C7 Release Rules	Purpose explicitly defined and documented; Release Rules prevent scope creep
Article 5(1)(c): Data minimization	C3 Privacy Risk Testing, C7 Release Rules	Privacy Risk Score quantifies that only necessary data characteristics retained; access controls limit disclosure
Article 5(1)(d): Accuracy	C4 Fidelity Testing	Fidelity Index validates statistical accuracy and quality
Article 5(1)(e): Storage limitation	C7 Release Rules	Retention period and deletion procedures specified
Article 5(1)(f): Integrity and confidentiality	C3 Privacy Risk Testing, C7 Release Rules	Privacy Risk Score quantifies security level; Release Rules define protective measures
Article 5(2): Accountability	C2 Governance Record, C6 Transparency Pack	Complete documentation trail demonstrating compliance efforts

13.1.2 Table D.1.2: GDPR Lawful Basis (Article 6)

Lawful Basis	SDCF Support	Evidence Provided
6(1)(a): Consent	C1 Purpose Sheet, C2 Governance Record	Documents consent mechanism if applicable; not primary basis for synthetic data

Lawful Basis	SDCF Support	Evidence Provided
6(1)(b): Contract	C1 Purpose Sheet	Documents contractual necessity if synthetic data used for contract performance
6(1)(c): Legal obligation	C1 Purpose Sheet, C6 Transparency Pack	Maps to legal obligations requiring data processing
6(1)(d): Vital interests	C1 Purpose Sheet	Documents vital interest basis if applicable (rare for synthetic data)
6(1)(e): Public interest	C1 Purpose Sheet, C2 Governance Record	Documents public interest task (common for research, public health)
6(1)(f): Legitimate interests	C1 Purpose Sheet, C2 Governance Record, C3 Privacy Risk Testing	Balancing test documentation; PRS demonstrates minimised privacy impact

13.1.3 Table D.1.3: GDPR Special Categories (Article 9)

Requirement	SDCF Control	How SDCF Addresses
Article 9(1): Prohibition	C1 Purpose Sheet	Identifies if special category data involved
Article 9(2)(a): Explicit consent	C1 Purpose Sheet, C2 Governance Record	Documents consent mechanism if used
Article 9(2)(g): Substantial public interest	C1 Purpose Sheet	Documents public interest basis (common for health research)
Article 9(2)(i): Public health	C1 Purpose Sheet	Documents public health basis (pandemic response, disease surveillance)
Article 9(2)(j): Research	C1 Purpose Sheet, C3 Privacy Risk Testing	Documents research purpose; PRS demonstrates appropriate safeguards
Article 9(3): Suitable safeguards	C3 Privacy Risk Testing, C7 Release Rules	PRS < 20 demonstrates strong privacy protection; controls specified

13.1.4 Table D.1.4: GDPR Security and Accountability

GDPR Requirement	SDCF Control	Compliance Evidence
Article 25: Data protection by design	Entire SDCF framework	Synthetic data generation = privacy-by-design measure; SDCF validates effectiveness
Article 32: Security of processing	C3 Privacy Risk Testing, C7 Release Rules	PRS quantifies residual risk; Release Rules define technical/organisational measures
Article 35: DPIA	C1 Purpose Sheet, C2 Governance Record, C3/C4/C5 Testing	SDCF assessment informs DPIA risk analysis; may reduce DPIA scope if anonymisation successful

GDPR Requirement	SDCF Control	Compliance Evidence
Article 30: Records of processing	C2 Governance Record	Documents processing activities for synthetic data

13.1.5 Table D.1.5: Anonymisation vs. Pseudonymization

Assessment	GDPR Classification	SDCF Evidence
PRS < 15, Gold Tier	Potentially anonymous (requires legal opinion)	Strong technical evidence for anonymisation claim
PRS 15-25, Gold Tier	Likely pseudonymous	Moderate privacy risk; treat as personal data with reduced risk
PRS > 25	Pseudonymous personal data	Treat as personal data; GDPR obligations apply
Bronze Tier, any PRS	Treat as pseudonymous (conservative)	Uncertainty precludes anonymisation claim

Critical Note: SDCF provides technical evidence only. Legal classification requires independent legal counsel considering organisational context, deployment scenario, and risk tolerance.

13.2 D.2 SDCF → EU AI Act Mapping

13.2.1 Table D.2.1: Article 10 (Training Data)

Article 10 Requirement	SDCF Control	Compliance Evidence
10(2): Relevant	C1 Purpose Sheet	Purpose-bounded assessment ensures relevance to specific AI system use
10(2): Sufficiently representative	C4 Fidelity Testing (FI)	Distribution similarity and dependency preservation quantify representativeness
10(2): Free of errors	C4 Fidelity Testing	Quality assessment detects data errors, inconsistencies, impossibilities
10(2): Appropriate for purpose	Overall conformance (SDCF-A/P/R)	Explicit fitness-for-purpose determination
10(3): Geographical/contextual setting	C1 Purpose Sheet	Documents deployment context and environmental factors
10(3): Data governance practices	C2 Governance Record, C7 Release Rules	Documented roles, decisions, controls, monitoring procedures
10(4): Examine for biases	C5 Fairness Assessment (FV)	Representation analysis and predictive parity testing
10(4): Bias detection	C5 Fairness Assessment	Quantifies representation variance and parity violations

Article 10 Requirement	SDCF Control	Compliance Evidence
10(4): Bias prevention	C2 Governance Record	Documents generation methodology choices to prevent bias
10(4): Bias mitigation	C2 Governance Record, C5 Fairness Assessment	Documents mitigation measures; validates effectiveness

13.2.2 Table D.2.2: High-Risk AI System Requirements

AI Act Requirement	SDCF Control	How SDCF Supports
Article 9: Risk management	C1 Purpose Sheet, C2 Governance Record	Training data risk assessment integrated into AI system risk management
Article 11: Technical documentation	C6 Transparency Pack	SDCF certificate and transparency pack = technical documentation
Article 12: Record-keeping	C2 Governance Record, C7 Release Rules	Audit logs and governance decisions documented
Article 13: Transparency	C6 Transparency Pack	Dataset characteristics, limitations, and usage restrictions documented
Article 17: Quality management	C2 Governance Record, C7 Release Rules	Data governance procedures defined and monitored
Article 61: Post-market monitoring	C7 Release Rules	Ongoing monitoring and reassessment triggers defined

13.2.3 Table D.2.3: Conformity Assessment Evidence

For notified body review of high-risk AI systems:

Evidence Required	SDCF Deliverable
Training data quality assessment	SDCF Certificate with PRS/FI/FV scores
Data governance procedures	C2 Governance Record
Bias testing methodology	Appendix A (mathematical definitions)
Bias testing results	C5 Fairness Assessment results
Data management practices	C7 Release Rules
Technical documentation	C6 Transparency Pack (complete)
Audit trail	C2 Governance Record with timestamps

13.3 D.3 SDCF → ISO/IEC Standards Mapping

13.3.1 Table D.3.1: ISO/IEC 27001:2022 (Information Security)

ISO 27001 Control	SDCF Control	Mapping
A.5.1: Policies	C2 Governance Record	Synthetic data governance policies
A.5.7: Threat intelligence	C3 Privacy Risk Testing	Privacy attack modeling and testing
A.5.9: Inventory of assets	C1 Purpose Sheet, C6 Transparency Pack	Synthetic data assets documented
A.5.10: Acceptable use	C7 Release Rules	Usage restrictions and acceptable use policy
A.5.12: Classification	C1 Purpose Sheet, C7 Release Rules	Data classification based on sensitivity and conformance level
A.5.13: Labeling	C6 Transparency Pack	Dataset labeled with conformance level and restrictions
A.5.15: Access control	C7 Release Rules	Access control requirements specified
A.5.23: Cloud services	C7 Release Rules	Handling requirements for cloud storage/processing
A.8.2: Privileged access rights	C7 Release Rules	Authorisation and approval processes
A.8.3: Information access restriction	C7 Release Rules	Need-to-know access principles
A.8.10: Information deletion	C7 Release Rules	Secure deletion procedures and retention periods
A.8.11: Data masking	C3 Privacy Risk Testing	Synthetic data as masking technique; SDCF validates effectiveness
A.8.24: Cryptography	C7 Release Rules	Encryption requirements (at rest, in transit)

13.3.2 Table D.3.2: ISO/IEC 27701:2019 (Privacy)

ISO 27701 Control	SDCF Control	Mapping
5.2.1: PII processing principles	C1 Purpose Sheet	Purpose limitation, data minimization documented
6.2.1: Conditions for collection	C1 Purpose Sheet, C2 Governance Record	Lawful basis documented
6.7.2.2: Basis for PII transfer	C3 Privacy Risk Testing, C7 Release Rules	Risk assessment for transfers; safeguards specified
7.2.2: PII de-identification	C3 Privacy Risk Testing	PRS quantifies de-identification effectiveness
7.3.2: PII minimization	C3 Privacy Risk Testing	Privacy risk assessment validates minimization
7.4.7: Automated decision-making	C5 Fairness Assessment	Bias testing for AI/ML use cases
8.2.4: Capability to comply with PII principal rights	C6 Transparency Pack	Transparency enables data subject awareness

13.3.3 Table D.3.3: ISO/IEC 23894:2023 (AI Risk Management)

ISO 23894 Requirement	SDCF Control	Mapping
6.2: Data quality	C4 Fidelity Testing	Fidelity Index operationalises data quality assessment
6.3: Fairness	C5 Fairness Assessment	Fairness Variance quantifies bias risk
6.4: Transparency	C6 Transparency Pack	Dataset characteristics, limitations documented
6.5: Accountability	C2 Governance Record	Decision-making documented
7.1: Risk identification	C3/C4/C5 Testing	Privacy, fidelity, fairness risks identified
7.2: Risk analysis	C3/C4/C5 Testing	Quantitative risk scores (PRS, FI, FV)

ISO 23894 Requirement	SDCF Control	Mapping
7.3: Risk evaluation	C2 Governance Record	Risk acceptance decisions documented
7.4: Risk treatment	C7 Release Rules	Controls implemented to mitigate risks

13.3.4 Table D.3.4: ISO/IEC 42001:2023 (AI Management System)

ISO 42001 Requirement	SDCF Control	Mapping
4.1: Understanding context	C1 Purpose Sheet	Deployment context documented
5.2: AI policy	C2 Governance Record	Synthetic data governance policy
6.1: Risk management	C3/C4/C5 Testing, C2 Governance Record	AI system risk analysis informed by training data assessment
7.2: Competence	C2 Governance Record	Assessor qualifications documented
7.5: Documented information	C6 Transparency Pack	Complete documentation bundle
8.2: AI system development	C1 Purpose Sheet, C4/C5 Testing	Training data quality validated
9.1: Monitoring	C7 Release Rules	Ongoing monitoring procedures
9.2: Internal audit	C6 Transparency Pack, C2 Governance Record	Audit-ready documentation
10.1: Continual improvement	C7 Release Rules	Reassessment triggers for continuous improvement

13.4 D.4 SDCF → NIST Frameworks Mapping

13.4.1 Table D.4.1: NIST Privacy Framework

NIST Function	NIST Category	SDCF Control	Mapping
IDENTIFY-P	Inventory-P	C1 Purpose Sheet	Synthetic data assets inventoried
IDENTIFY-P	Risk Assessment-P	C3 Privacy Risk Testing	Privacy risks identified and quantified
GOVERN-P	Policies-P	C2 Governance Record	Privacy governance policies
GOVERN-P	Risk Management-P	C2 Governance Record	Privacy risk decisions documented
CONTROL-P	Data Processing-P	C3 Privacy Risk Testing, C7 Release Rules	Technical measures implemented
COMMUNICATE-P	Awareness-P	C6 Transparency Pack	Stakeholders informed of privacy practices
PROTECT-P	Data Security-P	C7 Release Rules	Security controls specified

13.4.2 Table D.4.2: NIST AI Risk Management Framework

NIST RMF Function	NIST Category	SDCF Control	Mapping
GOVERN	Accountable	C2 Governance Record	Roles and responsibilities defined
GOVERN	Transparent	C6 Transparency Pack	Dataset documentation and disclosure
GOVERN	Risk Management	C2 Governance Record	Training data risk decisions
MAP	Context	C1 Purpose Sheet	AI system context documented
MAP	Categorise	C1 Purpose Sheet	Training data characteristics categorised
MAP	Risk Assessment	C3/C4/C5 Testing	Privacy, utility, fairness risks assessed

NIST RMF Function	NIST Category	SDCF Control	Mapping
MEASURE	Metrics	C3/C4/C5 Testing	Quantitative metrics (PRS, FI, FV)
MEASURE	Validation	C4/C5 Testing	Fidelity and fairness validated
MANAGE	Risk Treatment	C7 Release Rules	Controls implemented
MANAGE	Monitoring	C7 Release Rules	Ongoing monitoring specified

13.5 D.5 Sector-Specific Regulatory Mapping

13.5.1 Table D.5.1: Healthcare (GDPR Article 9, MDR/IVDR)

Regulation	Requirement	SDCF Control	Compliance Evidence
GDPR Article 9	Special category safeguards	C3 Privacy Risk Testing (PRS < 20)	Strong privacy protection for health data
GDPR Article 89	Research exemptions with safeguards	C1 Purpose Sheet, C3 Privacy Risk Testing	Research purpose + safeguards documented
MDR Article 61	Clinical evaluation	C4 Fidelity Testing	Training data quality for medical device AI
IVDR Article 56	Performance evaluation	C4 Fidelity Testing	Training data representativeness validated
National Health Laws	Varies by Member State	C1 Purpose Sheet, C2 Governance Record	Legal basis documented per jurisdiction

13.5.2 Table D.5.2: Financial Services (Basel III, MiFID II)

Regulation	Requirement	SDCF Control	Compliance Evidence
Basel III/CRR	Model risk management	C4 Fidelity Testing, C2 Governance Record	Training data validation for credit risk models
MiFID II	Algorithm testing	C4 Fidelity Testing	Synthetic transaction data quality validated
AML/KYC	Model effectiveness	C4 Fidelity Testing (rare events)	Preservation of fraud/AML patterns verified
GDPR + Finance	Legitimate interest for fraud	C1 Purpose Sheet, C3 Privacy Risk Testing	Balancing test documented

13.5.3 Table D.5.3: Insurance (Solvency II, IDD)

Regulation	Requirement	SDCF Control	Compliance Evidence
Solvency II	Internal model validation	C4 Fidelity Testing	Synthetic claims data statistical adequacy
IDD	Algorithmic pricing fairness	C5 Fairness Assessment	Bias testing for pricing models
GDPR + Insurance	Pricing transparency	C6 Transparency Pack	Training data characteristics disclosed

13.6 D.6 Cross-Reference: SDCF Controls to All Regulations

13.6.1 Table D.6.1: C1 Purpose Sheet

Regulation/Standard	Specific Requirement	How C1 Addresses
GDPR Article 5(1)(b)	Purpose limitation	Explicitly defines purpose upfront
GDPR Article 6	Lawful basis	Documents lawful basis consideration
EU AI Act Article 10(2)	Appropriate for purpose	Defines AI system deployment context
ISO 27001 A.5.9	Asset inventory	Identifies synthetic data asset
ISO 42001 4.1	Context understanding	Documents organisational context
NIST Privacy Framework	Inventory-P	Inventories privacy-sensitive data assets

13.6.2 Table D.6.2: C2 Governance Record

Regulation/Standard	Specific Requirement	How C2 Addresses
GDPR Article 5(2)	Accountability	Documents compliance decisions and evidence
GDPR Article 30	Records of processing	Maintains processing activity records
EU AI Act Article 17	Quality management	Documents data governance procedures
ISO 27001 A.5.1	Policies	Records governance policies and decisions
ISO 42001 6.1	Risk management	Documents risk acceptance decisions
NIST Privacy Framework	Risk Management-P	Records privacy risk decisions

13.6.3 Table D.6.3: C3 Privacy Risk Testing

Regulation/Standard	Specific Requirement	How C3 Addresses
GDPR Article 32	Security of processing	Quantifies residual privacy risk
GDPR Recital 26	Anonymisation test	Provides technical evidence for “reasonable means”
ISO 27701 7.2.2	De-identification	Validates de-identification effectiveness
ISO 27001 A.8.11	Data masking	Tests synthetic data as masking technique
NIST Privacy Framework	Risk Assessment-P	Identifies and quantifies privacy risks

13.6.4 Table D.6.4: C4 Fidelity Testing

Regulation/Standard	Specific Requirement	How C4 Addresses
GDPR Article 5(1)(d)	Accuracy	Validates statistical accuracy of synthetic data
EU AI Act Article 10(2)	Representative, free of errors	Quantifies representativeness and detects errors
ISO 23894 6.2	Data quality	Operationalises data quality assessment
ISO 42001 8.2	AI development	Validates training data quality
NIST AI RMF	Validation	Measures fidelity and utility

13.6.5 Table D.6.5: C5 Fairness Assessment

Regulation/Standard	Specific Requirement	How C5 Addresses
GDPR Article 5(1)(a)	Fairness	Quantifies representation and predictive fairness
EU AI Act Article 10(4)	Bias examination	Identifies, measures, documents bias
ISO 23894 6.3	Fairness	Quantifies fairness variance
ISO 27701 7.4.7	Automated decision-making	Tests for disparate impact
NIST AI RMF	Fairness metrics	Provides quantitative fairness metrics

13.6.6 Table D.6.6: C6 Transparency Pack

Regulation/Standard	Specific Requirement	How C6 Addresses
GDPR Article 12	Transparent information	Provides transparency documentation
EU AI Act Article 13	Transparency obligations	Documents system characteristics
ISO 27001 A.5.13	Information labeling	Labels datasets with conformance and restrictions
ISO 42001 7.5	Documented information	Provides comprehensive documentation bundle
NIST Privacy Framework	Awareness-P	Communicates privacy practices to stakeholders

13.6.7 Table D.6.7: C7 Release Rules

Regulation/Standard	Specific Requirement	How C7 Addresses
GDPR Article 32	Security measures	Specifies technical and organisational controls
GDPR Article 5(1)(e)	Storage limitation	Defines retention and deletion procedures
ISO 27001 A.5.10	Acceptable use	Documents usage restrictions
ISO 27001 A.5.15	Access control	Specifies access control requirements
ISO 27701 6.7.2.2	Transfer safeguards	Documents controls for data transfers
NIST Privacy Framework	Data Security-P	Implements protective measures

13.7 D.7 Using Mapping Tables for Compliance

13.7.1 How to Use These Mappings

For Internal Compliance: 1. Identify applicable regulations/standards 2. Review relevant mapping tables 3. Conduct SDCF assessment (generates required evidence) 4. Cross-reference SDCF deliverables to compliance obligations 5. Document gaps (if any) and additional measures needed

For Auditor/Regulator Presentation: 1. Present SDCF certificate and Transparency Pack 2. Use mapping tables to demonstrate coverage 3. Provide Governance Record as audit trail 4. Reference specific table entries for each requirement

For Procurement/RFP: 1. Include SDCF conformance requirements in vendor contracts 2. Specify required conformance level (SDCF-A for critical use) 3. Request mapping evidence from vendors 4. Use tables to validate vendor claims

13.7.2 Example: Demonstrating EU AI Act Article 10 Compliance

Step 1: Identify requirements from Table D.2.1

Step 2: Gather SDCF evidence: - Relevant: C1 Purpose Sheet shows purpose-bounded assessment - Representative: C4 results show FI > 80 (high fidelity) - Free of errors: C4 red flags check found no quality issues - Appropriate: Overall SDCF-A conformance - Data governance: C2 Governance Record + C7 Release Rules - Bias examination: C5 results show FV = 12 (low fairness concerns)

Step 3: Present to notified body: - Provide complete Transparency Pack (C6) - Reference Table D.2.1 to show systematic coverage - Highlight quantitative evidence (FI, FV scores) - Demonstrate ongoing monitoring (C7 reassessment triggers)

13.7.3 Gaps and Additional Requirements

SDCF does not replace: - Legal opinions on GDPR classification (requires counsel) - Full AI system risk assessment (SDCF covers training data only) - Domain-specific technical validation (clinical, actuarial, financial) - Penetration testing and security audits (operational security)

Organisations must supplement SDCF with: - Independent legal review - Domain expert validation - Security assessments - Full AI system lifecycle governance

End of Appendix D

Continue to Appendix E: Sample Outputs for additional certificate and report examples.

14 Appendix E: Sample Outputs

This appendix provides additional output examples beyond those in Appendix C. See Appendix C.5 for complete Bronze Tier certificate templates.

14.1 E.1 Gold Tier SDCF-A Certificate (Abbreviated)

```
=====
SYNTHETIC DATA COMPLIANCE FRAMEWORK (SDCF)
GOLD TIER ASSESSMENT CERTIFICATE
=====

DATASET INFORMATION
Dataset ID: SYNTH-HEALTH-2025-Q4
Assessment Tier: Gold (Full Source Data Access)
Assessment Date: 2025-11-20
Assessor: Dr. Sarah Johnson, Data Science Lead
Organisation: HealthTech Research Institute

PURPOSE STATEMENT
Intended Use: Training dataset for diabetic retinopathy detection AI system
(high-risk AI under EU AI Act Annex III)

CONFORMANCE LEVEL: SDCF-A (APPROVED)
```

ASSESSMENT RESULTS

Gold Privacy Risk Score (PRS): 18 (95% CI: 15-21) TARGET MET

Gold Fidelity Index (FI): 87 (95% CI: 84-90) TARGET MET

Gold Fairness Variance (FV): 11 (95% CI: 8-14) TARGET MET

INTERPRETATION

LOW PRIVACY RISK - Suitable for controlled research sharing

HIGH FIDELITY - Representative of clinical populations

LOW FAIRNESS CONCERNS - Balanced across demographics and pathology stages

APPROVED FOR STATED PURPOSE

Valid until: 2026-11-20 (12 months) or material change

=====

14.2 E.2 Assessment Report (Executive Summary)

SYNTHETIC DATA ASSESSMENT REPORT

Executive Summary

Dataset: Synthetic Customer Transaction Data (SYNTH-FIN-2025-Q3)

Organisation: Irish Financial Services Ltd.

Assessment: Silver Tier (Aggregate Statistics Available)

Date: 2025-10-15

EXECUTIVE SUMMARY

Purpose: Train fraud detection models for real-time transaction monitoring

Key Findings:

- Privacy Risk: MODERATE (PRS = 32) - Acceptable for internal ML training
- Fidelity: GOOD (FI = 78) - Distributions well-preserved, rare events adequate
- Fairness: EXCELLENT (FV = 9) - Balanced representation across customer segments

Recommendation: SDCF-P (Provisionally Approved)

The dataset is suitable for fraud detection model training with the following conditions:

1. Restrict to model development environment (no production deployment)
2. Enhanced monitoring for bias in model predictions
3. Reassessment before production use

Business Impact:

- Enables safe model development without production data exposure
- Reduces regulatory risk through documented assessment
- Supports Article 10 compliance for AI Act

Next Steps:

1. Implement Release Rules (access controls, monitoring)

2. Train fraud detection models
3. Validate model performance on holdout real data
4. Conduct Gold Tier assessment before production deployment

Full technical report: 45 pages including methodology, detailed results, risk analysis, and recommendations.

=====

14.3 E.3 Risk Statement (Gold Tier)

RISK STATEMENT

Dataset: SYNTH-HEALTH-2025-Q4

Conformance: SDCF-A (Approved)

Assessment Tier: Gold

RESIDUAL RISKS (Post-Mitigation)

1. PRIVACY RISK: LOW (PRS = 18)
 - Re-identification probability: <2% (membership inference testing)
 - Distance to closest record: All synthetic records >0.15 normalized distance
 - Attribute disclosure: No elevation above population baseline
 - Mitigation: Access restricted to secure research environment
2. UTILITY RISK: MINIMAL (FI = 87)
 - Distribution fidelity: 92% similarity to source
 - Correlation preservation: 89% maintained
 - Predictive utility: Models achieve 94% of baseline performance
 - Limitation: Rare pathology stages slightly underrepresented (3% vs. 5%)
3. FAIRNESS RISK: MINIMAL (FV = 11)
 - Representation: All demographics within ±8% of source
 - Predictive parity: Maximum performance gap 7% across groups
 - Limitation: Small sample size for transgender patients (n=45)

RISK ACCEPTANCE

HealthTech Research Institute accepts these residual risks for diabetic retinopathy detection model development subject to controls in Release Rules.

Risk Owner: Dr. Michael Chen, Chief AI Officer

Date: 2025-11-20

Signature: _____

MONITORING PLAN

- Quarterly bias audits on trained models
- Semi-annual reassessment of synthetic dataset
- Incident reporting for any re-identification attempts
- Annual review of technological developments in privacy attacks

14.4 E.4 JSON Metadata Schema

SDCF assessments can export machine-readable metadata:

```
{
  "sdcf_version": "1.0",
  "assessment": {
    "dataset_id": "SYNTH-FIN-2025-Q3",
    "assessment_date": "2025-10-15T14:30:00Z",
    "tier": "Silver",
    "assessor": {
      "name": "Jane Smith",
      "organisation": "Irish Financial Services Ltd.",
      "credentials": "Senior Data Scientist, SDCF Trained"
    },
    "purpose": {
      "use_case": "Fraud detection model training",
      "disclosure_scope": "Internal development environment",
      "regulatory_context": ["GDPR Article 6(1)(f)", "EU AI Act Article 10"]
    },
    "scores": {
      "prs": {
        "value": 32,
        "confidence_interval": [25, 39],
        "components": {
          "mir": 28,
          "rsr": 35,
          "adr": 30
        }
      },
      "interpretation": "MODERATE_RISK"
    },
    "fi": {
      "value": 78,
      "confidence_interval": [72, 84],
      "components": {
        "ds": 82,
        "dp": 76,
        "pu": 77
      }
    },
    "interpretation": "GOOD_FIDELITY"
  },
  "fv": {
    "value": 9,
    "confidence_interval": [6, 12],
    "components": {
      "rv": 8,
      "ppv": 10
    }
  },
}
```

```

        "interpretation": "LOW_FAIRNESS_CONCERNS"
    }
},
"conformance": {
    "level": "SDCF-P",
    "rationale": "Privacy risk slightly above target; acceptable with controls",
    "conditions": [
        "Restrict to development environment",
        "Enhanced bias monitoring",
        "Reassessment before production"
    ],
    "validity_period": "2026-04-15T00:00:00Z"
}
}
}

```

End of Appendix E

Continue to Appendix F: Reference Implementations for code examples and tool integration patterns.

15 Appendix F: Reference Implementations

This appendix provides practical code examples for implementing SDCF assessments using open-source tools.

15.1 F.1 Bronze Tier Assessment Using SDMetrics

15.1.1 F.1.1 Setup and Data Loading

```

# sdcf_bronze_assessment.py
"""
SDCF Bronze Tier Assessment Reference Implementation
Uses SDMetrics for metric computation
"""

import pandas as pd
import numpy as np
from sdmetrics.reports.single_table import QualityReport
from sklearn.neighbors import LocalOutlierFactor
from scipy.spatial.distance import pdist, squareform
from scipy.stats import chi2_contingency
import warnings
warnings.filterwarnings('ignore')

# Load synthetic data (Bronze Tier: no source data available)
synthetic_data = pd.read_csv('synthetic_data.csv')

```

```
print(f"Loaded synthetic data: {synthetic_data.shape[0]} rows, {synthetic_data.shape[1]} columns")
```

15.1.2 F.1.2 B-PRS Component 1: Membership Inference Proxy (Outlier Analysis)

```
def compute_b_mir(data, contamination=0.1):
    """
    Bronze Tier Membership Inference Risk (outlier proxy)

    Args:
        data: DataFrame with synthetic data
        contamination: Expected proportion of outliers (0.1 = 10%)

    Returns:
        B-MIR score (0-100, lower is better)
    """
    # Prepare numeric data for outlier detection
    numeric_cols = data.select_dtypes(include=[np.number]).columns
    X = data[numeric_cols].fillna(data[numeric_cols].median())

    # Local Outlier Factor
    lof = LocalOutlierFactor(n_neighbors=20, contamination=contamination)
    outlier_labels = lof.fit_predict(X)
    outlier_scores = -lof.negative_outlier_factor_ # Higher = more outlying

    # Normalize to 0-100
    outlier_normalized = (outlier_scores - outlier_scores.min()) / \
        (outlier_scores.max() - outlier_scores.min()) * 100

    # Compute B-MIR: percentage in top 10%
    threshold_90 = np.percentile(outlier_normalized, 90)
    high_outliers = (outlier_normalized > threshold_90).sum()
    b_mir_raw = (high_outliers / len(data)) * 100

    # Add conservative penalty
    b_mir_adjusted = b_mir_raw + 10

    print(f"B-MIR: {b_mir_adjusted:.1f}")
    print(f"  High outliers: {high_outliers} ({high_outliers/len(data)*100:.1f}%)")

    return b_mir_adjusted

b_mir = compute_b_mir(synthetic_data)
```

15.1.3 F.1.3 B-PRS Component 2: Record Similarity Proxy

```
def compute_b_rsr(data, threshold=0.05):
    """
```

Bronze Tier Record Similarity Risk (internal diversity)

Args:

data: DataFrame with synthetic data
threshold: Distance threshold for near-duplicates

Returns:

B-RSR score (0-100, lower is better)

```
"""
# Use numeric columns for distance computation
numeric_cols = data.select_dtypes(include=[np.number]).columns
X = data[numeric_cols].fillna(data[numeric_cols].median())

# Normalize
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_normalized = scaler.fit_transform(X)

# Compute pairwise distances
distances = pdist(X_normalized, metric='euclidean')
distance_matrix = squareform(distances)

# Find nearest neighbor distance (excluding self)
np.fill_diagonal(distance_matrix, np.inf)
nearest_distances = distance_matrix.min(axis=1)

# Compute B-RSR: percentage with low internal distance
near_duplicates = (nearest_distances < threshold).sum()
b_rsr_raw = (near_duplicates / len(data)) * 100

# Add conservative penalty
b_rsr_adjusted = b_rsr_raw + 15

print(f"B-RSR: {b_rsr_adjusted:.1f}")
print(f"  Near-duplicates: {near_duplicates} ({near_duplicates/len(data)*100:.1f}%)")

return b_rsr_adjusted

b_rsr = compute_b_rsr(synthetic_data)
```

15.1.4 F.1.4 B-PRS Component 3: Attribute Disclosure Proxy

```
def compute_b_adr(data, quasi_identifiers, sensitive_attributes):
    """
    Bronze Tier Attribute Disclosure Risk (correlation analysis)

    Args:
        data: DataFrame with synthetic data
```



```

    quasi_identifiers: List of QI column names
    sensitive_attributes: List of sensitive column names

Returns:
    B-ADR score (0-100, lower is better)
"""
from scipy.stats import pearsonr, spearmanr

max_correlations = []

for sensitive_attr in sensitive_attributes:
    correlations = []

    for qi in quasi_identifiers:
        # Check if both are numeric
        if pd.api.types.is_numeric_dtype(data[qi]) and \
            pd.api.types.is_numeric_dtype(data[sensitive_attr]):
            corr, _ = pearsonr(data[qi].fillna(0), data[sensitive_attr].fillna(0))
        else:
            # Use contingency table for categorical
            contingency = pd.crosstab(data[qi], data[sensitive_attr])
            chi2, p, dof, expected = chi2_contingency(contingency)
            # Cramér's V
            n = contingency.sum().sum()
            corr = np.sqrt(chi2 / (n * (min(contingency.shape) - 1)))

        correlations.append(abs(corr))

    max_correlations.append(max(correlations))

b_adr = np.mean(max_correlations) * 100

print(f"B-ADR: {b_adr:.1f}")
for i, attr in enumerate(sensitive_attributes):
    print(f" {attr} max correlation: {max_correlations[i]:.2f}")

return b_adr

# Define quasi-identifiers and sensitive attributes for your dataset
quasi_identifiers = ['age', 'zip_code', 'gender']
sensitive_attributes = ['income', 'health_condition']

b_adr = compute_b_adr(synthetic_data, quasi_identifiers, sensitive_attributes)

```

15.1.5 F.1.5 Composite B-PRS

```

def compute_b_prs(b_mir, b_rsr, b_adr, weights=(0.4, 0.4, 0.2), penalty=20):
    """

```

Composite Bronze Tier Privacy Risk Score

Args:

b_mir, b_rsr, b_adr: Component scores
weights: Component weights (default: 0.4, 0.4, 0.2)
penalty: Uncertainty penalty (default: 20)

Returns:

B-PRS composite score (0-100, lower is better)

```
"""
b_prs = (weights[0] * b_mir +
         weights[1] * b_rsr +
         weights[2] * b_adr +
         penalty)

b_prs = min(100, max(0, b_prs))  # Clamp to 0-100

print(f"\n=== BRONZE TIER PRIVACY RISK SCORE ===")
print(f"B-MIR: {b_mir:.1f} (weight: {weights[0]})")
print(f"B-RSR: {b_rsr:.1f} (weight: {weights[1]})")
print(f"B-ADR: {b_adr:.1f} (weight: {weights[2]})")
print(f"Uncertainty Penalty: +{penalty}")
print(f"B-PRS: {b_prs:.1f}")

# Interpretation
if b_prs < 20:
    interpretation = "LOW RISK"
elif b_prs < 50:
    interpretation = "MODERATE RISK"
elif b_prs < 80:
    interpretation = "HIGH RISK"
else:
    interpretation = "VERY HIGH RISK"

print(f"Interpretation: {interpretation}")

return b_prs
```

```
b_prs = compute_b_prs(b_mir, b_rsr, b_adr)
```

15.1.6 F.1.6 B-FI: Fidelity Assessment (Domain Validation)

```
def compute_b_fi_distribution(data, known_stats=None):
    """
    Bronze Tier Fidelity - Distribution validation against known statistics

    Args:
        data: Synthetic data
```

```

    known_stats: Dict of column -> expected distribution (e.g., from census)

Returns:
    B-DS score (0-100, higher is better)
"""
violations = []
similarity_scores = []

# Check for impossible values
if 'age' in data.columns:
    if (data['age'] < 0).any() or (data['age'] > 120).any():
        violations.append("Impossible age values")

    # Compare to known age distribution (example: use census data)
    if known_stats and 'age' in known_stats:
        obs_dist = data['age'].value_counts(normalize=True, bins=10).sort_index()
        exp_dist = known_stats['age']
        # Chi-square test
        chi2, p = chisquare(obs_dist, exp_dist)
        similarity = 1 - (chi2 / len(obs_dist))
        similarity_scores.append(max(0, similarity))

# Average similarity
public_similarity = np.mean(similarity_scores) if similarity_scores else 0.5
violation_penalty = len(violations) * 5
uncertainty_penalty = 15

b_ds = (public_similarity * 100) - violation_penalty - uncertainty_penalty
b_ds = max(0, b_ds)

print(f"\n=== BRONZE TIER FIDELITY (Distribution) ===")
print(f"B-DS: {b_ds:.1f}")
print(f"Violations: {len(violations)}")
if violations:
    for v in violations:
        print(f"    - {v}")

return b_ds

# Example: Provide known statistics (from census, domain knowledge, etc.)
known_stats = {
    'age': [0.12, 0.18, 0.20, 0.18, 0.15, 0.10, 0.05, 0.02] # Age distribution bins
}

b_ds = compute_b_fi_distribution(synthetic_data, known_stats)

```

15.1.7 F.1.7 Complete Bronze Assessment Function

```
def sdcf_bronze_assessment(synthetic_data, quasi_identifiers, sensitive_attributes,
                           known_stats=None, output_path='sdcf_bronze_report.json'):
    """
    Complete SDCF Bronze Tier Assessment

    Returns:
        Dictionary with all scores and conformance determination
    """
    print("="*60)
    print("SDCF BRONZE TIER ASSESSMENT")
    print("="*60)

    # Compute B-PRS
    b_mir = compute_b_mir(synthetic_data)
    b_rsr = compute_b_rsr(synthetic_data)
    b_adr = compute_b_adr(synthetic_data, quasi_identifiers, sensitive_attributes)
    b_prs = compute_b_prs(b_mir, b_rsr, b_adr)

    # Compute B-FI (simplified - distribution only for brevity)
    b_ds = compute_b_fi_distribution(synthetic_data, known_stats)
    b_fi = b_ds # In practice, would include B-DP and B-PU

    # Determine conformance (simplified logic)
    if b_prs < 50 and b_fi > 60:
        conformance = "SDCF-P"
        rationale = "Acceptable for lower-risk use with controls"
    elif b_prs >= 50 or b_fi < 40:
        conformance = "SDCF-R"
        rationale = "Not suitable for intended purpose"
    else:
        conformance = "SDCF-P"
        rationale = "Moderate quality, enhanced controls required"

    results = {
        "assessment_date": pd.Timestamp.now().isoformat(),
        "tier": "Bronze",
        "scores": {
            "b_prs": round(b_prs, 1),
            "b_mir": round(b_mir, 1),
            "b_rsr": round(b_rsr, 1),
            "b_adr": round(b_adr, 1),
            "b_fi": round(b_fi, 1)
        },
        "conformance": conformance,
        "rationale": rationale
    }
```

```

    # Save to JSON
    import json
    with open(output_path, 'w') as f:
        json.dump(results, f, indent=2)

    print(f"\n=== ASSESSMENT COMPLETE ===")
    print(f"Conformance: {conformance}")
    print(f"Report saved to: {output_path}")

    return results

# Execute full assessment
results = sdcf_bronze_assessment(
    synthetic_data=synthetic_data,
    quasi_identifiers=['age', 'zip_code', 'gender'],
    sensitive_attributes=['income', 'health_condition'],
    known_stats=known_stats
)

```

15.2 F.2 Gold Tier Assessment Using SDMetrics

15.2.1 F.2.1 Complete Gold Tier Implementation

```

# sdcf_gold_assessment.py
"""
SDCF Gold Tier Assessment Reference Implementation
Requires both source and synthetic data
"""

import pandas as pd
import numpy as np
from sdmetrics.reports.single_table import QualityReport
from sdmetrics.single_table import NewRowSynthesis, BoundaryAdherence
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, roc_auc_score

# Load both source and synthetic data
source_data = pd.read_csv('source_data.csv')
synthetic_data = pd.read_csv('synthetic_data.csv')

print(f"Source: {source_data.shape}, Synthetic: {synthetic_data.shape}")

def gold_tier_privacy_risk(source_data, synthetic_data):
    """
    Gold Tier Privacy Risk Score using SDMetrics
    """

```

```

# Use SDMetrics privacy metrics
from sdmetrics.single_table import NewRowSynthesis

# NewRowSynthesis measures if synthetic rows are novel (not copies)
nrs_score = NewRowSynthesis.compute(
    real_data=source_data,
    synthetic_data=synthetic_data
)

# Convert to risk score (lower novelty = higher risk)
mir_component = (1 - nrs_score) * 100

# Distance to closest record
from scipy.spatial.distance import cdist
distances = cdist(synthetic_data.select_dtypes(include=[np.number]),
                  source_data.select_dtypes(include=[np.number]),
                  metric='euclidean')
dcr = distances.min(axis=1)
rsr_component = (dcr < 0.1).sum() / len(synthetic_data) * 100

# Simplified PRS
prs = 0.5 * mir_component + 0.5 * rsr_component

print(f"Gold Tier PRS: {prs:.1f}")
return prs

def gold_tier_fidelity(source_data, synthetic_data, target_column):
    """
    Gold Tier Fidelity Index using SDMetrics
    """
    # Generate SDMetrics quality report
    metadata = {
        'columns': {col: {'sdtype': 'numerical' if pd.api.types.is_numeric_dtype(source_data[col])
                        else 'categorical'}
                   for col in source_data.columns}
    }

    report = QualityReport()
    report.generate(source_data, synthetic_data, metadata)

    # Extract overall quality score (serves as FI)
    fi = report.get_score() * 100

    # Predictive utility test
    X_real = source_data.drop(columns=[target_column])
    y_real = source_data[target_column]
    X_real_train, X_real_test, y_real_train, y_real_test = train_test_split(
        X_real, y_real, test_size=0.2, random_state=42

```

```

)

# Train on real
model_real = RandomForestClassifier(random_state=42)
model_real.fit(X_real_train.select_dtypes(include=[np.number]), y_real_train)
pred_real = model_real.predict(X_real_test.select_dtypes(include=[np.number]))
acc_real = accuracy_score(y_real_test, pred_real)

# Train on synthetic
X_synth = synthetic_data.drop(columns=[target_column])
y_synth = synthetic_data[target_column]
model_synth = RandomForestClassifier(random_state=42)
model_synth.fit(X_synth.select_dtypes(include=[np.number]), y_synth)
pred_synth = model_synth.predict(X_real_test.select_dtypes(include=[np.number]))
acc_synth = accuracy_score(y_real_test, pred_synth)

# Utility preservation
utility_ratio = acc_synth / acc_real if acc_real > 0 else 0
pu_component = utility_ratio * 100

# Composite FI (simplified)
fi_composite = 0.5 * fi + 0.5 * pu_component

print(f"Gold Tier FI: {fi_composite:.1f}")
print(f"  SDMetrics Quality: {fi:.1f}")
print(f"  Predictive Utility: {pu_component:.1f}")

return fi_composite

# Execute Gold Tier assessment
prs_gold = gold_tier_privacy_risk(source_data, synthetic_data)
fi_gold = gold_tier_fidelity(source_data, synthetic_data, target_column='target')

```

15.3 F.3 Tool Integration Patterns

15.3.1 F.3.1 mostlyai-qa Integration

```

# Integration with mostlyai-qa (if available)
"""
mostlyai-qa provides fidelity and novelty reports
SDCF can consume these reports for Silver/Bronze tier assessments
"""

# Example: Parse mostlyai-qa JSON output
import json

def parse_mostlyai_qa_report(qa_report_path):
    """

```

```

Parse mostlyai-qa report and map to SDCF scores
"""
with open(qa_report_path, 'r') as f:
    qa_report = json.load(f)

# Extract fidelity score
fidelity_score = qa_report.get('fidelity', {}).get('overall_score', 0)

# Extract novelty score (proxy for privacy)
novelty_score = qa_report.get('novelty', {}).get('score', 0)

# Map to SDCF
# High novelty = low privacy risk
prs_estimate = (1 - novelty_score) * 100 + 20 # Add Bronze penalty
fi_estimate = fidelity_score * 100 - 15 # Subtract Bronze penalty

return {
    'prs': prs_estimate,
    'fi': fi_estimate,
    'source': 'mostlyai-qa'
}

```

15.3.2 F.3.2 Generating SDCF Certificate

```

def generate_sdcf_certificate(assessment_results, output_path='certificate.txt'):
    """
    Generate formatted SDCF certificate from assessment results
    """
    cert = f"""
{' '*80}
SYNTHETIC DATA COMPLIANCE FRAMEWORK (SDCF)
{assessment_results['tier'].upper()} TIER ASSESSMENT CERTIFICATE
{' '*80}

DATASET INFORMATION
Dataset ID: {assessment_results.get('dataset_id', 'N/A')}
Assessment Date: {assessment_results['assessment_date']}
Tier: {assessment_results['tier']}

ASSESSMENT RESULTS
Privacy Risk Score (PRS): {assessment_results['scores']['b_prs']:.1f}
Fidelity Index (FI): {assessment_results['scores']['b-fi']:.1f}

CONFORMANCE LEVEL: {assessment_results['conformance']}

RATIONALE
{assessment_results['rationale']}

```



```

{'='*80}
"""

with open(output_path, 'w') as f:
    f.write(cert)

print(f"Certificate generated: {output_path}")
return cert

# Generate certificate
certificate = generate_sdcf_certificate(results, output_path='sdcf_certificate.txt')
print(certificate)

```

End of Appendix F

Continue to Supporting Materials: Glossary, References, Acknowledgments, and Version History.

16 Supporting Materials

16.1 Glossary

Anonymous Data: Data that does not relate to an identified or identifiable natural person, or has been rendered anonymous such that re-identification is not reasonably likely. Not subject to GDPR (Recital 26).

Attribute Disclosure: Privacy risk where adversary infers sensitive attributes for individuals they partially know (quasi-identifier attack).

Bronze Tier: SDCF assessment conducted without source data access; relies on synthetic-only analysis with conservative risk estimates.

Conformance Level: SDCF determination of fitness for purpose: SDCF-A (Approved), SDCF-P (Provisional), SDCF-R (Restricted).

Data Protection Impact Assessment (DPIA): Systematic assessment of privacy risks required under GDPR Article 35 for high-risk processing.

Differential Privacy: Mathematical framework providing formal privacy guarantees through controlled noise addition (characterised by ϵ , δ parameters).

EU AI Act: European Union regulation establishing requirements for AI systems, including training data governance (Article 10) for high-risk systems.

Fairness Variance (FV): SDCF composite metric quantifying bias and representation issues; lower scores indicate lower fairness concerns.

Fidelity Index (FI): SDCF composite metric quantifying statistical similarity between synthetic and source data; higher scores indicate higher fidelity.

GDPR: General Data Protection Regulation (EU 2016/679), establishing comprehensive data protection requirements.

Gold Tier: SDCF assessment with full source data access, enabling rigorous privacy, fidelity, and fairness testing.

Membership Inference: Privacy attack determining whether specific individual’s data was in training dataset.

Personal Data: Information relating to identified or identifiable natural person (GDPR Article 4(1)); subject to full GDPR obligations.

Privacy Risk Score (PRS): SDCF composite metric quantifying re-identification and disclosure risk; lower scores indicate lower privacy risk.

Pseudonymous Data: Personal data processed such that it cannot be attributed to data subject without additional information kept separately (GDPR Article 4(5)); remains personal data.

Purpose-Bounded Assessment: SDCF philosophy requiring explicit fitness-for-purpose determination rather than generic quality scoring.

Quasi-Identifier: Attribute that alone doesn’t identify individuals but combined with other attributes or auxiliary data enables identification (e.g., age + gender + zip code).

Silver Tier: SDCF assessment with partial source data access (aggregates, samples, metadata); moderate confidence level.

Synthetic Data: Artificially generated data that preserves statistical properties of source data while not directly corresponding to real individuals.

Transparency Pack: SDCF deliverable (C6) providing comprehensive documentation of assessment results, limitations, and usage guidance.

16.2 References

16.2.1 Regulatory Documents

1. **European Union.** Regulation (EU) 2016/679 (General Data Protection Regulation). April 27, 2016.
2. **European Union.** Regulation (EU) 2024/1689 (Artificial Intelligence Act). May 21, 2024.
3. **European Data Protection Board.** Guidelines 01/2025 on Pseudonymisation. January 17, 2025.
4. **Article 29 Data Protection Working Party.** Opinion 05/2014 on Anonymisation Techniques. April 10, 2014.
5. **European Commission.** Commission Implementing Regulation on AI Act harmonised standards (in development, expected 2026).

16.2.2 Standards

6. **ISO/IEC 27001:2022.** Information security, cybersecurity and privacy protection — Information security management systems — Requirements.

7. **ISO/IEC 27701:2019.** Security techniques — Extension to ISO/IEC 27001 and ISO/IEC 27002 for privacy information management.
8. **ISO/IEC 23894:2023.** Information technology — Artificial intelligence — Guidance on risk management.
9. **ISO/IEC 42001:2023.** Information technology — Artificial intelligence — Management system.
10. **NIST.** Privacy Framework: A Tool for Improving Privacy Through Enterprise Risk Management, Version 1.0. January 16, 2020.
11. **NIST.** Artificial Intelligence Risk Management Framework (AI RMF 1.0). January 26, 2023.

16.2.3 Technical Literature

12. **Jordon, J., Yoon, J., van der Schaar, M.** “Synthetic Data – what, why and how?” arXiv:2205.03257, May 2022.
13. **Stadler, T., Oprisanu, B., Troncoso, C.** “Synthetic Data – Anonymisation Groundhog Day.” USENIX Security 2022.
14. **Torkzadehmahani, R., Kairouz, P., Paten, B.** “DP-CGAN: Differentially Private Synthetic Data and Label Generation.” CVPR Workshop 2020.
15. **Yale Privacy Lab.** “Synthetic Data Vault (SDV) – Documentation and User Guide.” MIT Data to AI Lab, 2023.
16. **MOSTLY AI.** “Quality Assurance for Synthetic Data – Technical Documentation.” April 2025.

16.2.4 AI Training on Synthetic Data

17. **Pleias (French AI Lab).** “Baguettotron: Efficient Reasoning with Synthetic Training Data.” November 2025. [Referenced in Sections 2.2, 2.5]
18. **Zhang, Y., et al.** “Synthetic Data Prevents Model Collapse in Large Language Models.” arXiv preprint, 2024.

16.2.5 Privacy and Security

19. **Shokri, R., Stronati, M., Song, C., Shmatikov, V.** “Membership Inference Attacks Against Machine Learning Models.” IEEE S&P 2017.
20. **Stadler, T., et al.** “Synthetic Data – what, why and how?” Royal Society Open Science, 2022.

16.2.6 Fairness and Bias

21. **Mehrabi, N., Morstatter, F., Saxena, N., Lerman, K., Galstyan, A.** “A Survey on Bias and Fairness in Machine Learning.” ACM Computing Surveys, 2021.
22. **Mitchell, S., et al.** “Algorithmic Fairness: Choices, Assumptions, and Definitions.” Annual Review of Statistics, 2021.

16.2.7 Tools and Software

23. **SDMetrics (SDV)**. <https://docs.sdv.dev/sdmetrics/> (Accessed November 2025)
24. **mostlyai-qa**. <https://github.com/mostly-ai/mostlyai-qa> (Accessed November 2025)
25. **Gretel.ai Platform**. <https://gretel.ai/> (Accessed November 2025)

16.3 Acknowledgments

16.3.1 Development and Contributions

The Synthetic Data Compliance Framework (SDCF) was developed by **Wayne Kearns** (Senior Technology Program Leader, Kaionix Labs) based on: - 8+ years of experience in pharmaceutical and healthcare technology sectors - Doctoral research on healthcare cybersecurity governance at Ireland's Health Service Executive (HSE) - Practitioner-researcher methodology informed by operational resilience and AI governance programs

16.3.2 Intellectual Foundations

SDCF builds upon: - Privacy research from academic institutions (membership inference, anonymisation techniques) - Open-source synthetic data quality tools (SDMetrics/SDV, mostlyai-qa) - International standards (ISO/IEC, NIST frameworks) - European regulatory guidance (EDPB, EU AI Act implementation materials)

16.3.3 Community Feedback

The framework will continue to evolve based on: practitioner implementation experiences, academic peer review and validation studies, regulatory developments and enforcement patterns, and tool ecosystem evolution

16.3.4 Contact and Contributions

Technical Questions: wayne.kearns@nortesconsulting.com

Framework Repository: <https://www.kaionix.com/kaionix-labs>

Feedback and Contributions: Welcomed via email or GitHub issues (if repository established)

Empirical Validation: Researchers interested in validating SDCF thresholds, conducting comparison studies, or contributing domain-specific calibrations are encouraged to reach out.

16.4 Version History

16.4.1 Version 1.95 (December 2025) - Reference Corrections

Version 1.95 corrects citation key errors identified during systematic reference verification:

- **Citation year corrections:** Fixed `kleinberg2016inherent` → `kleinberg2017inherent` (ITCS 2017, not 2016); fixed `yale2019generation` → `yale2020generation` (Neurocomputing 2020, not 2019)
- **Bibliography file:** Updated to verified bibliography with corrected BibTeX entry types (7 corrections: journal vs. conference classifications)
- **No methodology changes:** All thresholds, metrics, assessment procedures, and empirical results unchanged

- **Rationale:** Ensures bibliographic accuracy for peer review; all 34 references web-verified against original sources

16.4.2 Version 1.9 (December 2025) - HAL Optimised

Version 1.9 eliminates all remaining overclaim language and strengthens legal/regulatory framing for HAL preprint submission:

- **Abstract rewrite (HAL-optimised):** Replaced abstract with legally conservative, journal-ready version; eliminated "performs as designed" and "demonstrates" in favour of "provides indicative support" and "is consistent with"; improved HAL indexability by front-loading key terms (synthetic data, GDPR, EU AI Act, audit-ready)
- **Section 8.6 Key Findings (conservative rewrite):** Removed all confirmation language ("confirms", "validates", "demonstrates"); replaced with "indicative observations", "suggests", "supports"; changed from 5 findings format to cohesive narrative; maintained technical accuracy while eliminating overclaim risk
- **Section C.8 renamed & simplified:** "Validated Performance Expectations" → "Observed Performance Ranges (Bronze Tier)"; removed detailed practitioner guidance tables; condensed to essential indicative ranges only; added explicit caveat: "not normative regulatory thresholds"
- **Appendix B legal disclaimer (EU-hardened):** Strengthened opening with consolidated legal framing; separated anonymisation and EU AI Act into distinct subsections; removed quasi-legal guidance tone; emphasised technical vs. regulatory qualification distinction
- **Version consistency:** Updated all version references from 1.8 to 1.9 throughout document
- **No methodology changes:** All thresholds, metrics, assessment procedures, and empirical results unchanged
- **Rationale:** Eliminates all remaining overclaim phrases identified in simulated peer review; optimises for HAL moderation approval; ensures long-term defensibility; positions cleanly for future journal submission

Version 1.8 strengthens methodological transparency with honest framing of validation scope:

- **Preliminary validation framing:** Updated Abstract, Section 1, Section 7 to clarify that empirical validation ($n=10$ datasets) provides preliminary evidence, not comprehensive statistical generalisation
- **Future work explicit:** Added statement throughout that statistical expansion to $n>50$ datasets with adversarial validation (membership inference, linkage attacks) and ROC-based threshold calibration is planned for v2.0
- **Conservative claims:** Changed "confirms methodology performs as designed" to "initial evidence suggests" and "comprehensive validation" to "preliminary empirical validation"
- **TVAE specificity:** Added ($n=2$ demographic datasets) qualifier to TVAE superiority claim
- **Document status:** Updated title from "Version 1.7" to "Version 1.8 (Preprint - Preliminary Validation)"
- **Version consistency:** Updated all version references from 1.7 to 1.8 throughout document
- **No methodology changes:** All thresholds, metrics, assessment procedures, and empirical results unchanged
- **Rationale:** Positions framework for HAL preprint submission with academically defensible claims; enables community feedback to inform v2.0 journal submission

Version 1.7 completes UK English standardisation to 100%:

- **Final corrections:** Converted remaining 2 instances of “Authorized” to “Authorised” in template labels (lines 2983, 7200)
- **Result:** 100% UK English consistency (117+ total conversions)
- **Version consistency:** Updated all version references from 1.6 to 1.7 throughout document
- **No methodology changes:** All thresholds, metrics, assessment procedures, and empirical results unchanged

Version 1.6 completes UK English standardisation to 99.5% consistency:

- **Critical fix:** Corrected mixed US/UK spelling in same sentence (“contextualize” → “contextualise”, line 763)
- **Final -ize conversions:** authorised (6x), Unauthorised (2x), authorises (1x), Standardised (2x), Operationalises (2x), Equalised (3x), Formalise (1x), Categorise (1x), Harmonised (1x)
- **Mathematical terms:** Retained normalise, minimise, maximise (internationally accepted in technical literature)
- **Version consistency:** Updated all version references from 1.5 to 1.6 throughout document
- **Result:** 99.5% UK English consistency (115+ total conversions)
- **No methodology changes:** All thresholds, metrics, assessment procedures, and empirical results unchanged

Version 1.5 resolves LaTeX compilation issues identified in external review:

- **lmodern.sty:** Made optional using `\IfFileExists` to prevent compilation failure when package unavailable
- **Extra closing brace:** Removed extraneous brace from Version 1.0 label (line 9825)
- **UK English completion:** Added final -ize → -ise conversions (operationalise, prioritise, standardise, harmonise, generalise, categorise, characterise, optimise) for 100% consistency
- **Version consistency:** Updated all version references from 1.4 to 1.5 throughout document
- **No methodology changes:** All thresholds, metrics, assessment procedures, and empirical results unchanged

Version 1.4 applies formatting and consistency fixes identified in peer review:

- **List formatting:** Corrected inline lists to proper LaTeX itemize structure (lines 494–534)
- **Version consistency:** Updated all version references from 1.1 to 1.4 throughout document
- **UK English:** Confirmed consistent use of UK English spelling (organisation, specialised, materialise)
- **Typography:** Fixed hypertarget/label mismatch, removed extraneous closing brace
- **Capitalization:** Corrected sentence-initial capitalization
- **No methodology changes:** All thresholds, metrics, assessment procedures, and empirical results unchanged

Version 1.3 improved presentation quality without changing methodology:

- **Abstract:** Condensed from 2,046 to approximately 1,800 characters for arXiv compliance
- **Typography:** Improved mathematical notation (k -anonymity), proper en-dashes for ranges, non-breaking spaces
- **Clarity:** Minor prose improvements for readability

Version 1.1 added empirical validation of the Bronze Tier methodology:

- **Section 7:** Complete Bronze Tier retrospective validation study (10 datasets, 7 domains)
- **Appendix G:** Detailed validation results and statistical analysis
- **Appendix C.8-C.9:** Validated performance expectations and synthesis method benchmarks
- **Empirical findings:** Framework performs as designed; provides evidence-based guidance for practitioners

16.4.3 Version 1.0 (November 2025) - Initial Public Release

Status: Request for Comments (RFC)

What's Included: - Complete methodology for Gold, Silver, and Bronze Tier assessments - Mathematical definitions for Privacy Risk Score (PRS), Fidelity Index (FI), Fairness Variance (FV) - Seven control sets (C1-C7) with implementation guidance - Provisional thresholds pending empirical validation - Regulatory mapping for GDPR, EU AI Act, ISO/IEC standards, NIST frameworks - Reference implementations using SDMetrics and Python - Certificate templates and sample outputs

Known Limitations: - Thresholds are provisional (require domain-specific calibration and validation) - Limited to structured tabular data (unstructured data out of scope) - Bronze Tier methodology is novel and untested in practice - No large-scale empirical validation studies conducted - Regulatory landscape continues to evolve (AI Act implementation ongoing)

Anticipated Evolution: - **v1.1 (Q1 2026):** Incorporate initial community feedback, refine Bronze Tier based on early implementations - **v1.2 (Q2 2026):** Add validated threshold calibrations for specific domains (healthcare, finance, insurance) - **v2.0 (2026-2027):** Major update incorporating empirical validation, potential software tooling, expanded use cases

Feedback Encouraged: Organisations implementing SDCF are encouraged to share: threshold calibration findings, domain-specific adaptations - Implementation challenges and solutions - Comparison with alternative methodologies - Regulatory acceptance experiences

END OF DOCUMENT

17 Document Summary

Title: Synthetic Data Compliance Framework (SDCF) Version 1.1

Author: Wayne Kearns, Kaionix Labs

Date: December 2025

Pages: ~105 pages

Licence: CC BY-SA 4.0

Core Components: - Purpose-bounded assessment methodology - Three pillars: Privacy (PRS), Fidelity (FI), Fairness (FV) - Three tiers: Gold, Silver, Bronze (handles synthetic-only scenarios) - Three conformance levels: SDCF-A, SDCF-P, SDCF-R - Seven control sets: C1-C7 (Purpose, Governance, Privacy, Fidelity, Fairness, Transparency, Release)

Key Innovation: Bronze Tier methodology for assessing synthetic data without source access - addresses most common real-world constraint ignored by existing frameworks.

Primary Use Cases: - EU GDPR compliance (anonymisation vs. pseudonymization determination) - EU AI Act Article 10 (training data governance for high-risk AI systems) - Vendor synthetic data evaluation and procurement - Internal synthetic data quality validation - Regulatory compliance evidence and audit trail

Status: Request for Comments - Provisional thresholds pending empirical validation

Contact: wayne.kearns@nortesconsulting.com

Website: <https://www.kaionix.com/kaionix-labs>

18 Appendix G: Bronze Tier Validation Detailed Results

This appendix provides complete quantitative results from the Bronze Tier retrospective validation study presented in Section 7. All data, code, and reproducibility materials are available in ancillary files.

18.1 G.1 Complete Results Table

Table 38 presents all validation metrics for all 10 datasets.

Table 38: Complete Bronze Tier Validation Results (All Datasets, All Metrics)									
ID	Dataset Name	Records	Features	B-PRS	Level	B-FI	Level	B-FV	Conform
D1	PLEIAs SYNTH	10,000	14	0.777	Critical	0.927	Excellent	0.909	SDCF-R
D2	SDV Adult - GC	32,561	15	0.639	High	0.999	Excellent	0.693	SDCF-R
D3	SDV Adult - CTGAN	32,561	15	0.637	High	0.998	Excellent	0.548	SDCF-R
D4	SDV Adult - TVAE	32,561	15	0.534	High	0.994	Excellent	0.724	SDCF-R
D5	Gretel Safety	8,361	14	0.789	Critical	0.999	Excellent	0.100	SDCF-R
D6	MostlyAI Census	48,842	15	0.631	High	0.997	Excellent	0.692	SDCF-R
D7	MostlyAI CDNOW	69,659	4	0.360	Moderate	0.999	Excellent	0.100	SDCF-A
D8	CMS SynPUF	5,000	14	0.390	Moderate	0.997	Excellent	0.100	SDCF-A
D9	Census SynLBD	10,000	12	0.360	Moderate	1.000	Excellent	0.100	SDCF-A
D10	Jupyter Agent	377	11	0.090	Low	0.999	Excellent	0.128	SDCF-A
Minimum		377	4	0.090	—	0.927	—	0.100	—
Maximum		69,659	15	0.789	—	1.000	—	0.909	—
Mean		21,936	12.9	0.516	Mod-High	0.991	Excellent	0.396	—
Median		10,000	14	0.584	High	0.998	Excellent	0.336	—
Std Dev		23,115	3.3	0.229	—	0.021	—	0.334	—

18.2 G.2 Statistical Summary

18.2.1 Descriptive Statistics by Metric

Table 39 provides comprehensive statistical summary.

Table 39: Statistical Summary: All Metrics						
Metric	Min	Max	Mean	Median	Std Dev	Range
B-PRS	0.090	0.789	0.516	0.584	0.229	0.699
B-FI	0.927	1.000	0.991	0.998	0.021	0.073
B-FV	0.100	0.909	0.396	0.336	0.334	0.809
Records	377	69,659	21,936	10,000	23,115	69,282
Features	4	15	12.9	14	3.3	11

18.2.2 Distribution by Conformance Level

Table 40 aggregates metrics by conformance determination.

Key Observations:

Table 40: Metric Statistics by Conformance Level

Conformance	N	Avg B-PRS	Avg B-FI	Avg B-FV	Avg Records
SDCF-R-Bronze	6	0.684	0.992	0.611	28,447
SDCF-A-Bronze	4	0.300	0.989	0.107	12,169
SDCF-P-Bronze	0	—	—	—	—
Overall	10	0.516	0.991	0.396	21,936

- SDCF-R datasets average B-PRS 0.684 (High-Critical) vs. SDCF-A 0.300 (Low-Moderate) — 2.3x difference
- SDCF-R datasets average B-FV 0.611 (Problematic) vs. SDCF-A 0.107 (Fair) — 5.7x difference
- B-FI shows minimal difference: 0.992 vs. 0.989 (both Excellent)
- Conformance primarily driven by B-FV (5/6 R datasets) and B-PRS (1/6 R datasets)

18.2.3 Correlation Analysis

Table 41 presents Pearson correlation coefficients between metrics and dataset characteristics.

Table 41: Correlation Matrix: Metrics and Dataset Characteristics

	B-PRS	B-FI	B-FV	Features
B-PRS	1.000	−0.412	0.517	0.628
B-FI	−0.412	1.000	−0.301	−0.189
B-FV	0.517	−0.301	1.000	0.743
Features	0.628	−0.189	0.743	1.000
Records	−0.156	0.089	−0.223	−0.391

Interpretation:

- **B-FV ↔ Features:** Strong positive correlation (0.743) — More columns = higher representation variance
- **B-PRS ↔ Features:** Moderate positive correlation (0.628) — Complexity increases privacy risk
- **B-PRS ↔ B-FV:** Moderate positive correlation (0.517) — High risk datasets often have high fairness variance (demographic data pattern)
- **B-FI ↔ others:** Weak correlations — Fidelity largely independent of other factors (modern tools consistently high quality)
- **Records ↔ metrics:** Weak correlations — Dataset size not predictive of metric scores

18.3 G.3 Dataset-Specific Notes

18.3.1 D1: PLEIAs SYNTH

Source: Pleias/common-sense-reasoning-t0.3 on HuggingFace **Description:** Frontier AI reasoning pre-training data, 10K sample from 150M dataset **Domain:** AI Training **Synthesis Method:** LLM-generated (iterative synthesis) **Year:** 2025

Key Characteristics:

- Unique reasoning examples (high content diversity)
- 14 columns: premise, hypothesis, label, reasoning steps, etc.
- Mix of categorical (7 cols) and text (7 cols)

Metric Results:

- B-PRS 0.777 (Critical): Highest uniqueness score (0.892) due to diverse reasoning content
- B-FI 0.927 (Excellent): Lowest in portfolio but still excellent; minor missing values (2.1%)
- B-FV 0.909 (Problematic): Highest variance; diverse AI training content → unbalanced representation

Conformance Trigger: B-FV > 0.30 (Problematic fairness variance)

Interpretation: High B-PRS and B-FV reflect unique, diverse AI training content. This is **expected and appropriate** for training data (diversity is valuable). Restricted conformance flags uncertainty for expert review, not necessarily poor quality.

18.3.2 D2–D4: SDV Adult Variants

Source: UCI Adult Income dataset (real), synthesized using SDV library **Description:** US Census 1994 data (demographics + income), 32,561 records **Domain:** Demographic **Methods:** Gaussian-Copula (D2), CTGAN (D3), TVAE (D4) **Year:** 2024 (synthesized for validation)

Controlled Method Comparison:

- **D2 (GaussianCopula):** B-PRS 0.639, B-FI 0.999, B-FV 0.693 → SDCF-R (B-FV trigger)
- **D3 (CTGAN):** B-PRS 0.637, B-FI 0.998, B-FV 0.548 → SDCF-R (B-FV trigger)
- **D4 (TVAE):** B-PRS 0.534, B-FI 0.994, B-FV 0.724 → SDCF-R (B-FV trigger)

Key Finding: TVAE achieves 16–20% lower B-PRS than CTGAN/GaussianCopula while maintaining B-FI > 0.99. All three trigger Restricted due to B-FV > 0.30 (demographic complexity inherent to source data).

Schema: 15 columns: age, workclass, education, marital-status, occupation, relationship, race, sex, capital-gain, capital-loss, hours-per-week, native-country, income

Why High B-FV? Census demographic data with 9 categorical columns (workclass, education, marital-status, occupation, relationship, race, sex, native-country, income) creates many quasi-identifiers and representation groups → high variance is population characteristic, not synthesis failure.

18.3.3 D5: Gretel Safety Alignment

Source: gretelai/synthetic-gsm8k-reflection-405b on HuggingFace **Description:** LLM safety alignment dataset, 8,361 records **Domain:** AI Safety **Synthesis Method:** LLM-generated (Llama 3.1 405B) **Year:** 2024

Key Characteristics:

- Math reasoning + safety constraints
- 14 columns: question, answer, reasoning, safety_label, etc.
- High text content (10 text cols, 4 categorical)

Metric Results:

- B-PRS 0.789 (Critical): **Highest in portfolio** due to unique safety scenarios
- B-FI 0.999 (Excellent): Near-perfect internal quality
- B-FV 0.100 (Fair): Baseline penalty only (intentionally balanced by design)

Conformance Trigger: B-PRS > 0.70 (Critical privacy risk)

Interpretation: Critical B-PRS reflects unique safety alignment content (each example distinct). Framework correctly flags high uniqueness as privacy risk uncertainty without source comparison. Low B-FV indicates intentional balance across safety categories.

18.3.4 D6: MostlyAI Census

Source: MostlyAI public GitHub repository **Description:** US Census 1994 synthetic, 48,842 records **Domain:** Demographic **Synthesis Method:** Commercial GAN (MostlyAI proprietary) **Year:** 2023

Key Characteristics:

- Commercial-grade demographic synthesis
- 15 columns: similar to SDV Adult variants
- Largest demographic dataset in portfolio

Metric Results:

- B-PRS 0.631 (High): Consistent with demographic pattern
- B-FI 0.997 (Excellent): Commercial vendor quality
- B-FV 0.692 (Problematic): Demographic complexity (same as SDV variants)

Conformance Trigger: B-FV > 0.30

Comparison to SDV Adult: Similar B-PRS (0.631 vs. 0.534–0.639), B-FI (0.997 vs. 0.994–0.999), B-FV (0.692 vs. 0.548–0.724) despite different synthesis method and vendor. Confirms demographic data pattern is robust across methods.

18.3.5 D7: MostlyAI CDNOW Purchases

Source: MostlyAI public GitHub repository **Description:** E-commerce purchase history, 69,659 transactions **Domain:** E-commerce **Synthesis Method:** Commercial GAN (MostlyAI) **Year:** 2023

Key Characteristics:

- **Simplest schema:** Only 4 columns (customer_id, date, quantity, amount)
- **Largest dataset:** 69,659 records
- Transactional time-series data

Metric Results:

- B-PRS 0.360 (Moderate): **Lowest among GAN methods** due to simple schema
- B-FI 0.999 (Excellent): Commercial vendor quality
- B-FV 0.100 (Fair): Baseline (4 cols, minimal categorical complexity)

Conformance: SDCF-A-Bronze (clean pass all thresholds)

Key Insight: Simple transactional schemas achieve Acceptable Bronze conformance. Demonstrates schema complexity is primary driver of B-PRS and B-FV.

18.3.6 D8: CMS DE-SynPUF Demo

Source: CMS (Centers for Medicare & Medicaid Services) **Description:** Medicare synthetic claims, 5,000 beneficiaries **Domain:** Healthcare **Synthesis Method:** Multi-method (government synthesis) **Year:** 2023 (demo variant for validation)

Key Characteristics:

- Government-produced synthetic healthcare data
- 14 columns: beneficiary demographics, claims, costs, procedures
- Multi-year longitudinal structure

Metric Results:

- B-PRS 0.390 (Moderate): Healthcare claims less complex than census demographics
- B-FI 0.997 (Excellent): Government data quality standards
- B-FV 0.100 (Fair): Balanced beneficiary demographics

Conformance: SD CF-A-Bronze (clean pass)

Note: Demo variant used for validation (full CMS DE-SynPUF is 2.2M beneficiaries). Results demonstrate healthcare claims data achieves Acceptable Bronze conformance when demographic complexity is moderate.

18.3.7 D9: US Census SynLBD Demo

Source: US Census Bureau **Description:** Synthetic Longitudinal Business Database, 10,000 establishments **Domain:** Business **Synthesis Method:** Synthetically augmented (Census methodology) **Year:** 2024 (demo variant)

Key Characteristics:

- Government business establishment data
- 12 columns: industry, location, employment, payroll, revenue
- Economic census derived

Metric Results:

- B-PRS 0.360 (Moderate): Business data less personally identifiable
- B-FI 1.000 (Excellent): **Perfect fidelity** (only dataset achieving 1.000)
- B-FV 0.100 (Fair): Simple industry/location representation

Conformance: SD CF-A-Bronze (clean pass)

Perfect B-FI: Reflects Census Bureau data quality standards, comprehensive validation, clean schema design. Demonstrates B-FI 1.000 is achievable (not merely theoretical).

18.3.8 D10: Jupyter Agent Dataset

Source: datadreamer-dev/kaggle_notebook_verifier on HuggingFace **Description:** Synthetic Q&A pairs from Kaggle notebooks, 377 records **Domain:** Code/Data **Synthesis Method:** LLM-generated **Year:** 2025

Key Characteristics:

- **Smallest dataset:** Only 377 records
- **Text-heavy:** 11 columns, mostly code snippets and explanations
- Jupyter notebook agent training data

Metric Results:

- B-PRS 0.090 (Low): **Lowest in portfolio** (only Low risk dataset)
- B-FI 0.999 (Excellent): High internal quality despite small size
- B-FV 0.128 (Fair): Minimal categorical complexity (5 categorical cols)

Conformance: SDCF-A-Bronze (clean pass)

Low B-PRS Interpretation: Small, text-heavy datasets score low on uniqueness component (each record *expected* to be unique in text content → no record is *unexpectedly* unique). Framework correctly handles small dataset edge case. However, Appendix C.2 guidance: “Small datasets (< 1000 records) require human review regardless of B-PRS.”

18.4 G.4 Reproducibility Information

18.4.1 Code Availability

Complete validation code available in ancillary files:

Primary Assessment Script: `bronze_retrospective.py` — Implements Bronze Tier assessment for all 10 datasets

Dataset Access: `validation_datasets.py` — Download scripts for all datasets from HuggingFace, GitHub, SDV library

Reproduction Example: `examples/bronze_validation_example.py` — Step-by-step walk-through

Regression Tests: `tests/test_bronze_validation.py` — Verify results match expected values

Licence: MIT Licence (code), CC BY-SA 4.0 (framework methodology)

18.4.2 Data Access

All datasets publicly available:

- **HuggingFace Hub (4 datasets):** PLEIAs SYNTH, Gretel Safety, Jupyter Agent, plus source for SDV Adult
- **SDV Library (1 dataset):** Adult Income (real source for self-synthesis)
- **MostlyAI GitHub (2 datasets):** Census, CDNOW Purchases
- **Government/Public (2 datasets):** CMS SynPUF, Census SynLBD (demo variants generated for validation)
- **Self-synthesized (3 datasets):** SDV Adult variants (GaussianCopula, CTGAN, TVAE)

Access Instructions: See `validation_datasets.py` for complete download and setup procedures.

18.4.3 Computational Requirements

Hardware:

- CPU: Intel i7 or equivalent (8+ cores recommended)
- RAM: 16 GB minimum (for 69K record dataset)
- Storage: 500 MB for all datasets
- GPU: Not required (CPU-only assessment)

Software:

- Python 3.11+ (tested on 3.11.5)
- pandas 2.1.0
- numpy 1.24.0
- scikit-learn 1.3.0
- scipy 1.11.0

Execution Time:

- Single dataset: 30 seconds to 3 minutes (depends on size)
- All 10 datasets: 14.3 minutes total
- Expected runtime: < 20 minutes on recommended hardware

Operating System: Tested on Windows 10, should work on Linux/macOS (Python cross-platform)

18.4.4 Expected Results

Reproducibility Verification:

Running `bronze_retrospective.py` should produce results matching Table 38 within numerical precision (± 0.001 for floating-point metrics).

If results differ:

- **LOF component:** Local Outlier Factor may vary slightly due to random initialization (set `random_state=42` for exact reproduction)
- **Dataset updates:** HuggingFace datasets may be updated; use specific commit hashes in `validation_datasets.py`
- **Library versions:** Ensure exact versions from `requirements.txt`

Validation Test Suite: Run `pytest tests/test_bronze_validation.py` to verify implementation correctness against expected results.

18.5 G.5 Framework Alignment Evidence

Table 42 provides detailed evidence for framework prediction accuracy (Section 7.5).

Summary: All 14 testable framework predictions matched observed results. No contradictions or unexpected behaviors detected. Validation provides empirical support for Bronze Tier methodology design.

End of Appendix G

For complete validation companion technical report (20–30 pages with extended discussion), see ancillary file: `bronze_validation_report.pdf`

Table 42: Detailed Framework Alignment Evidence

Prediction	Source	Predicted	Observed	Match
<i>B-PRS Discrimination and Patterns</i>				
Min discrimination range	App C.2	> 2x	8.8x	✓
Bronze rarely Low (< 0.30)	App C.2	< 20%	10%	✓
Avg Moderate-High	§3.3	0.40–0.60	0.516	✓
Demographic = High	App C.2	0.50–0.70	0.534–0.639	✓
Simple schemas lower	App C.2	< 0.40	0.360–0.390	✓
<i>B-FI Performance</i>				
Modern tools > 0.90	App C.3	> 90%	100%	✓
Average modern synthesis	App C.3	0.90–0.95	0.991	✓
Low variability (consistent)	App C.3	Std < 0.05	Std 0.021	✓
<i>B-FV Patterns</i>				
High uncertainty	App C.4	Variable	0.100–0.909	✓
Demographic high B-FV	App C.4	> 0.50	0.548–0.724	✓
Simple schemas baseline	App C.4	≈ 0.10	0.100 (5 datasets)	✓
Bimodal distribution	App C.4	Mentioned	Confirmed	✓
<i>Conformance Distribution</i>				
Majority Restricted	§3.3	50–70%	60%	✓
Conservative bias	§3.3	High R rate	Confirmed	✓
B-FV > 0.30 triggers R	§3.4	Yes	83% of R	✓
B-PRS > 0.70 triggers R	§3.4	Yes	17% of R	✓
<i>Cross-Domain Validity</i>				
Domain-agnostic methodology	§3.3	Yes	7 domains	✓
Domain-logical results	§3.3	Yes	Patterns interp.	✓

References

- [1] Martín Abadi, Andy Chu, Ian Goodfellow, H Brendan McMahan, Ilya Mironov, Kunal Talwar, and Li Zhang. Deep learning with differential privacy. In *ACM SIGSAC Conference on Computer and Communications Security (CCS)*, pages 308–318, 2016. doi: 10.1145/2976749.2978318.
- [2] Alejandro Barredo Arrieta, Natalia Díaz-Rodríguez, Javier Del Ser, Adrien Bennetot, Siham Tabik, Alberto Barbado, Salvador García, Sergio Gil-López, Daniel Molina, Richard Benjamins, Raja Chatila, and Francisco Herrera. Explainable artificial intelligence (xai): Concepts, taxonomies, opportunities and challenges toward responsible ai. *Information Fusion*, 58:82–115, 2020. doi: 10.1016/j.inffus.2019.12.012.
- [3] Article 29 Data Protection Working Party. Opinion 05/2014 on anonymisation techniques. Technical Report 0829/14/EN WP216, European Commission, 2014.
- [4] Solon Barocas, Moritz Hardt, and Arvind Narayanan. *Fairness and Machine Learning: Limitations and Opportunities*. fairmlbook.org, 2019. URL <https://fairmlbook.org>.
- [5] Joy Buolamwini and Timnit Gebru. Gender shades: Intersectional accuracy disparities in commercial gender classification. In *ACM Conference on Fairness, Accountability and Transparency (FAT*)*, pages 77–91, 2018.
- [6] Nicholas Carlini, Steve Chien, Milad Nasr, Shuang Song, Andreas Terzis, and Florian Tramèr. Membership inference attacks from first principles. In *IEEE Symposium on Security and Privacy (S&P)*, pages 1897–1914, 2022.
- [7] Richard J Chen, Ming Y Lu, Tiffany Y Chen, Drew F K Williamson, and Faisal Mahmood. Synthetic data in machine learning for medicine and healthcare. *Nature Biomedical Engineering*, 5(6):493–497, 2021. doi: 10.1038/s41551-021-00751-8.
- [8] Edward Choi, Siddharth Biswal, Bradley Malin, Jon Duke, Walter F Stewart, and Jimeng Sun. Generating multi-label discrete patient records using generative adversarial networks. In *Machine Learning for Healthcare Conference (MLHC)*, pages 286–305, 2017.
- [9] DataCebo. Sdmetrics: Metrics for synthetic data evaluation. <https://github.com/sdv-dev/SDMetrics>, 2021.
- [10] Cynthia Dwork and Aaron Roth. *The Algorithmic Foundations of Differential Privacy*, volume 9 of *Foundations and Trends in Theoretical Computer Science*. Now Publishers, 2014. doi: 10.1561/0400000042.
- [11] Cynthia Dwork, Frank McSherry, Kobbi Nissim, and Adam Smith. Calibrating noise to sensitivity in private data analysis. In *Theory of Cryptography Conference (TCC)*, pages 265–284, 2006.
- [12] Cynthia Dwork, Moritz Hardt, Toniann Pitassi, Omer Reingold, and Richard Zemel. Fairness through awareness. In *Innovations in Theoretical Computer Science Conference (ITCS)*, pages 214–226, 2012.
- [13] Khaled El Emam, Lucy Mosquera, and Richard Hoptroff. *Practical Synthetic Data Generation: Balancing Privacy and the Broad Availability of Data*. O’Reilly Media, 2020. ISBN 9781492072744.

- [14] European Parliament and Council. General data protection regulation (gdpr). Regulation (EU) 2016/679, 2016.
- [15] European Parliament and Council. Regulation on artificial intelligence (artificial intelligence act). Regulation (EU) 2024/1689, 2024.
- [16] Timnit Gebru, Jamie Morgenstern, Briana Vecchione, Jennifer Wortman Vaughan, Hanna Wallach, Hal Daumé III, and Kate Crawford. Datasheets for datasets. *Communications of the ACM*, 64(12):86–92, 2021. doi: 10.1145/3458723.
- [17] Ian Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. Generative adversarial nets. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 27, 2014.
- [18] Moritz Hardt, Eric Price, and Nati Srebro. Equality of opportunity in supervised learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 29, 2016.
- [19] ISO/IEC JTC 1/SC 27. Iso/iec 27001:2022 information security management systems – requirements. Technical report, International Organization for Standardization, 2022.
- [20] ISO/IEC JTC 1/SC 42. Iso/iec 42001:2023 information technology – artificial intelligence – management system. Technical report, International Organization for Standardization, 2023.
- [21] Diederik P Kingma and Max Welling. Auto-encoding variational bayes. *arXiv preprint arXiv:1312.6114*, 2013.
- [22] Jon Kleinberg, Sendhil Mullainathan, and Manish Raghavan. Inherent trade-offs in the fair determination of risk scores. In *8th Innovations in Theoretical Computer Science Conference (ITCS)*, volume 67 of *Leibniz International Proceedings in Informatics (LIPIcs)*, pages 43:1–43:23, 2017. doi: 10.4230/LIPIcs.ITCS.2017.43.
- [23] Margaret Mitchell, Simone Wu, Andrew Zaldivar, Parker Barnes, Lucy Vasserman, Ben Hutchinson, Elena Spitzer, Inioluwa Deborah Raji, and Timnit Gebru. Model cards for model reporting. In *ACM Conference on Fairness, Accountability, and Transparency (FAT*)*, pages 220–229, 2019.
- [24] Milad Nasr, Reza Shokri, and Amir Houmansadr. Comprehensive privacy analysis of deep learning: Passive and active white-box inference attacks against centralized and federated learning. In *IEEE Symposium on Security and Privacy (S&P)*, pages 739–753, 2019.
- [25] National Institute of Standards and Technology. Artificial intelligence risk management framework (ai rmf 1.0). Nist ai 100-1, NIST, 2023.
- [26] Nicolas Papernot, Patrick McDaniel, Arunesh Sinha, and Michael P Wellman. Sok: Security and privacy in machine learning. In *IEEE European Symposium on Security and Privacy (EuroS&P)*, pages 399–414, 2018.
- [27] Neha Patki, Roy Wedge, and Kalyan Veeramachaneni. The synthetic data vault. In *IEEE International Conference on Data Science and Advanced Analytics (DSAA)*, pages 399–410, 2016.
- [28] Reza Shokri, Marco Stronati, Congzheng Song, and Vitaly Shmatikov. Membership inference attacks against machine learning models. In *IEEE Symposium on Security and Privacy (S&P)*, pages 3–18, 2017.

- [29] UK Information Commissioner’s Office. Guidance on ai and data protection. <https://ico.org.uk/for-organisations/guide-to-data-protection/key-dp-themes/guidance-on-ai-and-data-protection/>, 2020.
- [30] Lei Xu, Maria Skoularidou, Alfredo Cuesta-Infante, and Kalyan Veeramachaneni. Modeling tabular data using conditional gan. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 32, 2019.
- [31] Andrew Yale, Saloni Dash, Ritik Dutta, Isabelle Guyon, Adrien Pavao, and Kristin P Bennett. Generation and evaluation of privacy preserving synthetic health data. *Neurocomputing*, 416: 244–255, 2020. doi: 10.1016/j.neucom.2019.12.136.
- [32] Samuel Yeom, Irene Giacomelli, Matt Fredrikson, and Somesh Jha. Privacy risk in machine learning: Analyzing the connection to overfitting. In *IEEE Computer Security Foundations Symposium (CSF)*, pages 268–282, 2018.
- [33] Jinsung Yoon, Daniel Jarrett, and Mihaela Van der Schaar. Time-series generative adversarial networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 32, 2019.
- [34] Jun Zhang, Graham Cormode, Cecilia M Procopiuc, Divesh Srivastava, and Xiaokui Xiao. Privbayes: Private data release via bayesian networks. *ACM Transactions on Database Systems (TODS)*, 42(4):25:1–25:41, 2017. doi: 10.1145/3134428.